

# Boosting algorithms

---

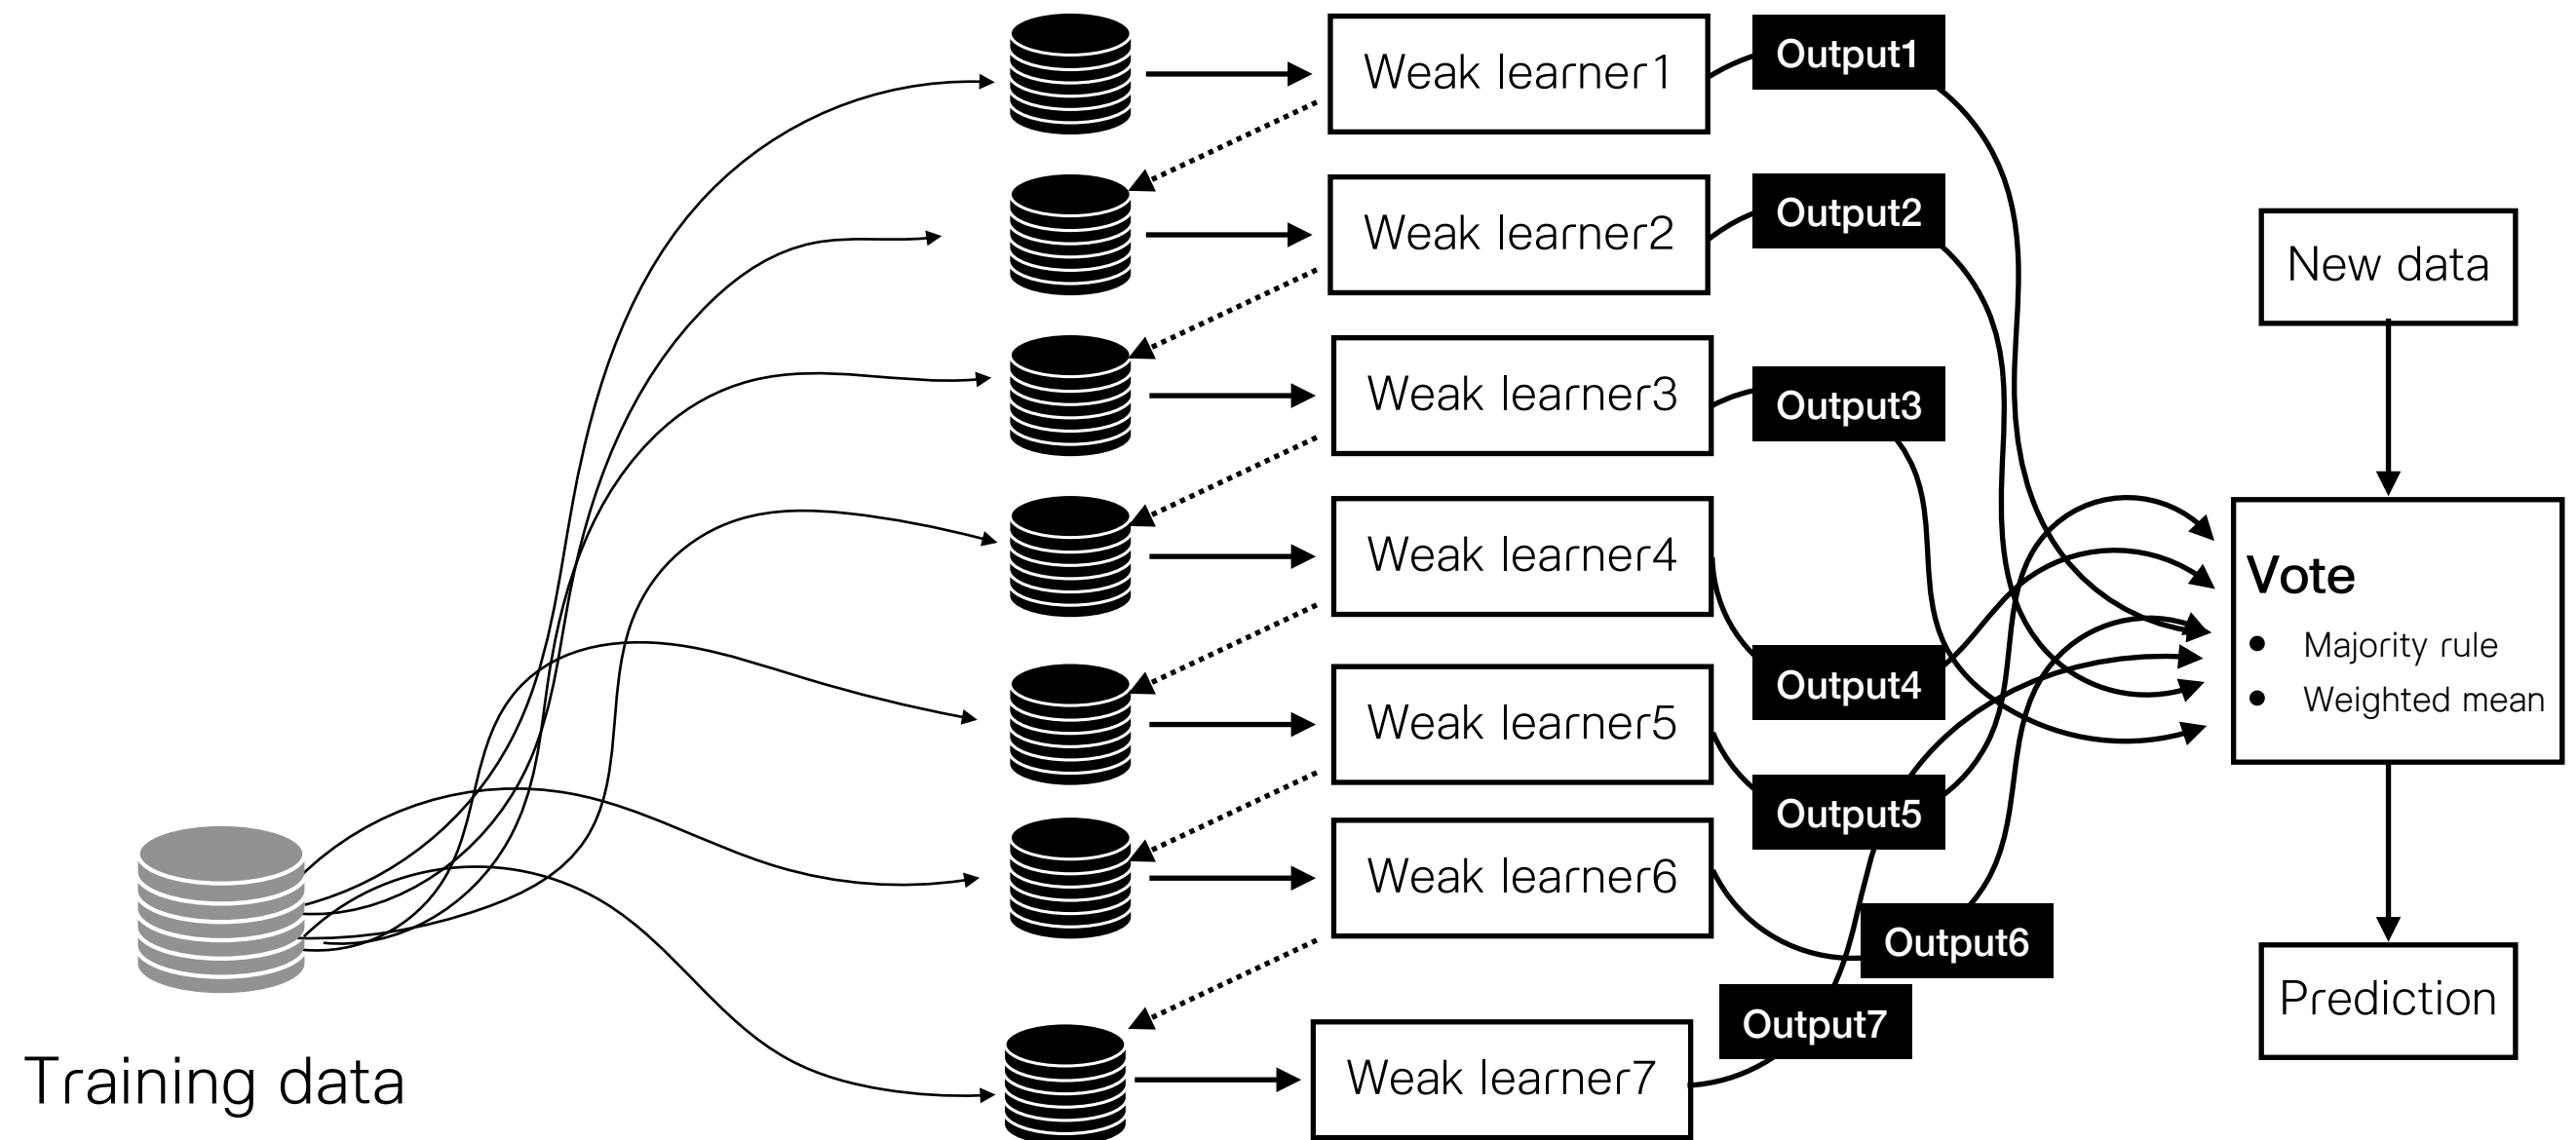
What's Boosting in ML?

How Boosting algorithms work?

Types of Boosting algorithm.

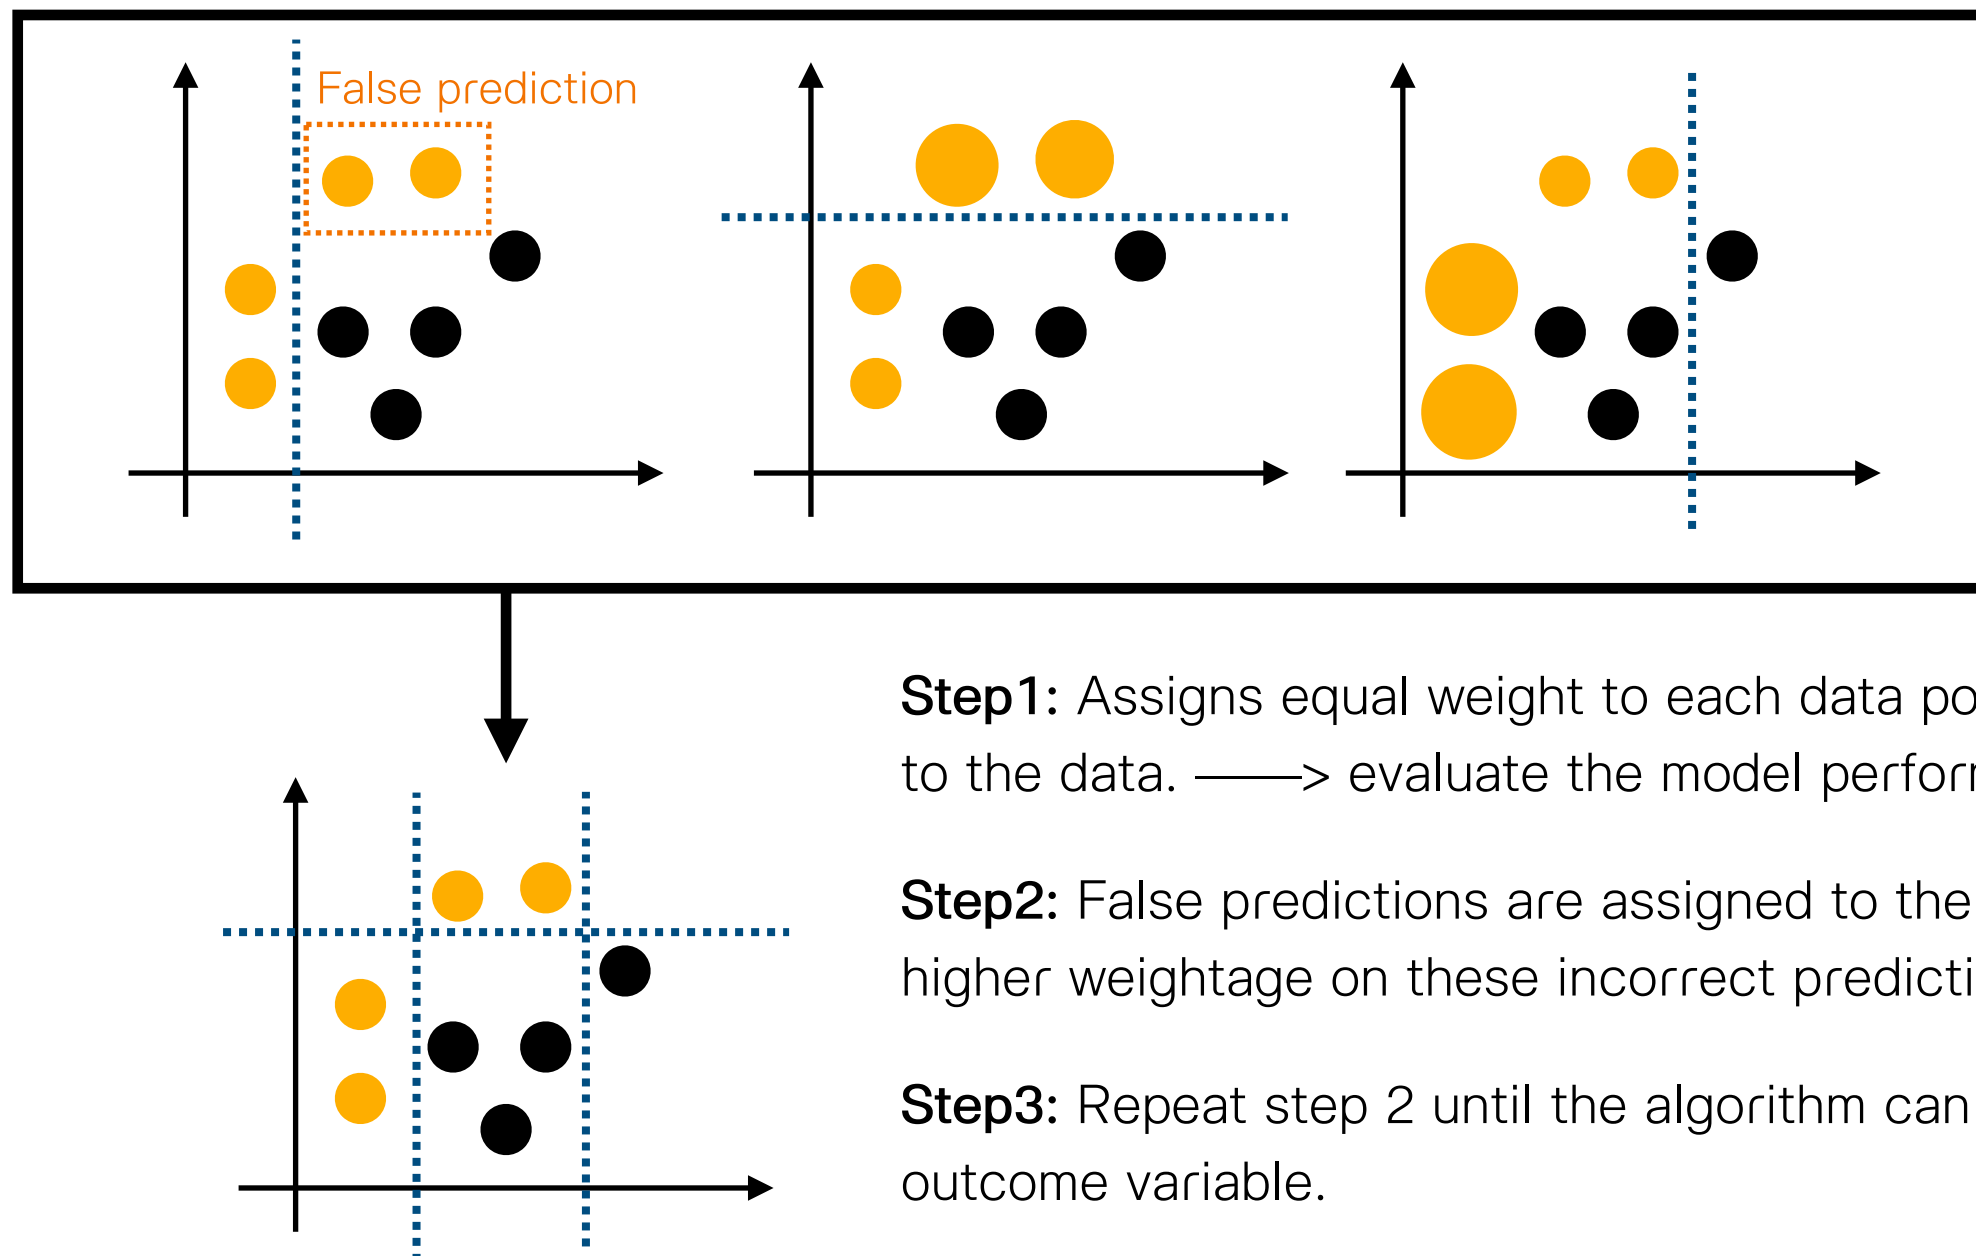
# Boosting algorithms

- Boosting is an ensemble learning techniques that uses a set of ML algorithms to **combines weak learner** (or base learner) to **form a strong learner** in order to increase the accuracy of the model.
- During the training phase, the performance of the model is improved by assigning a higher weight to the previous incorrectly classified subjects.



# How does Boosting algorithm work?

Basic concept behind the working of the boosting algorithm is to generate multiple weak learner and combine their predictions to form one strong rule



**Step1:** Assigns equal weight to each data point and fit the ML model to the data. —> evaluate the model performance

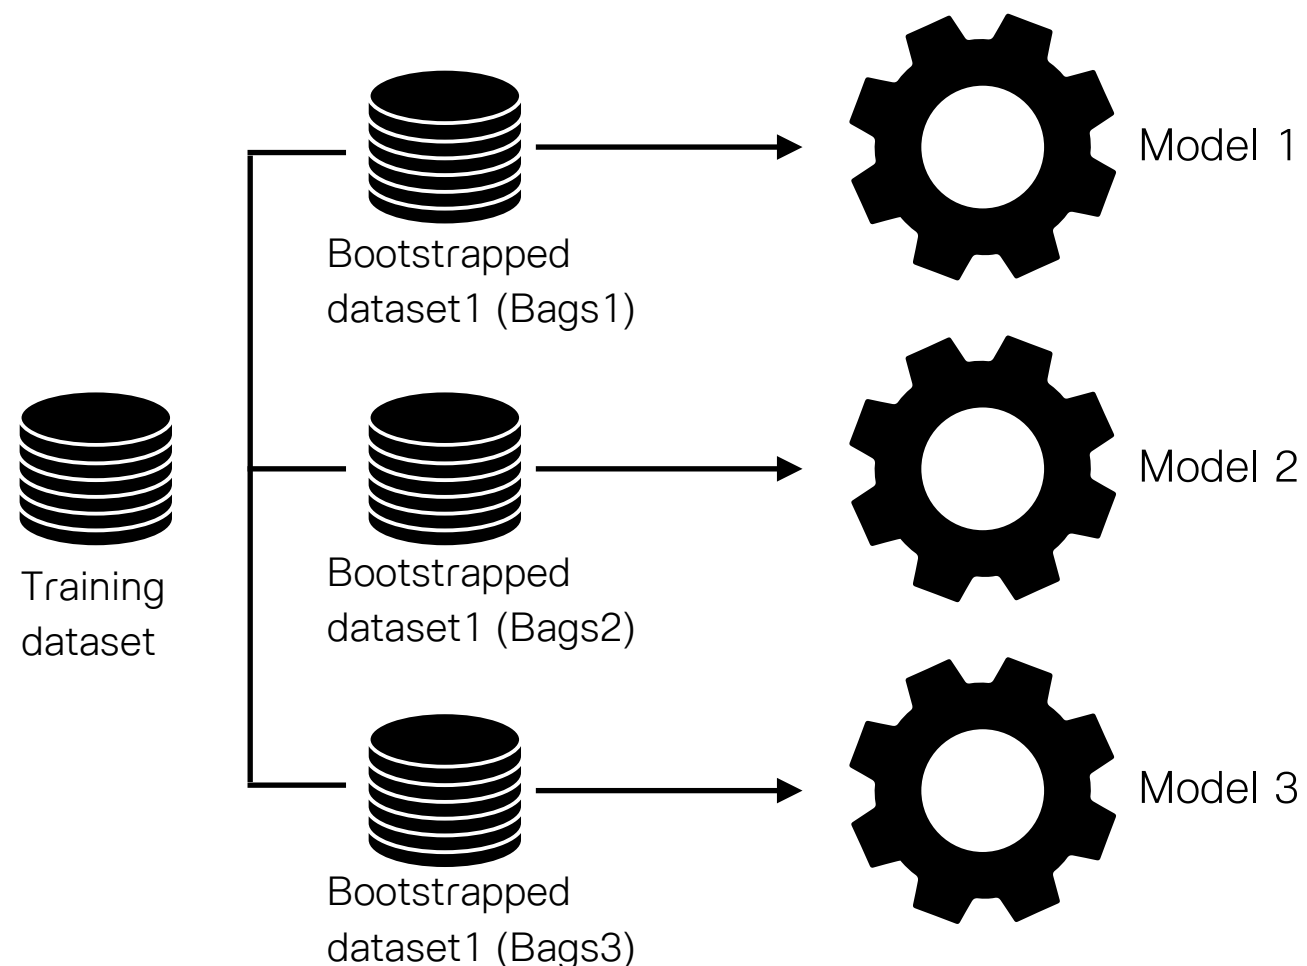
**Step2:** False predictions are assigned to the next dataset with higher weightage on these incorrect predictions.

**Step3:** Repeat step 2 until the algorithm can correctly predict the outcome variable.

# Ensemble learning

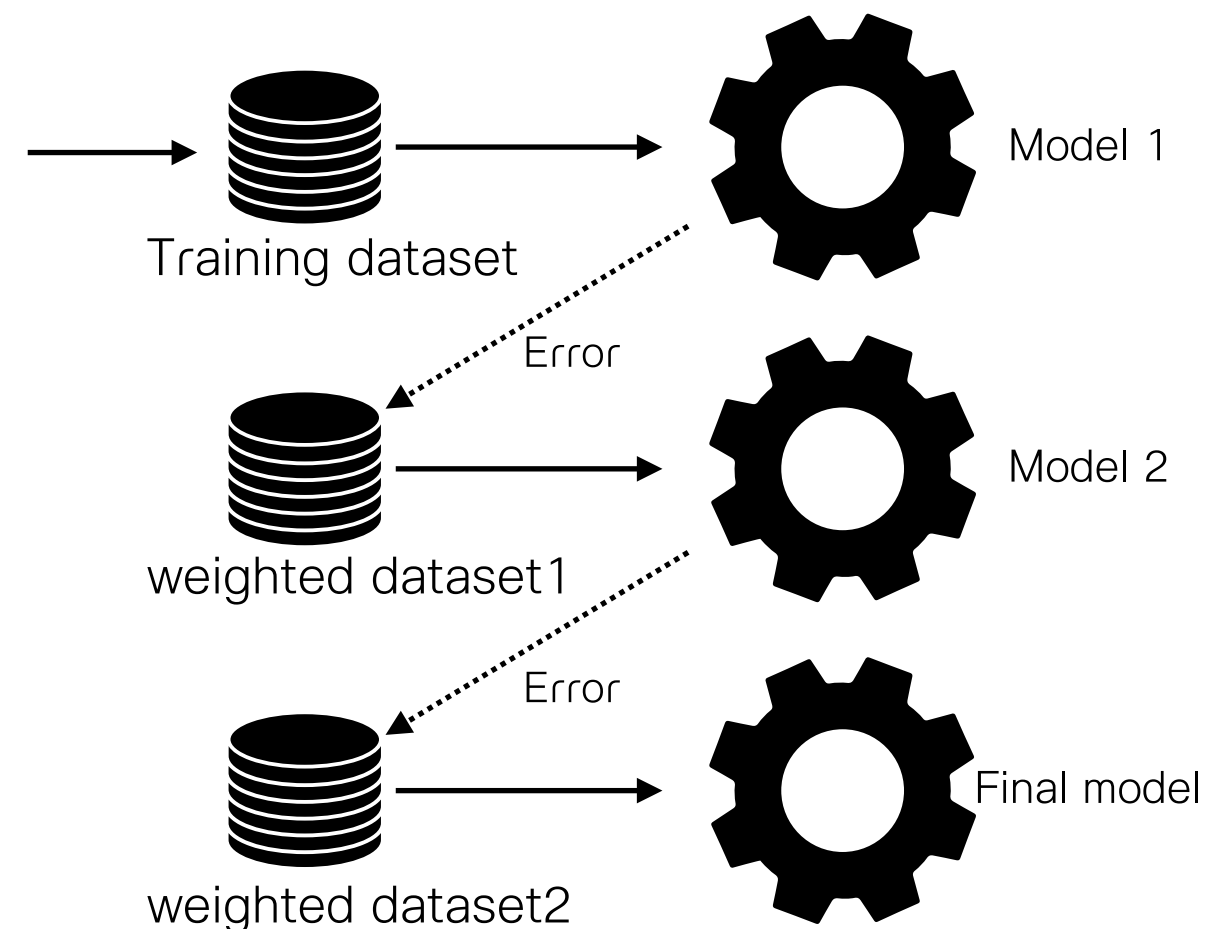
## Bagging - Parallel

- Build multiple ML models using same algorithm with subset of training dataset randomly selected from the full training dataset



## Boosting - Sequential

- Perform iterative process to reduce errors of previous models by selecting points which give wrong predictions and try to predict them with successive model.



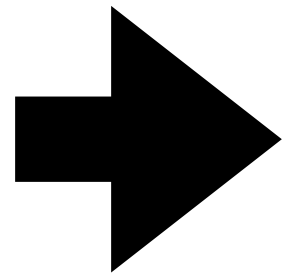
# Types of Boosting algorithms

1. AdaBoost (Adaptive Boosting)
2. Gradient Boosting
3. XGBoost

# AdaBoost

# AdaBoost (Adaptive boosting)

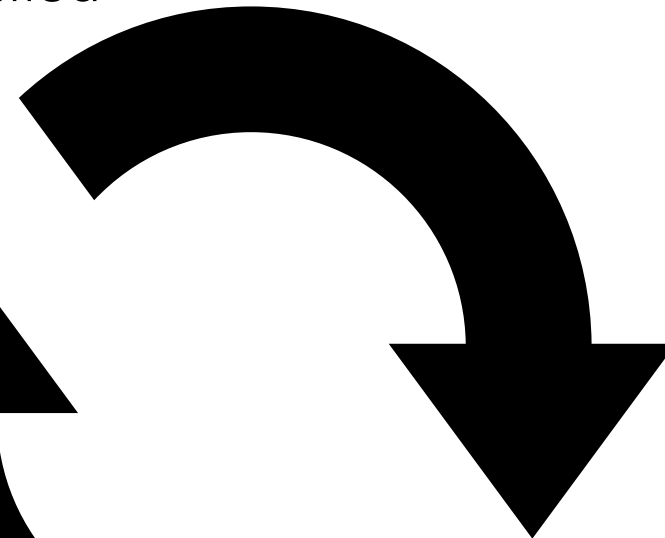
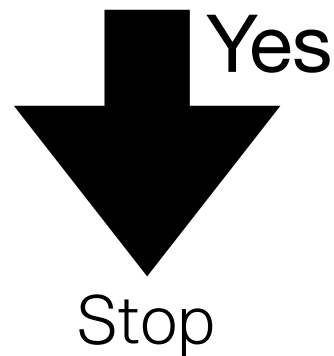
**Step1:** Each data point is weighted equally for the first decision stump.



**Step2:** Mis-classified data points are assigned higher weights

No

Is all data points fall into the correct class?



**Step3:** A new decision stump is drawn by considering the data points with higher weights as more significant.

**Note:** AdaBoost can also be used on a regressions problem, but it's most commonly seen in classification problems.

<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.31.314&rep=rep1&type=pdf>

<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.100.4560&rep=rep1&type=pdf>

# Original AdaBoost algorithm

1. Given a training dataset  $(X, y)$  which contains data points  $(x_1, y_1), (x_2, y_2), \dots, (x_m, y_m)$
2. Construct distribution of  $W_t$  on  $\{1, 2, 3, \dots, m\}$ , where  $W_t(i)$  is the weight attributed to subject  $i$  on the iteration  $t$
3. Initial weight are calculated by  $W_1(i) = 1/m$ , and the next weight are calculated by

$$W_{t+1}(i) = W_t(i)F_i/Z_t$$

- Where  $F_i = e^{-\alpha_t}$  if  $M_t(i) \neq y_i$ , and  $F_i = e^{\alpha_t}$  if  $M_t(i) = y_i$
- $\epsilon_t = \sum_{i=1}^m W_t(i) \times \delta_i$  where  $\delta_i = 0$  if  $M_t(i) = y_i$ , and  $\delta_i = 1$  if  $M_t(i) \neq y_i$   

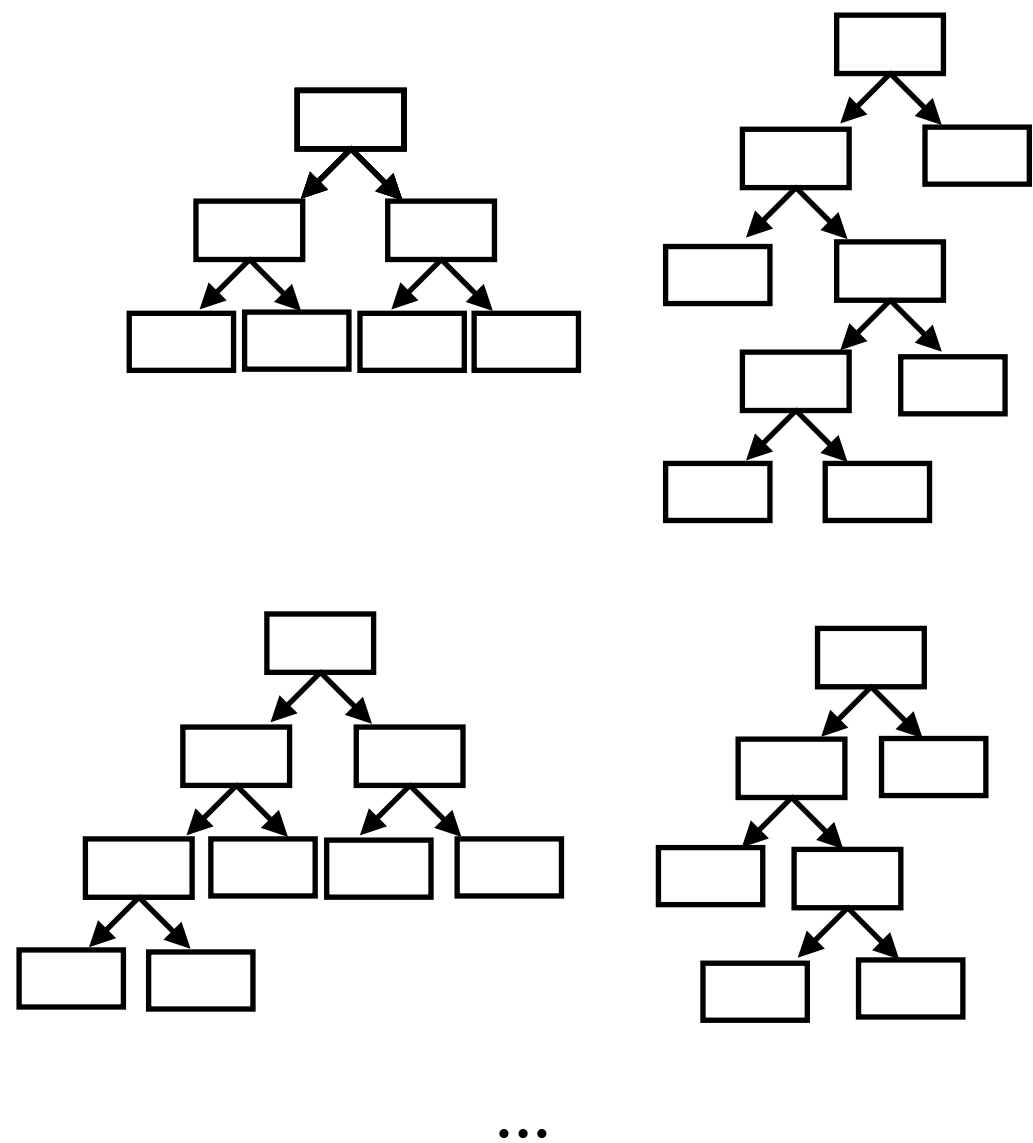
Correct classifiedMis-classified
- $\alpha_t = \frac{1}{2} \ln\left(\frac{1 - \epsilon_t}{\epsilon_t}\right)$

4. Build up a classifier (weak learner)  $M_t : X \rightarrow \{-1, 1\}$
5. The final classification considered using the weighted average of the classifier

$$\text{sign}\left(\sum_{t=1}^T \alpha_t M_t(x)\right)$$

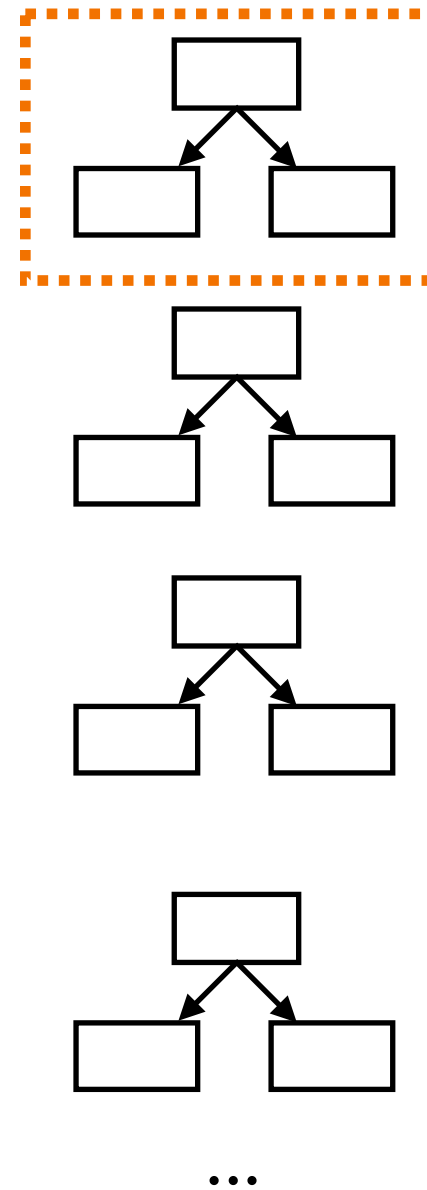


# Main concepts behind AdaBoost



Random forest

Called “stump”



AdaBoost

- Technically, stumps are weak learner, that is not great at making accurate classifications. **AdaBoost** combined a lot of weak learners to make classification.
- In a random forest, each decision tree is made independently of the others. In contrast, Forest of stumps made with **Adaboost**, **order is matter**. (the error that the first stump makes influence how the second stump is made, and .... so on ...)
- Hence some stump will get more weight in the classification than others.

STA	LOC	SYS	AGE
Dead	No	36	27
Dead	No	48	59
Dead	Yes	44	77
Dead	No	62	54
Alive	Yes	112	87
Alive	Yes	108	69
Alive	No	140	63
Alive	Yes	138	30

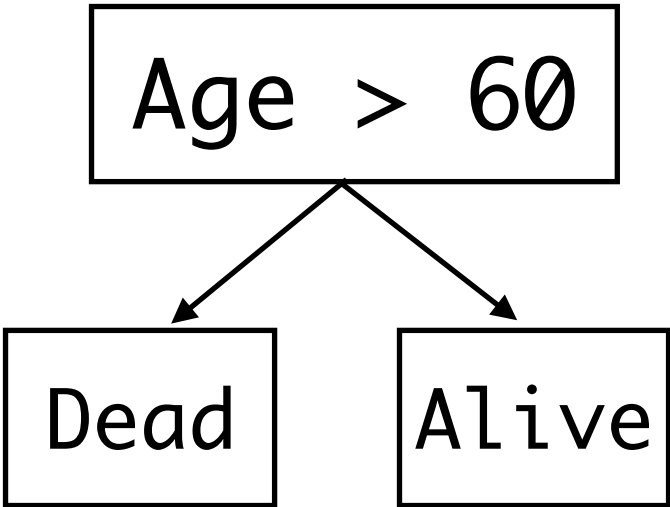
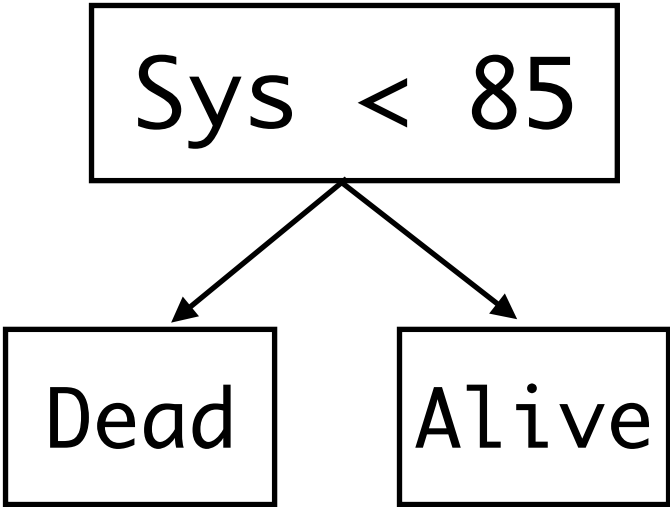
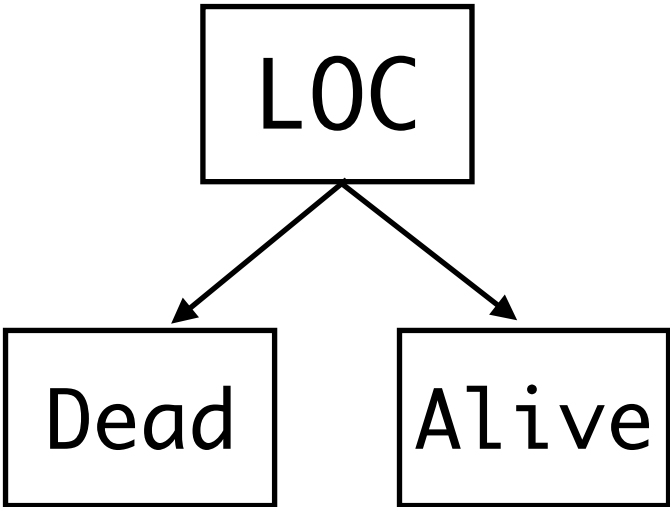
- STA - life status
- LOC - level of consciousness
- Sys - systolic blood pressure
- Age - patient's age

Step1: All subjects get the same weight:  $W_1(i) = 1/m = 1/8$

STA	LOC	SYS	AGE	Subject weight
Dead	No	36	27	1/8
Dead	No	48	59	1/8
Dead	Yes	44	77	1/8
Dead	No	62	54	1/8
Alive	Yes	112	87	1/8
Alive	Yes	108	69	1/8
Alive	No	140	63	1/8
Alive	Yes	138	30	1/8

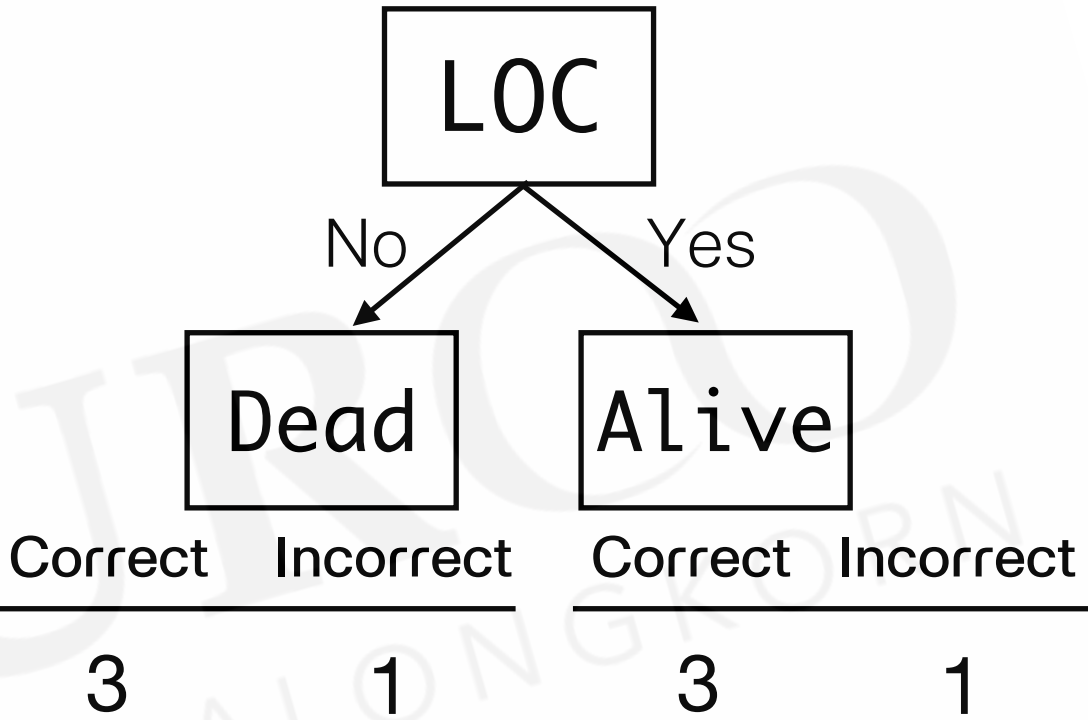
Step2: Build up the first stump

STA	LOC	SYS	AGE	Subject weight
Dead	No	36	27	1/8
Dead	No	48	59	1/8
Dead	Yes	44	77	1/8
Dead	No	62	54	1/8
Alive	Yes	112	87	1/8
Alive	Yes	108	69	1/8
Alive	No	140	63	1/8
Alive	Yes	138	30	1/8



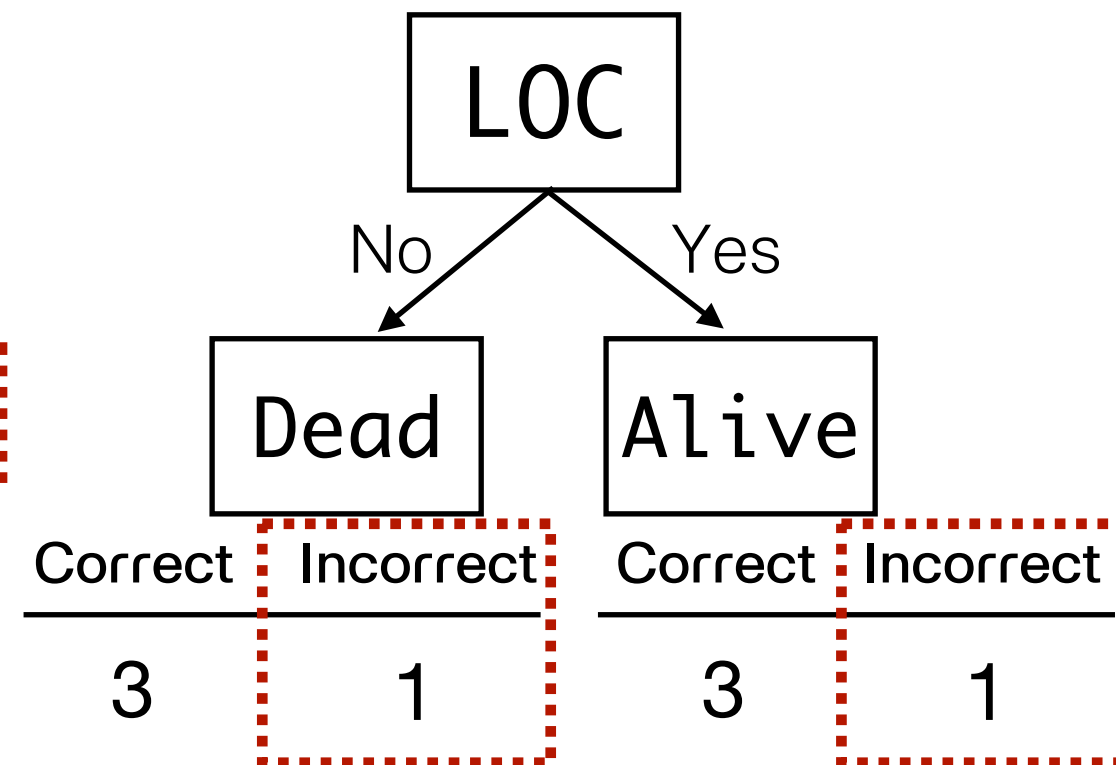
Step2: Build up the first stump

STA	LOC	SYS	AGE	Subject weight
Dead	No	36	27	1/8
Dead	No	48	59	1/8
Dead	Yes	44	77	1/8
Dead	No	62	54	1/8
Alive	Yes	112	87	1/8
Alive	Yes	108	69	1/8
Alive	No	140	63	1/8
Alive	Yes	138	30	1/8



### Step3: Evaluate the prediction error of the stump

STA	LOC	SYS	AGE	Subject weight
Dead	No	36	27	1/8
Dead	No	48	59	1/8
Dead	Yes	44	77	1/8
Dead	No	62	54	1/8
Alive	Yes	112	87	1/8
Alive	Yes	108	69	1/8
Alive	No	140	63	1/8
Alive	Yes	138	30	1/8



$$\text{Total Error} = \epsilon_t = \sum_{i=1}^m W_t(i) \times \delta_i = (1/8)(6)(0) + (1/8)(2)(1) = 2/8$$

where  $\delta_i = 0$  if  $M_t(i) = y_i$ , and  $\delta_i = 1$  if  $M_t(i) \neq y_i$

Step4: Calculate **the amount of say** (weight:  $\alpha_t$ ) of the stump

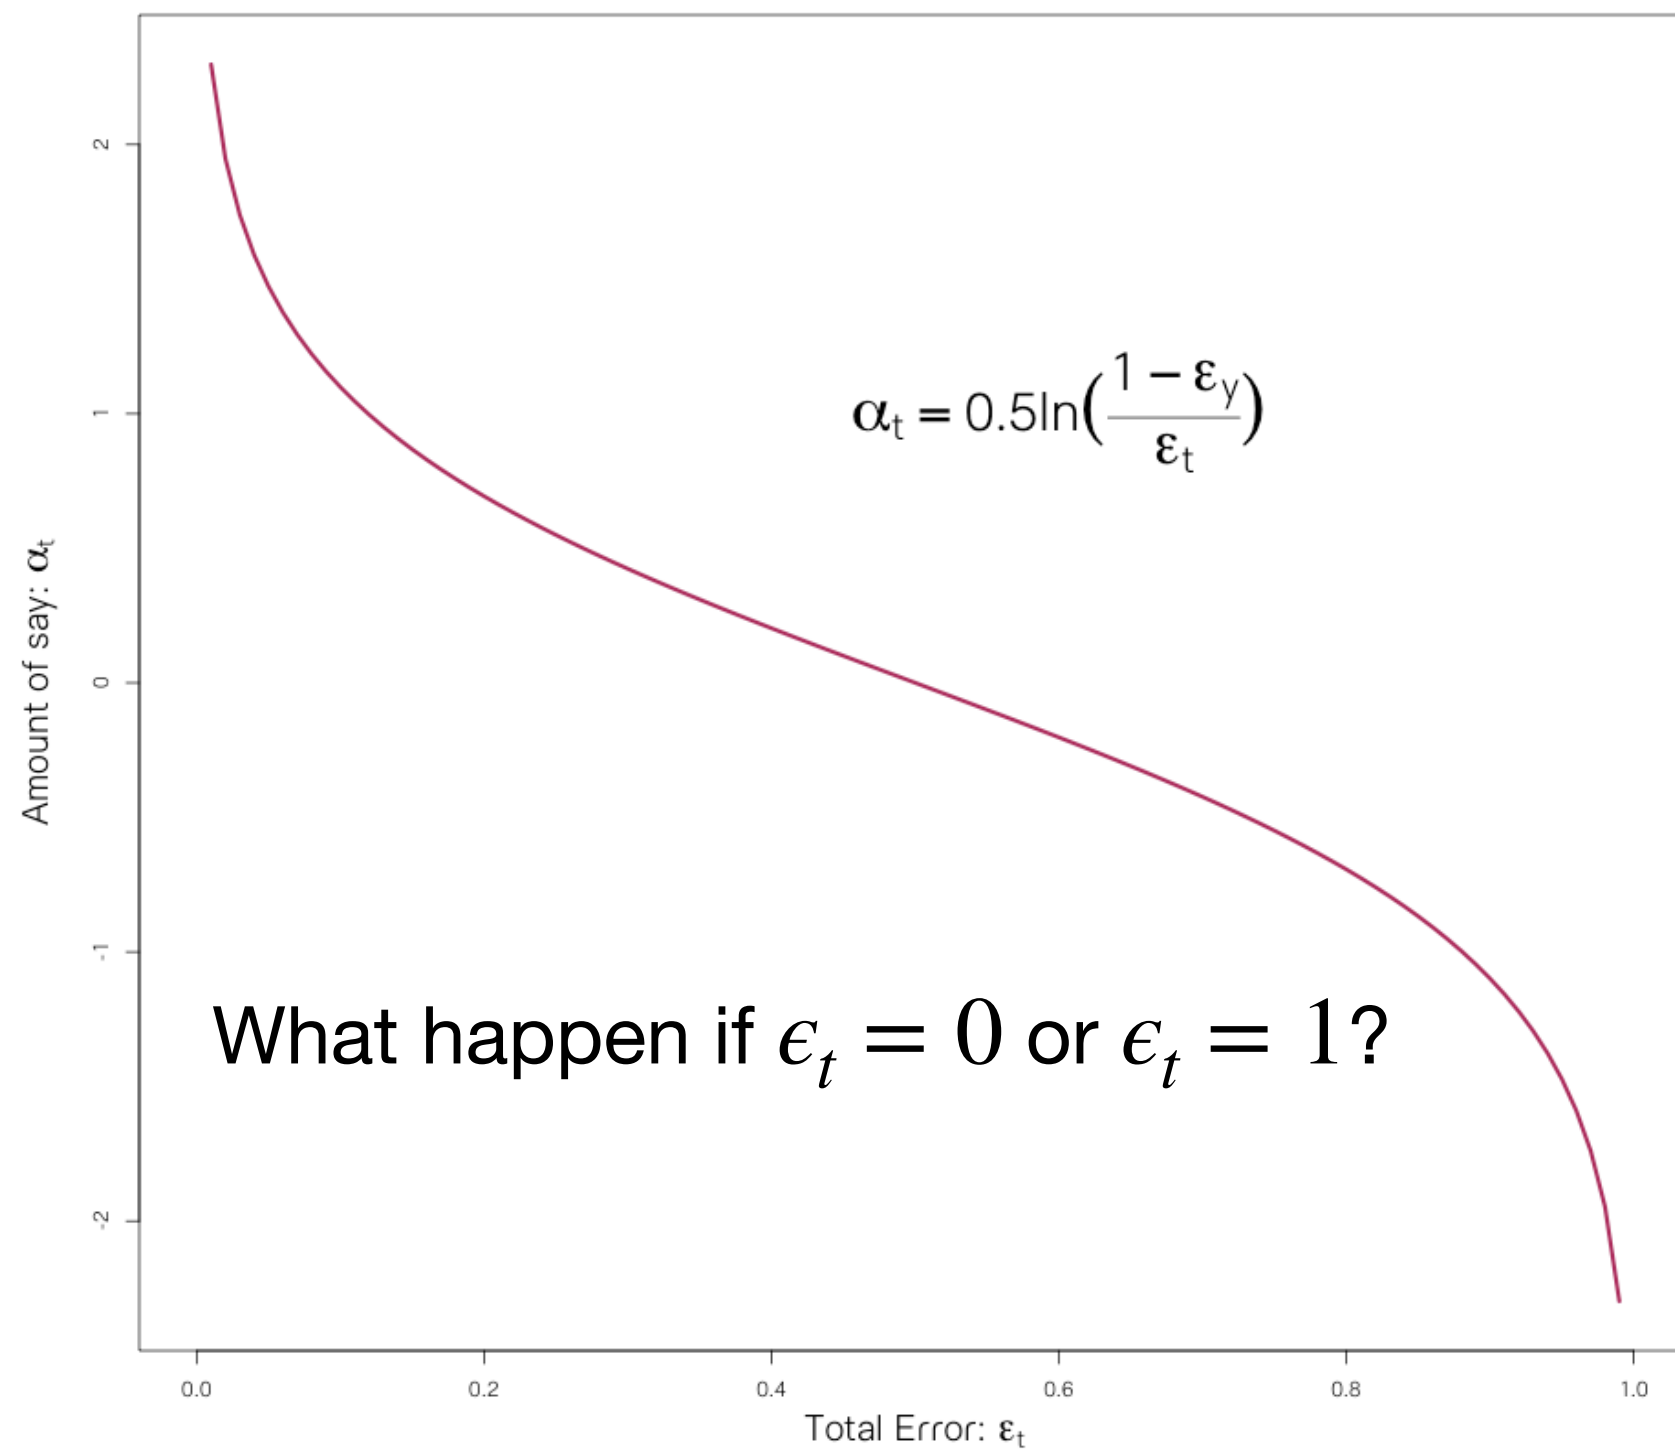
STA	LOC	SYS	AGE	Subject weight
Dead	No	36	27	1/8
Dead	No	48	59	1/8
Dead	Yes	44	77	1/8
Dead	No	62	54	1/8
Alive	Yes	112	87	1/8
Alive	Yes	108	69	1/8
Alive	No	140	63	1/8
Alive	Yes	138	30	1/8

$$\alpha_t = \frac{1}{2} \ln\left(\frac{1 - \epsilon_t}{\epsilon_t}\right)$$

$$\text{Total Error} = \epsilon_t = \sum_{i=1}^m W_t(i) \times \delta_i = (1/8)(6)(0) + (1/8)(2)(1) = 2/8$$

where  $\delta_i = 0$  if  $M_t(i) = y_i$ , and  $\delta_i = 1$  if  $M_t(i) \neq y_i$

**The amount of say:**  $\alpha_t = \frac{1}{2} \ln\left(\frac{1 - \epsilon_t}{\epsilon_t}\right)$





Step4: Calculate **the amount of say** (weight:  $\alpha_t$ ) of the stump

STA	LOC	SYS	AGE	Subject weight
Dead	No	36	27	1/8
Dead	No	48	59	1/8
Dead	Yes	44	77	1/8
Dead	No	62	54	1/8
Alive	Yes	112	87	1/8
Alive	Yes	108	69	1/8
Alive	No	140	63	1/8
Alive	Yes	138	30	1/8

$$\begin{aligned}
 \alpha_t &= \frac{1}{2} \ln\left(\frac{1 - \epsilon_t}{\epsilon_t}\right) \\
 &= \frac{1}{2} \ln\left(\frac{1 - 2/8}{2/8}\right) \\
 &= 0.549
 \end{aligned}$$

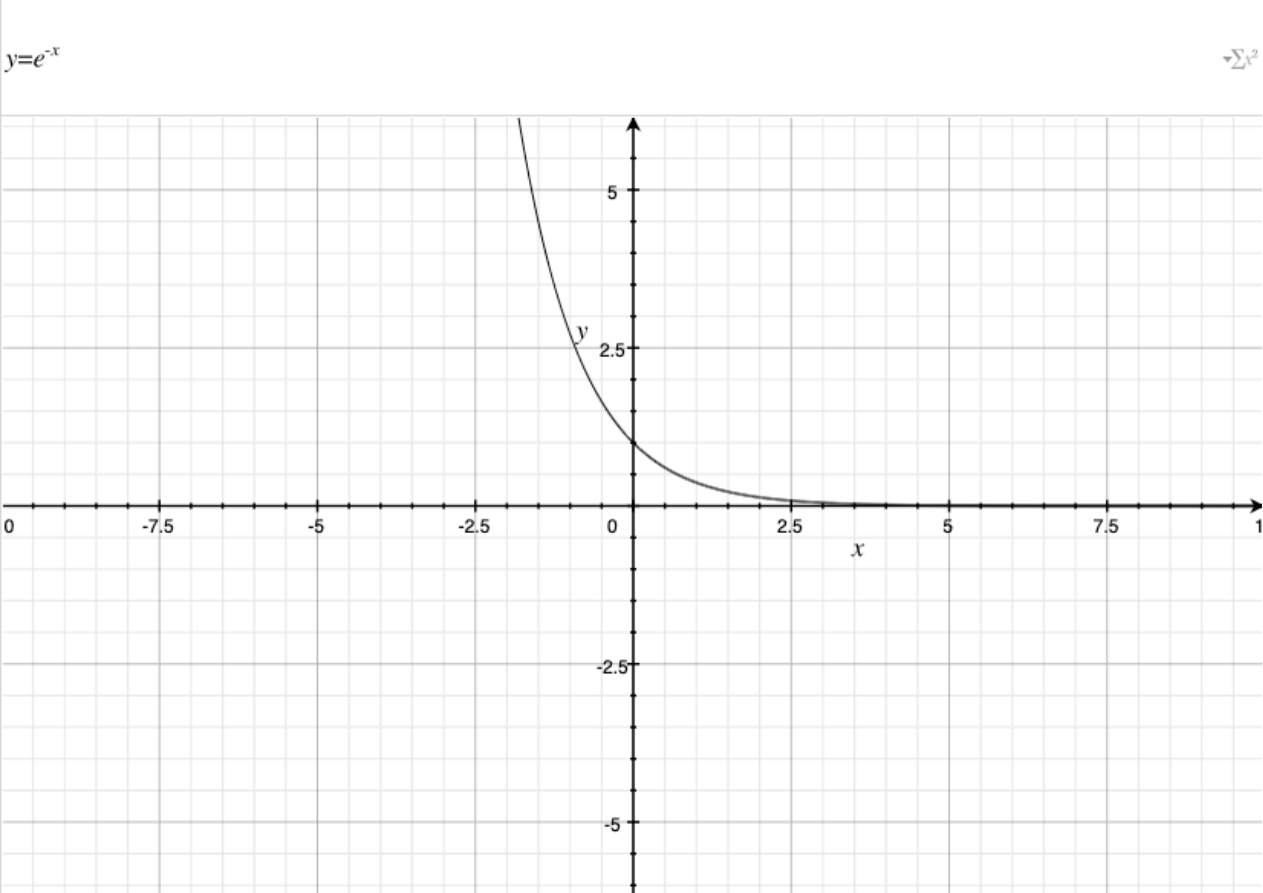
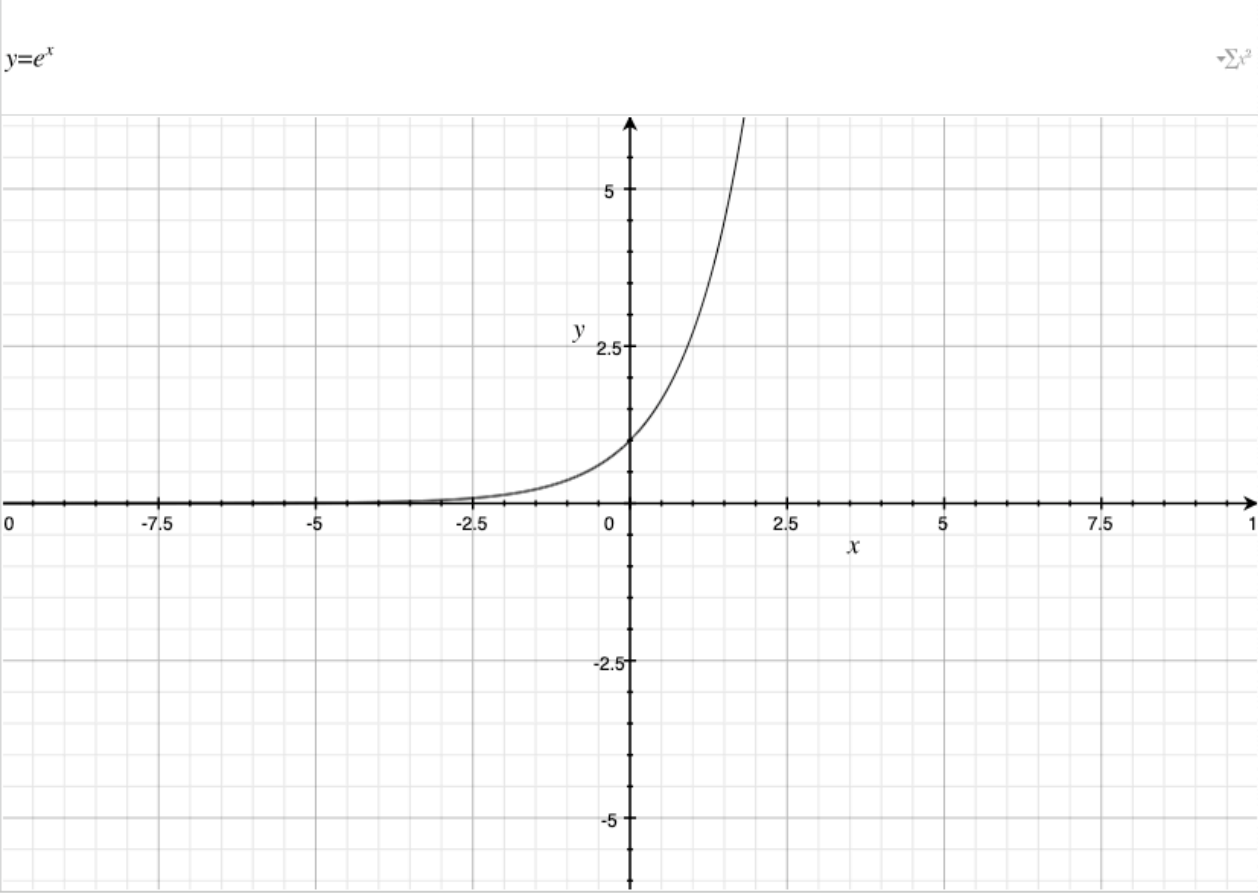
$$\text{Total Error} = \epsilon_t = \sum_{i=1}^m W_t(i) \times \delta_i = (1/8)(6)(0) + (1/8)(2)(1) = 2/8$$

where  $\delta_i = 0$  if  $M_t(i) = y_i$ , and  $\delta_i = 1$  if  $M_t(i) \neq y_i$

Get back to the step1: Reweight the subjects using previous stump.

STA	LOC	SYS	AGE	Subject weight	New weight*
Dead	No	36	27	1/8	0.072
Dead	No	48	59	1/8	0.072
Dead	Yes	44	77	1/8	0.216
Dead	No	62	54	1/8	0.072
Alive	Yes	112	87	1/8	0.072
Alive	Yes	108	69	1/8	0.072
Alive	No	140	63	1/8	0.216
Alive	Yes	138	30	1/8	0.072

$$\begin{aligned}
 \text{New weight}^* &= W_{t+1}(i) = W_t(i) \times e^{\alpha_t} ; \text{ if } M_t(i) \neq y_i \\
 &= W_t(i) \times e^{-\alpha_t} ; \text{ if } M_t(i) = y_i
 \end{aligned}$$



Get back to the step1: Reweight the subjects using previous stump.

STA	LOC	SYS	AGE	Subject weight	New weight	Normalized weight
Dead	No	36	27	1/8	0.072	0.083
Dead	No	48	59	1/8	0.072	0.083
Dead	Yes	44	77	1/8	0.216	0.250
Dead	No	62	54	1/8	0.072	0.083
Alive	Yes	112	87	1/8	0.072	0.083
Alive	Yes	108	69	1/8	0.072	0.083
Alive	No	140	63	1/8	0.216	0.250
Alive	Yes	138	30	1/8	0.072	0.083

$$\begin{aligned}
 \text{New weight (final)} = W_{t+1}(i) &= W_t(i) \times \frac{e^{\alpha_t}}{Z_t} ; \text{ if } M_t(i) \neq y_i \\
 &= W_t(i) \times \frac{e^{-\alpha_t}}{Z_t} ; \text{ if } M_t(i) = y_i
 \end{aligned}$$

## Get back to the step2: Build up the next stump

STA	LOC	SYS	AGE	Adjusted subject weight
Dead	No	36	27	0.083
Dead	No	48	59	0.083
Dead	Yes	44	77	0.250
Dead	No	62	54	0.083
Alive	Yes	112	87	0.083
Alive	Yes	108	69	0.083
Alive	No	140	63	0.250
Alive	Yes	138	30	0.083

Instead of using a weighted Gini or weighted Entropy index, we can make a new collection of data points that contains duplicate copies of the data points correspond to their weights.

STA	LOC	SYS	AGE	Adjusted subject weight
Dead	No	36	27	0.083
Dead	No	48	59	0.083
Dead	Yes	44	77	0.250
Dead	No	62	54	0.083
Alive	Yes	112	87	0.083
Alive	Yes	108	69	0.083
Alive	No	140	63	0.250
Alive	Yes	138	30	0.083

Weighted  
random  
sampling with  
replacement

[illegible]

STA	LOC	SYS	AGE	Adjusted subject weight
Dead	No	36	27	0.083
Dead	No	48	59	0.083
Dead	Yes	44	77	0.250
Dead	No	62	54	0.083
Alive	Yes	112	87	0.083
Alive	Yes	108	69	0.083
Alive	No	140	63	0.250
Alive	Yes	138	30	0.083

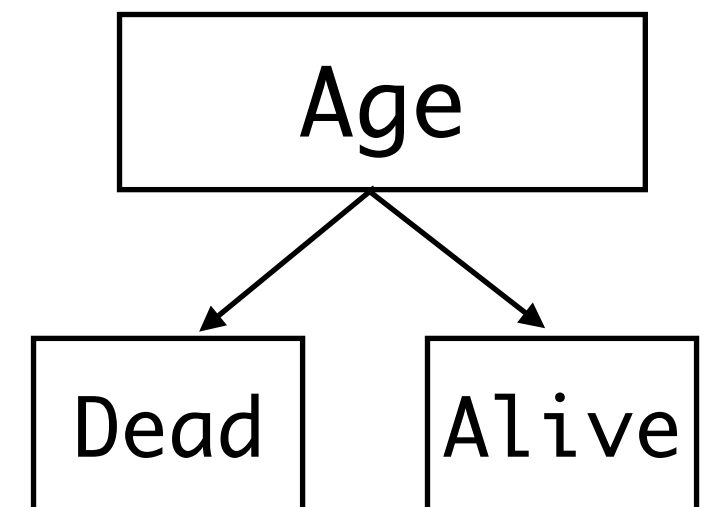
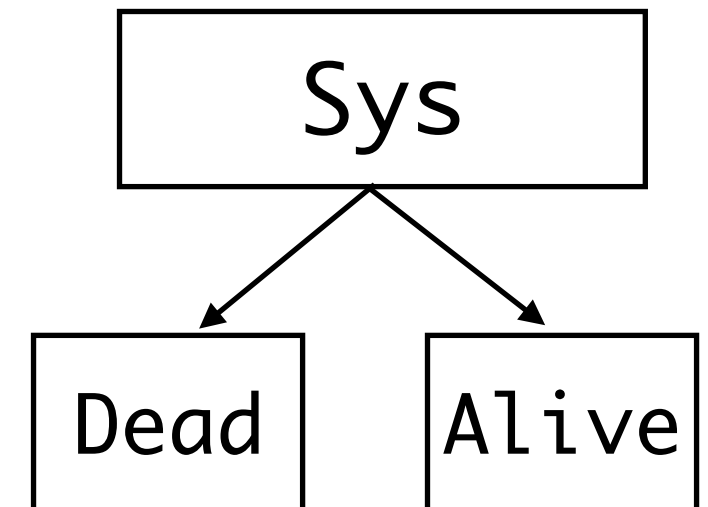
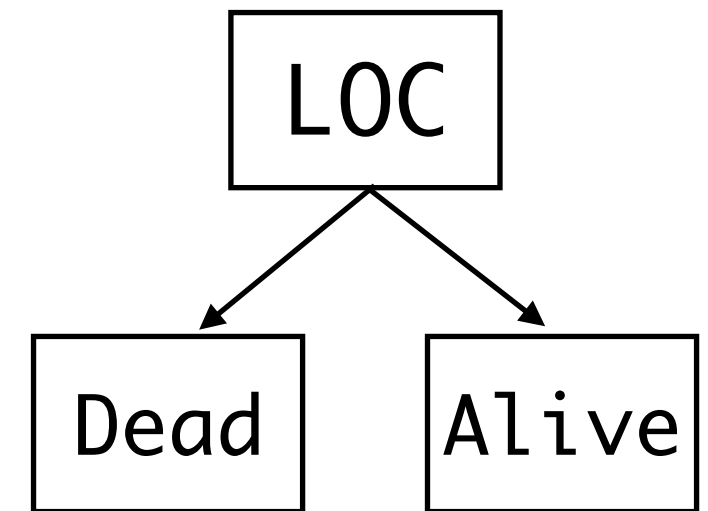
Weighted  
random  
sampling with  
replacement  
→

STA	LOC	SYS	AGE
Dead	Yes	44	77
Dead	Yes	44	77
Alive	No	140	63
Dead	No	36	27
Dead	No	48	59
Dead	Yes	44	77
Alive	No	140	63
Dead	No	36	27

## Get back to the step2: Build up the next stump

Use the new collection of data points as a training data

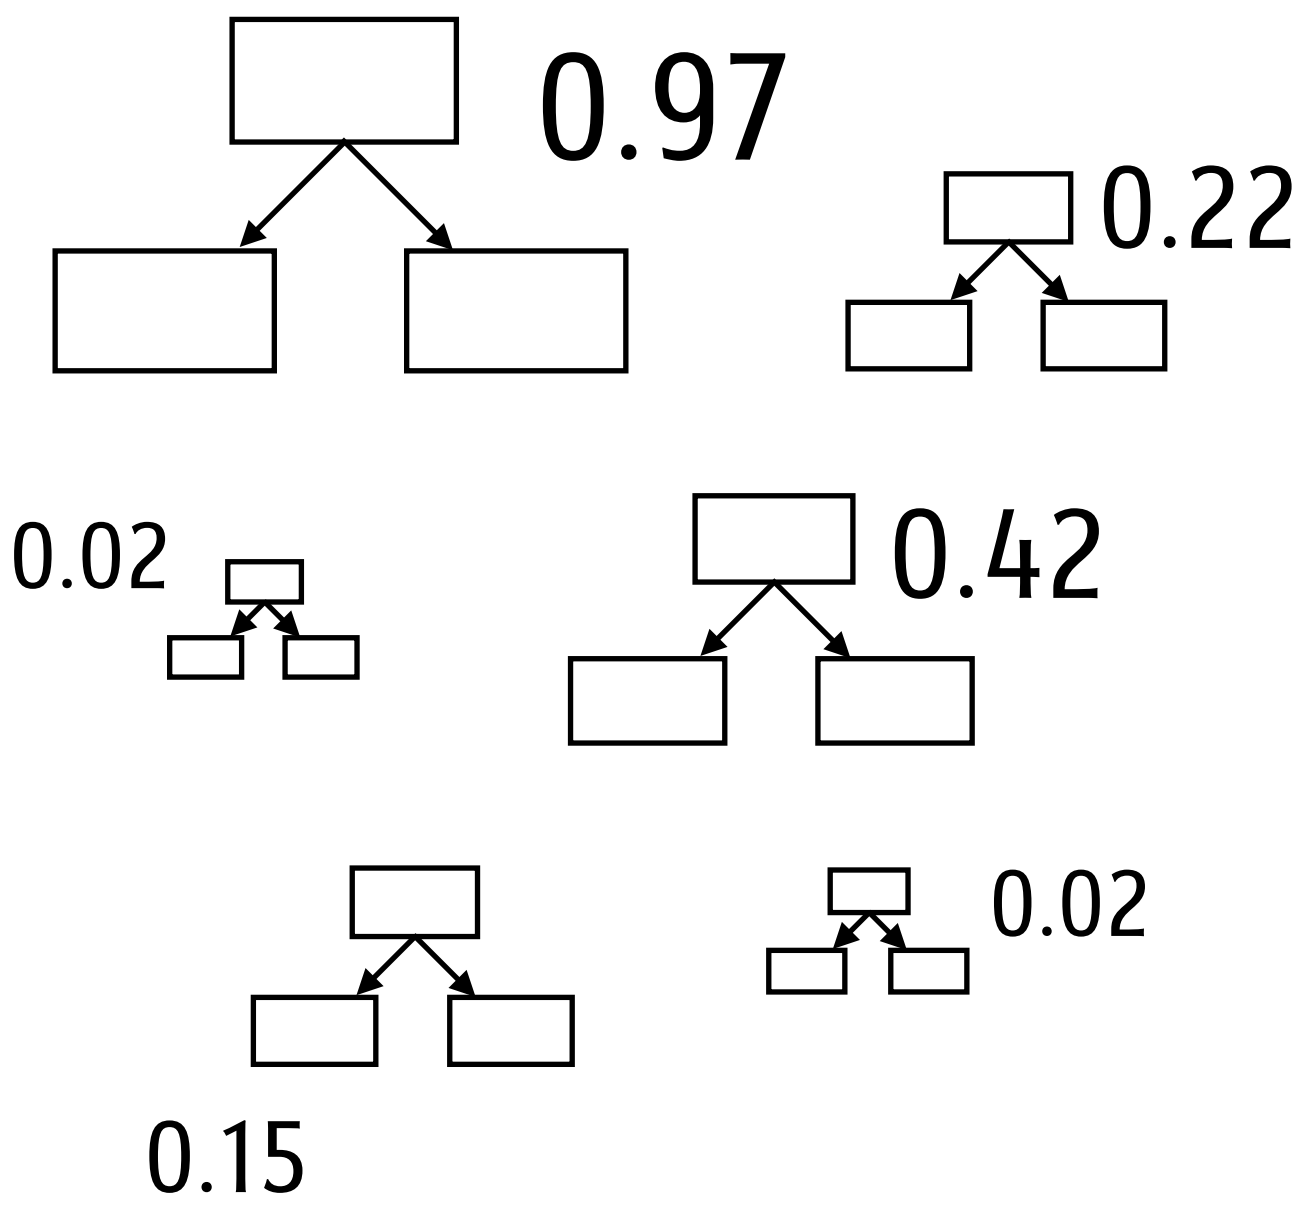
STA	LOC	SYS	AGE	Weight
Dead	Yes	44	77	1/8
Dead	Yes	44	77	1/8
Alive	No	140	63	1/8
Dead	No	36	27	1/8
Dead	No	48	59	1/8
Dead	Yes	44	77	1/8
Alive	No	140	63	1/8
Dead	No	36	27	1/8





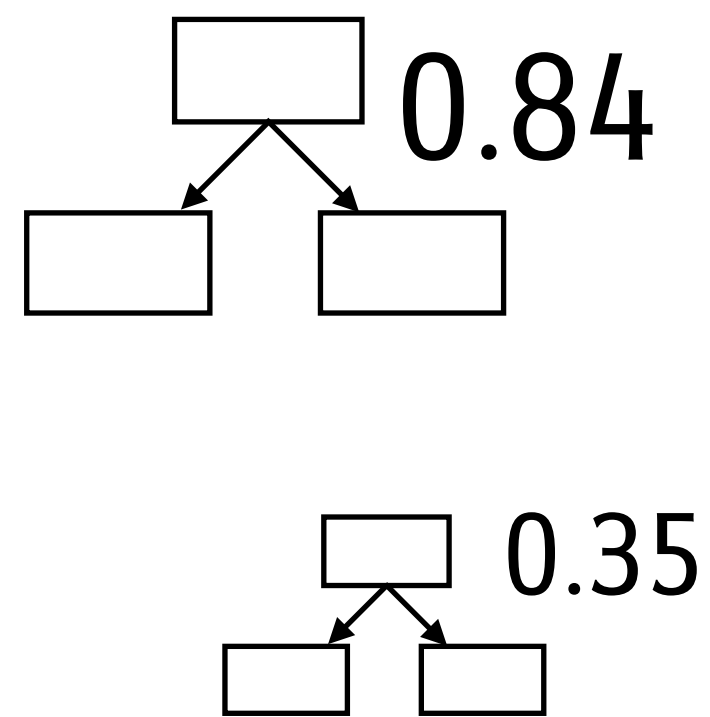
repeat ...

STA = Alive (-1)



Total amount of say = -1.8

STA = Dead (1)



Total amount of say = 1.19

# AdaBoosting using R

fit.ada

AdaBoost.M1

120 samples

4 predictor

3 classes: 'setosa', 'versicolor', 'virginica'

No pre-processing

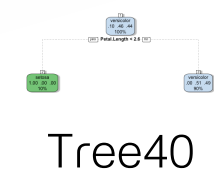
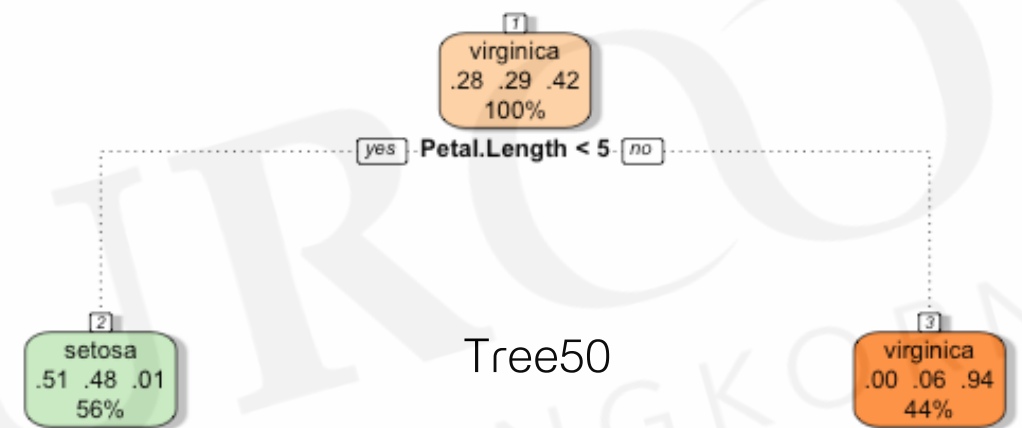
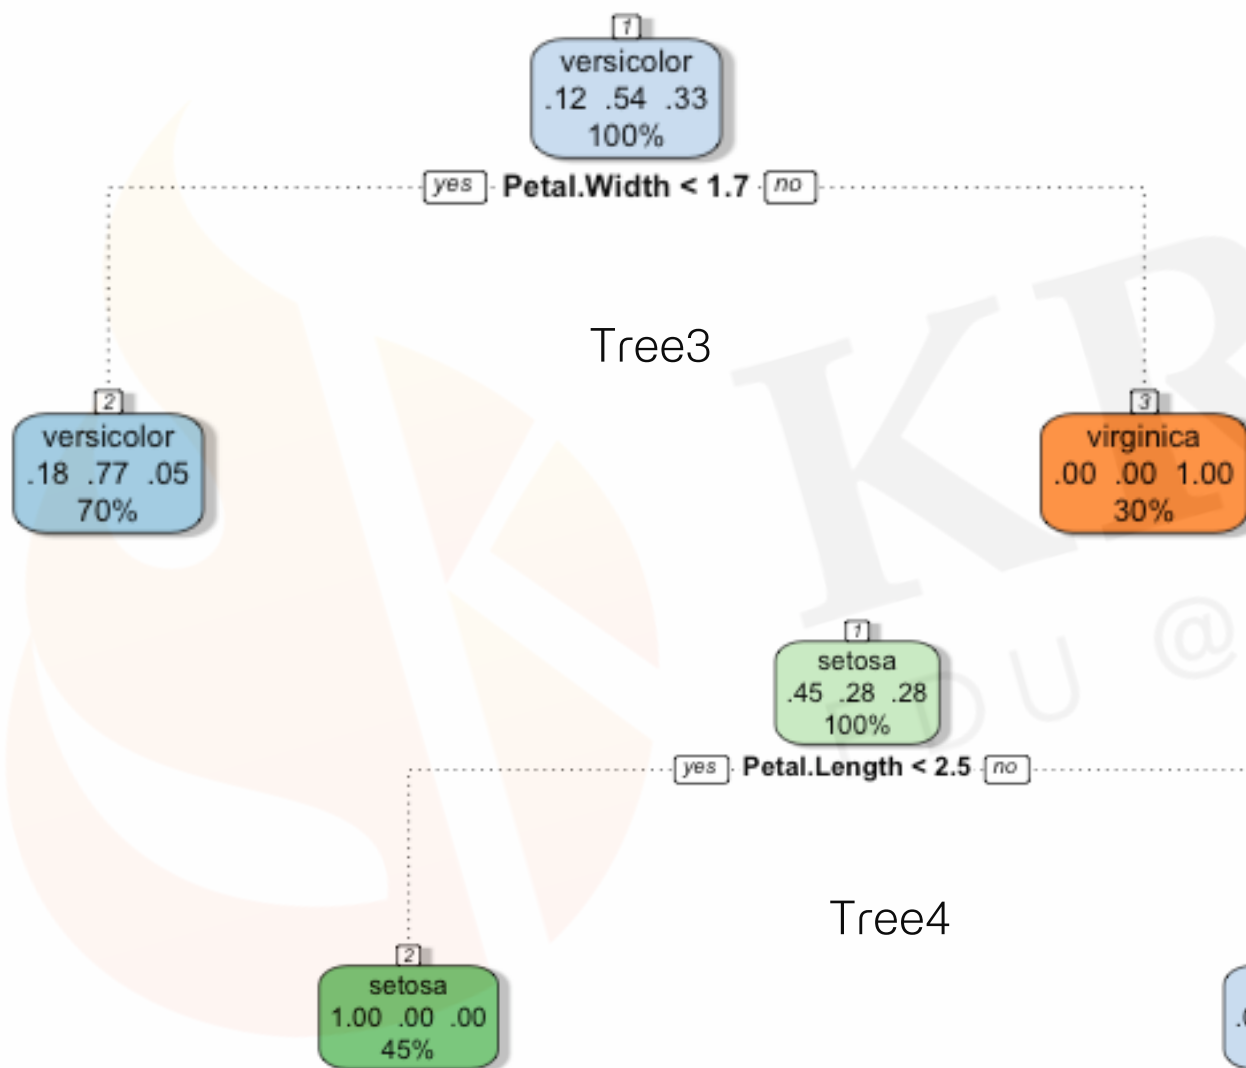
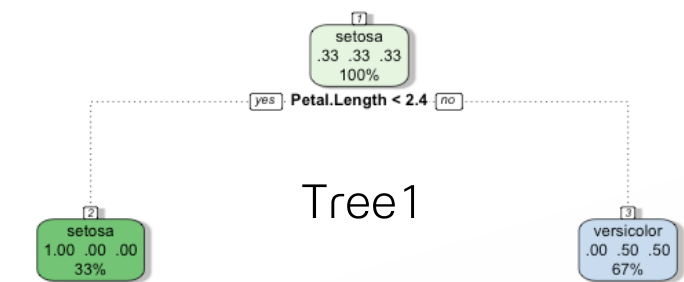
Resampling: Bootstrapped (5 reps)

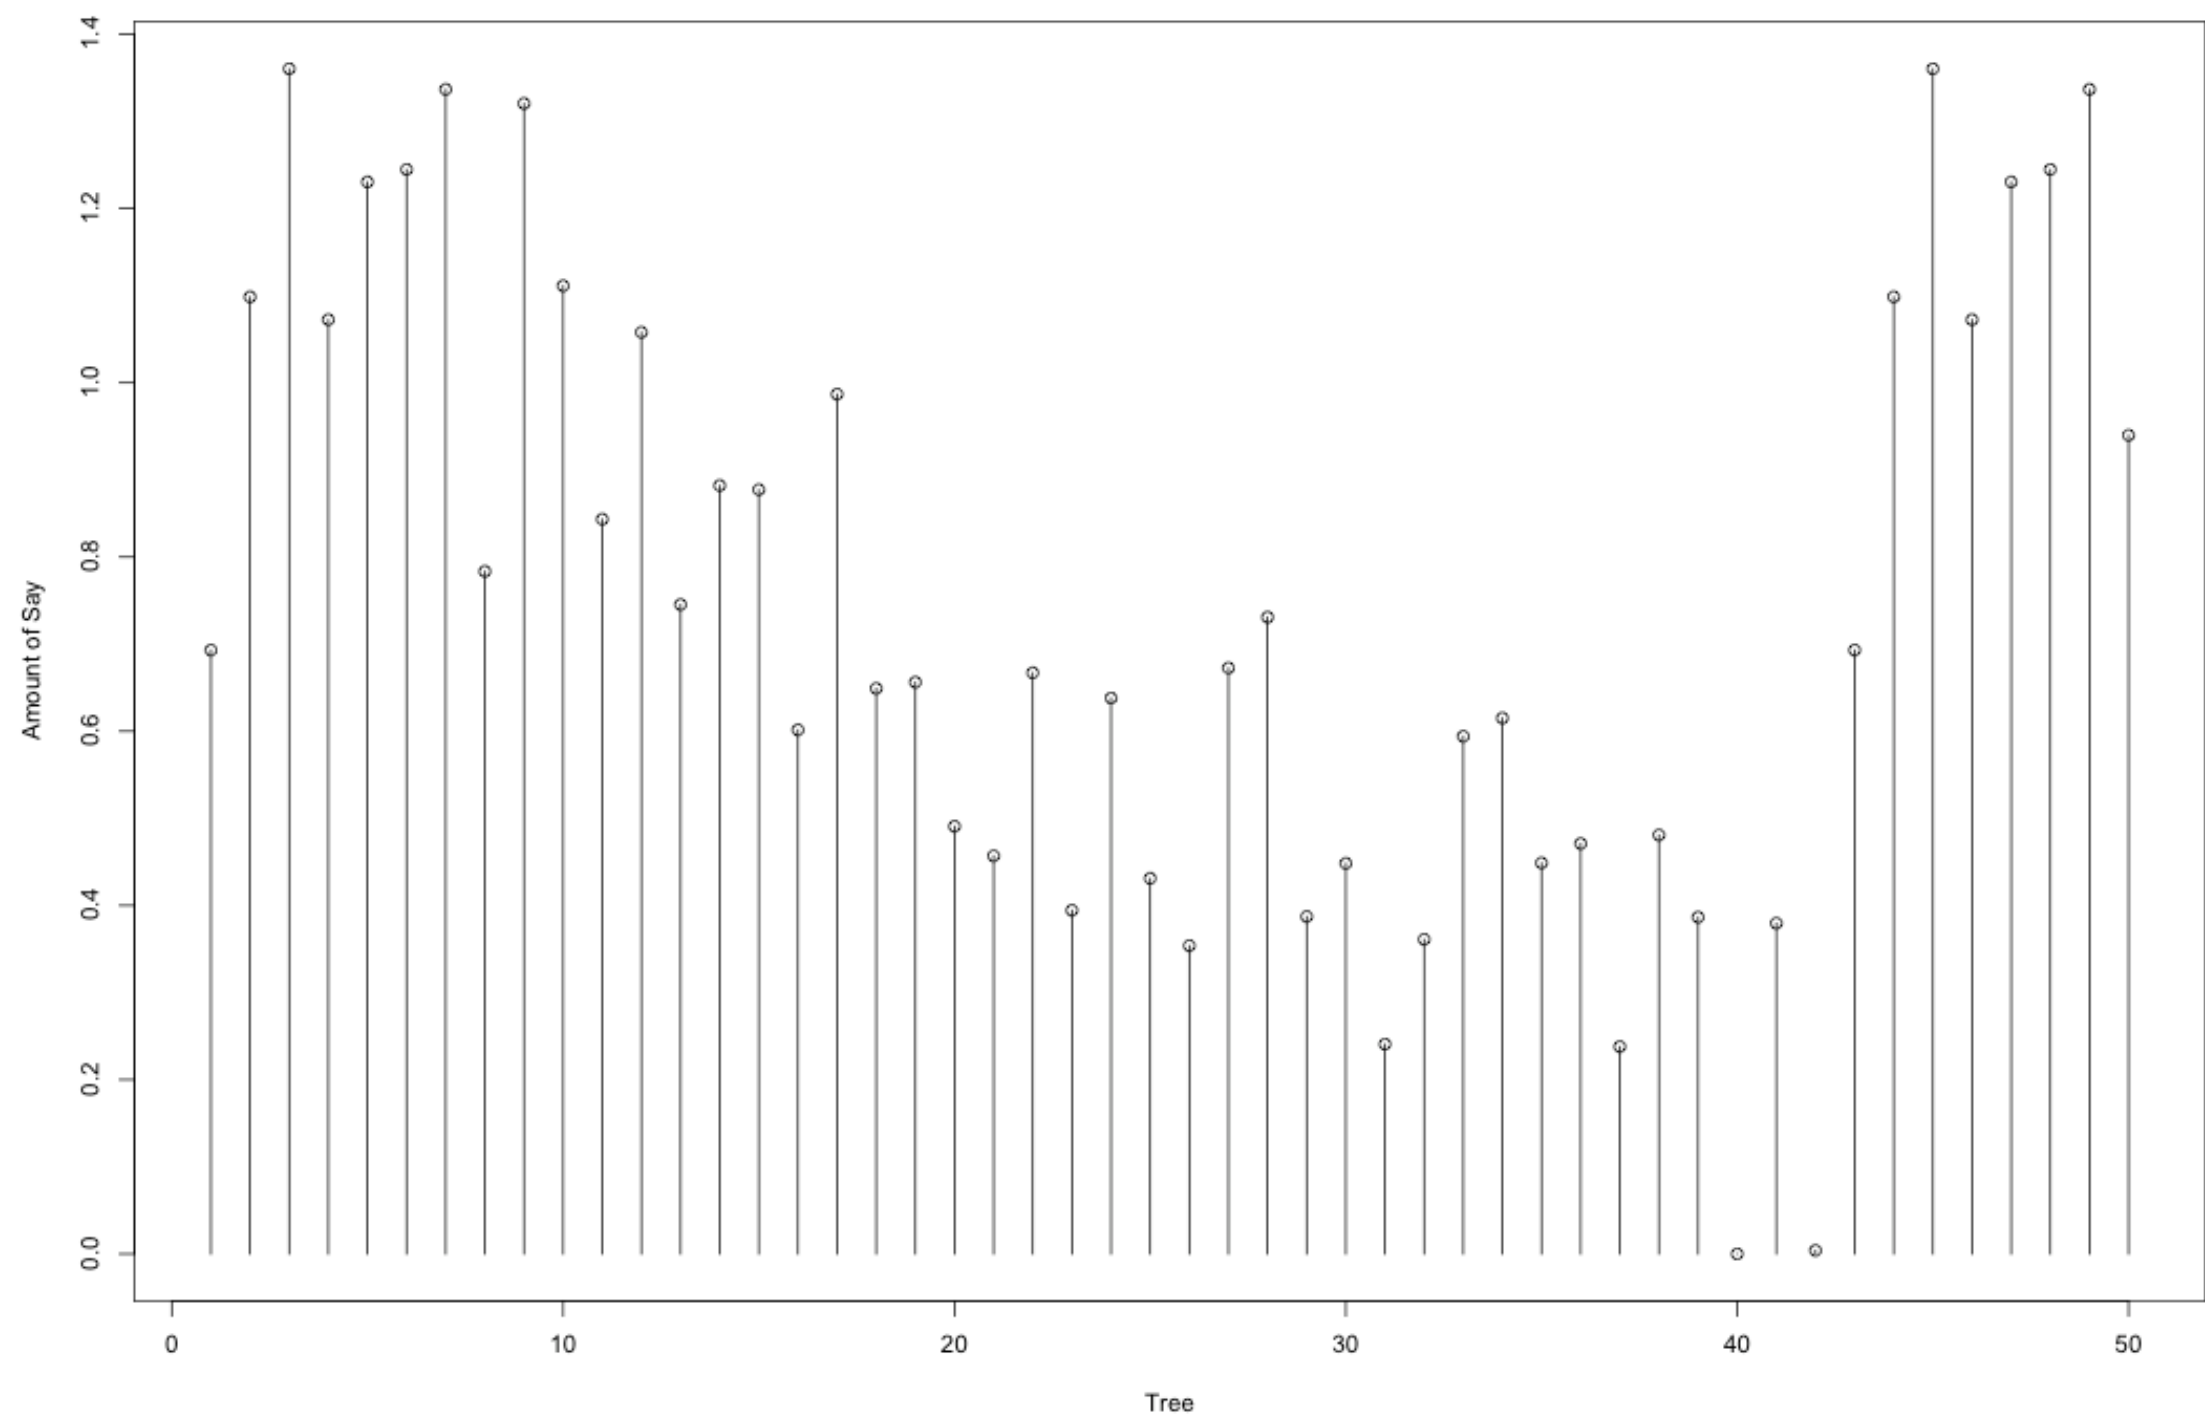
Summary of sample sizes: 24, 24, 24, 24, 24

Resampling results across tuning parameters:

coflearn	maxdepth	mfinal	logLoss	AUC	prAUC	Accuracy	Kappa	Mean_F1
Breiman	1	50	0.4750813	0.9697754	0.3008532	0.9375000	0.906250	0.9373596
Breiman	1	100	0.4767168	0.9803060	0.4699178	0.9375000	0.906250	0.9373596
Breiman	1	150	0.4756446	0.9783854	0.4825737	0.9354167	0.903125	0.9352625
Breiman	2	50	0.3659022	0.9818685	0.7150678	0.9354167	0.903125	0.9352625
Breiman	2	100	0.3745246	0.9817708	0.7789088	0.9354167	0.903125	0.9352625
Breiman	2	150	0.3806631	0.9822754	0.7898933	0.9354167	0.903125	0.9352625

# AdaBoosting using R





# Gradient Boosting (BGM)

---

- Regression
- Classification

# Gradient Boosting for Regression

Let  $(X, y) = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$  be a training dataset, and  $L(y_i, F(\mathbf{x}_i))$  be differentiable Loss function.

**Step 1:** fit the null model (model with only constant term)  $\implies F_0(\mathbf{X}) = \operatorname{argmin}_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$

**Step2:** for  $m$  in  $1 : M$

1. Compute  $e_{im} = - \left[ \frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(x)=F_{m-1}(x)}$  for  $i = 1, 2, \dots, n$   
Pseudo residual
2. Fit a regression tree to the  $e_{im}$  values and create terminal region  $R_{jm}$  for  $j = 1, \dots, J_m$
3. For  $j = 1, 2, \dots, J_m$  compute  $\gamma_{jm} = \operatorname{argmin}_x \sum_{x_i \in R_{ij}} L(y_i, F_{m-1}(x_i) + \gamma)$
4. Update  $F_m(\mathbf{x}) = F_{m-1}(\mathbf{x}) + \underset{\substack{\text{Learning rate} \\ [0,1]}}{\nu} \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$

**Output:**  $F_M(\mathbf{X})$

Let  $(X, \mathbf{y}) = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$  be a training dataset, and  $L(y_i, F(\mathbf{x}_i))$  be differentiable Loss function.

## Training dataset

SES	Level	Gender	Discipline
-1.0	Primary	Male	27
0.5	Secondary	Female	59
1.0	UnderGrad	Female	77
1.5	Primary	Male	54
1.5	UnderGrad	Male	87
1.0	Secondary	Female	69

## Loss Function

$$L(y_i, F(\mathbf{x}_i) = \hat{y}_i) = \frac{1}{2}(y_i - \hat{y}_i)^2$$

$$\frac{\partial L}{\partial \hat{y}_i} = -(y_i - \hat{y}_i)$$

where  $F(\mathbf{x}_i)$  is a function that give us the predicted values.

$$\sum_{i=1}^n \frac{1}{2}(y_i - \hat{y}_i)^2$$

$$\text{SSE} = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$



**Step 1:** fit the null model (model with only constant term)  $\implies F_0(\mathbf{X}) = \underset{\gamma}{\operatorname{argmin}} \sum_{i=1}^n L(y_i, \gamma)$

$L(y_i, \gamma) = \frac{1}{2}(y_i - \gamma)^2$  is loss function; where  $y_i$  and  $\gamma$  refer to observed and predicted value respectively.

$\sum_{i=1}^n L(y_i, \gamma)$  is the overall Loss function

$\underset{\gamma}{\operatorname{argmin}}$  means we need to find a predicted value that minimized the above sum.

$$\frac{\partial \sum_{i=1}^n L(y_i, \gamma)}{\partial \gamma} = - \sum_{i=1}^n (y_i - \gamma) = 0$$

$$\implies \gamma = \frac{\sum_{i=1}^n y_i}{n}$$

Hence, given the loss function  $L(y_i, \gamma) = \frac{1}{2}(y_i - \gamma)^2$  the value of  $\gamma$  that minimizes the above sum is the average of the observed outcome.  $\implies F_0(\mathbf{X}) = \gamma = \frac{\sum_{i=1}^n y_i}{n}$

**Step 1:** fit the null model (model with only constant term)  $\Rightarrow F_0(\mathbf{X}) = \operatorname{argmin}_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$

SES	Level	Gender	Discipline
-1.0	Primary	Male	27
0.5	Secondary	Female	59
1.0	UnderGrad	Female	77
1.5	Primary	Male	54
1.5	UnderGrad	Male	87
1.0	Secondary	Female	69

$$\text{Hence, } F_0(\mathbf{X}) = \frac{27 + 59 + \dots + 69}{6} = 62.17$$

That means the null model (or initial predicted value) is a tree with just one leaf.

62.17

- Gradient boost starts by making a single leaf, instead of a tree or stump.
- This leaf represent an prior or initial guess for the weight of all subjects

**Step2:** for  $m$  in  $1 : M$  ← We will produce  $M$  trees. (i.e.  $M=100$ )

1. Compute  $e_{im} = - \left[ \frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(x)=F_{m-1}(x)}$  for  $i = 1, 2, \dots, n$
2. Fit a regression tree to the  $e_{im}$  values and create terminal region  $R_{jm}$  for  $j = 1, \dots, J_m$
3. For  $j = 1, 2, \dots, J_m$  compute  $\gamma_{jm} = \operatorname{argmin}_x \sum_{x_i \in R_{ij}} L(y_i, F_{m-1}(x_i) + \gamma)$
4. Update  $F_m(\mathbf{x}) = F_{m-1}(\mathbf{x}) + \nu \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$

## Step2: for $m$ in $1 : M$

Let  $m = 1$

1. Compute  $e_{im} = - \left[ \frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(x)=F_{m-1}(x)}$  for  $i = 1, 2, \dots, n$

Since the loss function is  $L(y_i, F(\mathbf{x}_i) = \hat{y}_i) = \frac{1}{2}(y_i - \hat{y}_i)^2$ ,

$$\begin{aligned} e_{i(m=1)} &= - \left[ \frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(x)=F_0(x)} \\ &= - \left[ \frac{\partial L(y_i, F_0(\mathbf{x}_i))}{\partial F_0(\mathbf{x}_i)} \right] ; \text{ where } F_0(\mathbf{x}_i) = \mu = 62.17 \\ &= (y_i - 62.17) \end{aligned}$$

hence,  $e_{im}$  is a residual of  $i$ -th sample, and  $m$ -th trees

$$e_{1,1} = 27 - 62.17 = -35.17$$

...

$$e_{6,1} = 69 - 62.17 = 6.83$$

Step2: for  $m$  in  $1 : M$

SES	Level	Gender	Discipline	(Pseudo) Residual=actual-predicted ( $e_{i,1}$ )
-1.0	Primary	Male	27	-35.17
0.5	Secondary	Female	59	-3.17
1.0	UnderGrad	Female	77	14.83
1.5	Primary	Male	54	-8.17
1.5	UnderGrad	Male	87	24.83
1.0	Secondary	Female	69	6.83

Average value of  
Discipline score

62.17

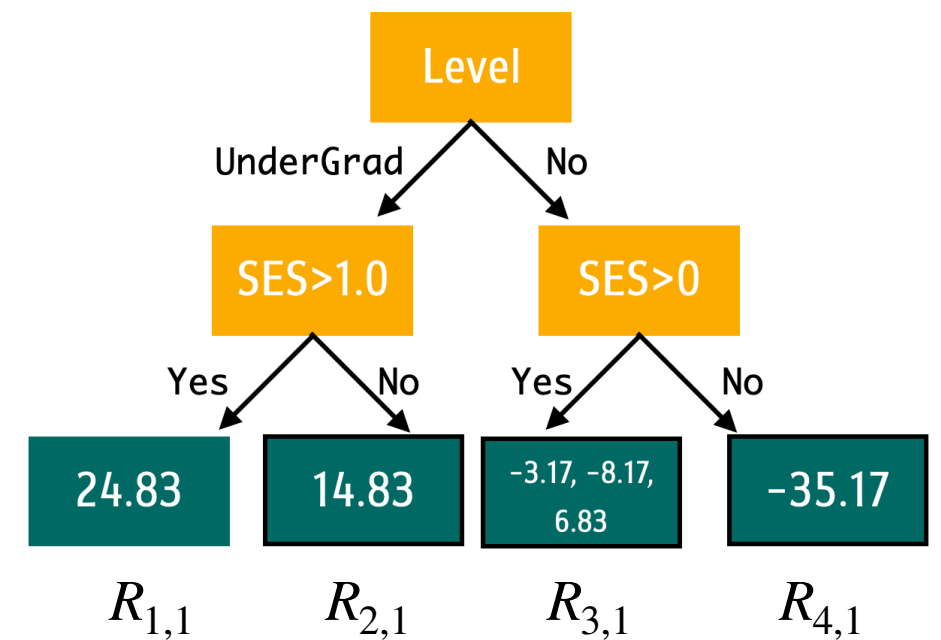
## Step2: for $m$ in $1 : M$

Let  $m = 1$

1. Compute  $e_{im} = - \left[ \frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(x)=F_{m-1}(x)}$  for  $i = 1, 2, \dots, n$
2. Fit a regression tree to the  $e_{im}$  values and create terminal region  $R_{jm}$  for  $j = 1, \dots, J_m$

- We will build a regression tree to predict the residual instead of outcome.

SES	Level	Gender	Discipline	$e_{i,1}$
-1.0	Primary	Male	27	-35.17
0.5	Secondary	Female	59	-3.17
1.0	UnderGrad	Female	77	14.83
1.5	Primary	Male	54	-8.17
1.5	UnderGrad	Male	87	24.83
1.0	Secondary	Female	69	6.83



Predict

The leaves are the terminal regions  $R_{jm}$  for  $j = 1, 2, \dots, J_m = 4$

## Step2: for $m$ in $1 : M$

Let  $m = 1$

1. Compute  $e_{im} = - \left[ \frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(x)=F_{m-1}(x)}$  for  $i = 1, 2, \dots, n$
2. Fit a regression tree to the  $e_{im}$  values and create terminal region  $R_{jm}$  for  $j = 1, \dots, J_m$
3. For  $j = 1, 2, \dots, J_m$  compute  $\gamma_{jm} = \operatorname{argmin}_x \sum_{x_i \in R_{ij}} L(y_i, F_{m-1}(x_i) + \gamma)$



For the each leaf in the new regression tree, we calculate the new estimate value ( $\gamma_{jm}$ ) of the output value for each leaf

SES	Level	Gender	Discipline	New prediction
-1.0	Primary	Male	27	
0.5	Secondary	Female	59	
1.0	UnderGrad	Female	77	
1.5	Primary	Male	54	
1.5	UnderGrad	Male	87	62.17+24.83 = 87
1.0	Secondary	Female	69	

$$\gamma_{1,1} = \operatorname{argmin}_{\gamma} \sum_{x_i \in R_{ij}} L(y_i, F_{m-1}(x_i) + \gamma)$$

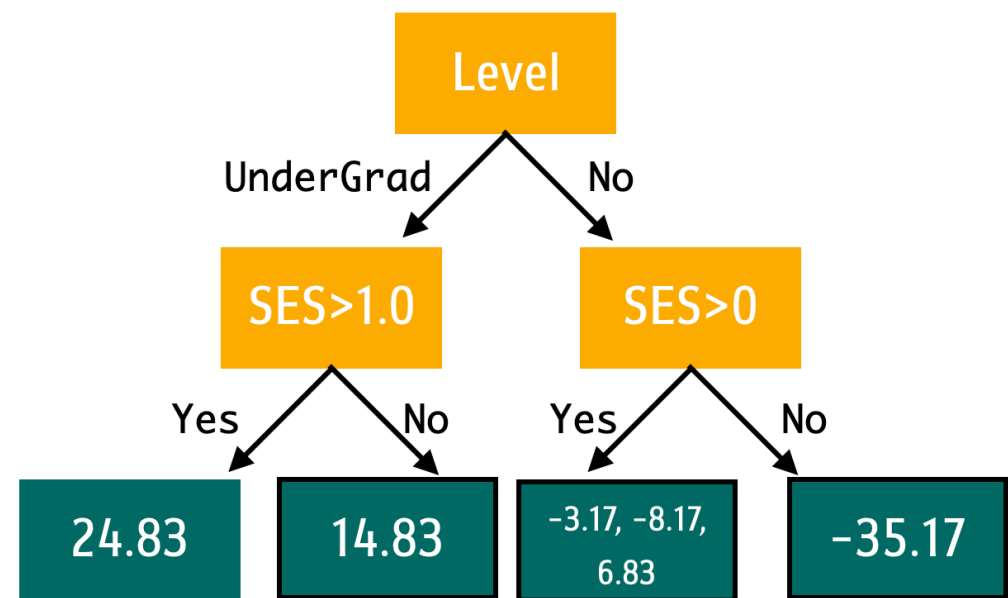
$$= \operatorname{argmin}_{\gamma} \sum_{x_i \in R_{ij}} \frac{1}{2} (y_i - (F_{m-1}(\mathbf{x}_i) + \gamma))^2$$

$$= \operatorname{argmin}_{\gamma} \frac{1}{2} (87 - (62.17 + \gamma))^2$$

$$= \operatorname{argmin}_{\gamma} \frac{1}{2} (24.83 - \gamma)^2 \implies \frac{\partial \frac{1}{2} (24.83 - \gamma)^2}{\partial \gamma} = 0$$

$$\implies \gamma = 24.83$$





$$R_{1,1}$$

$$\gamma_{1,1} = 24.83$$

$$R_{2,1}$$

$$\gamma_{2,1} = 14.83$$

$$R_{3,1}$$

$$\gamma_{3,1} = \operatorname{argmin}_{\gamma} \left\{ \frac{1}{2}(59 - (62.17 + \gamma))^2 + \frac{1}{2}(54 - (62.17 + \gamma))^2 + \frac{1}{2}(69 - (62.17 + \gamma))^2 \right\}$$

$$= \frac{-3.17 - 8.17 + 6.83}{3} = -1.50$$

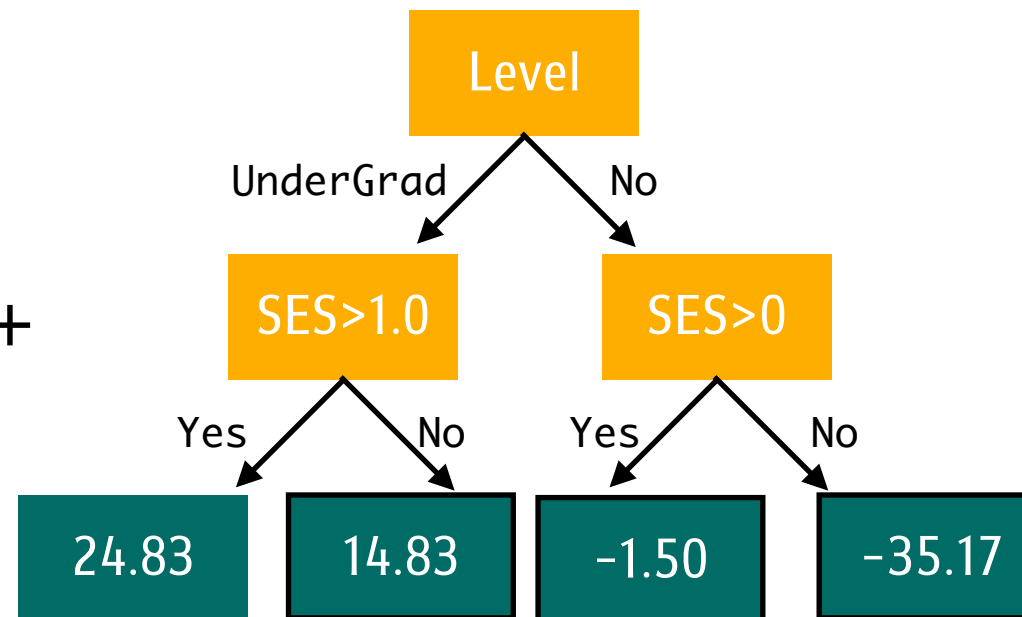
$$R_{4,1}$$

$$\gamma_{4,1} = 14.83 - 35.17$$

Average value of  
Discipline score

62.17

+



SES	Level	Gender	Discipline	New prediction
-1.0	Primary	Male	27	$62.17 - 35.17 = 27$
0.5	Secondary	Female	59	$62.17 - 1.50 = 60.67$
1.0	UnderGrad	Female	77	$62.17 + 14.83 = 77$
1.5	Primary	Male	54	$62.17 - 1.50 = 60.67$
1.5	UnderGrad	Male	87	$62.17 + 24.83 = 87$
1.0	Secondary	Female	69	$62.17 - 1.50 = 60.67$

} Overfit

## Step2: for $m$ in $1 : M$

Let  $m = 1$

1. Compute  $e_{im} = - \left[ \frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(x)=F_{m-1}(x)}$  for  $i = 1, 2, \dots, n$
2. Fit a regression tree to the  $e_{im}$  values and create terminal region  $R_{jm}$  for  $j = 1, \dots, J_m$
3. For  $j = 1, 2, \dots, J_m$  compute  $\gamma_{jm} = \operatorname{argmin}_x \sum_{x_i \in R_{ij}} L(y_i, F_{m-1}(x_i) + \gamma)$
4. Update  $F_m(\mathbf{x}) = F_{m-1}(\mathbf{x}) + \nu \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$



This is the new prediction

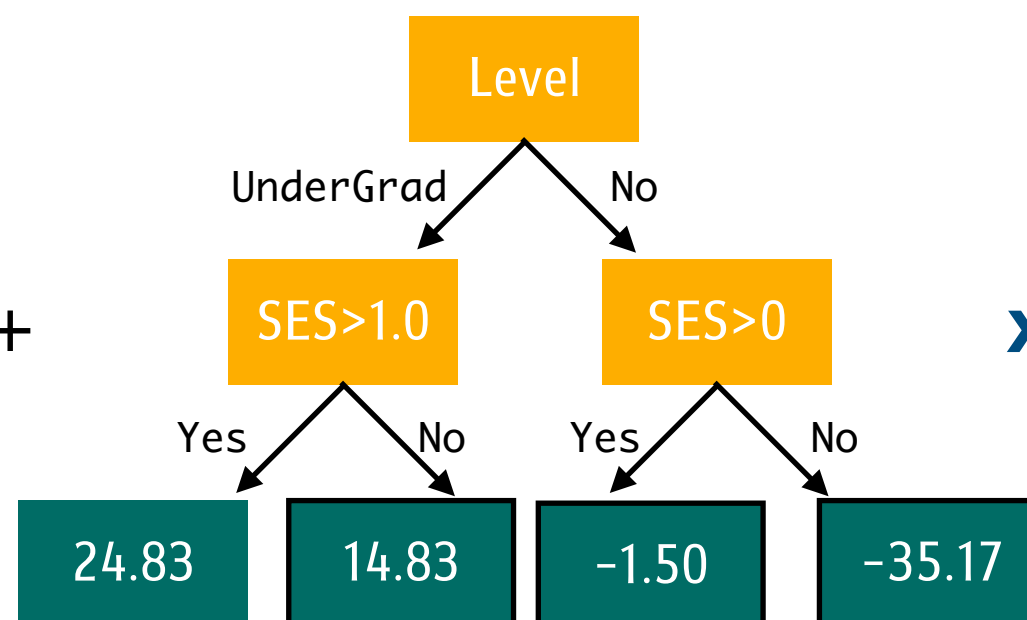
$$F_1(\mathbf{x}) = 62.17 + \nu \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$$

- $\nu \in (0,1)$  is the learning rate
- Learning rate is an effect of the regression tree on the final prediction

Average value of  
Discipline score

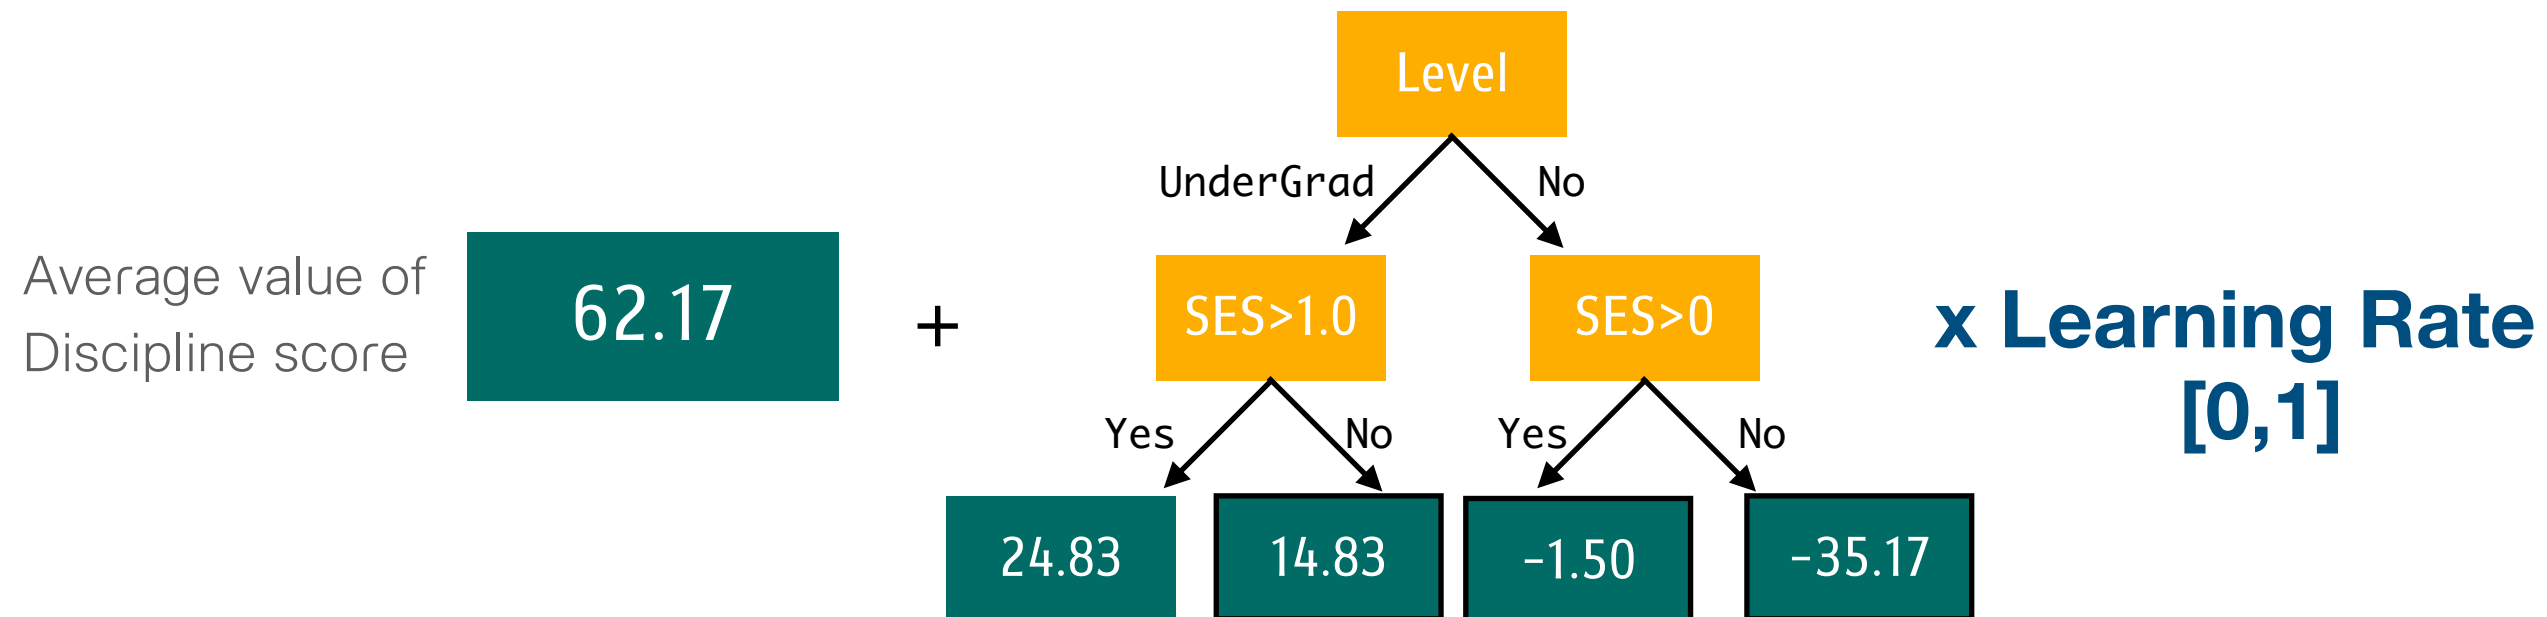
62.17

+



**x Learning Rate**  
**[0,1]**

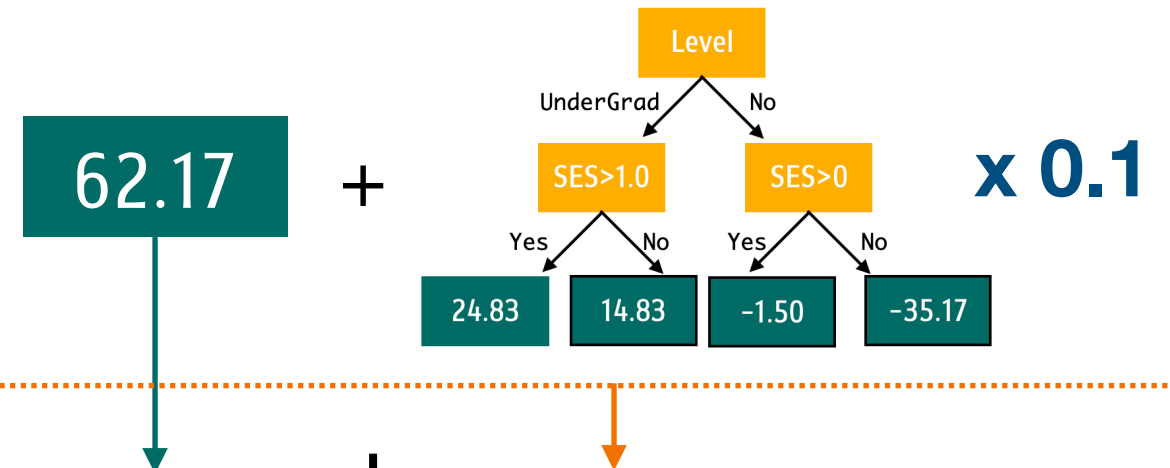
SES	Level	Gender	Discipline	New prediction
-1.0	Primary	Male	27	$62.17 - 35.17(0.1) = 58.65$
0.5	Secondary	Female	59	$62.17 - 1.50(0.1) = 62.02$
1.0	UnderGrad	Female	77	$62.17 + 14.83(0.1) = 63.65$
1.5	Primary	Male	54	$62.17 - 1.50(0.1) = 62.02$
1.5	UnderGrad	Male	87	$62.17 + 24.83(0.1) = 64.65$
1.0	Secondary	Female	69	$62.17 - 1.50(0.1) = 62.02$



- We will see that with the learning rate = 0.1, the prediction values are not as good as the prediction values with learning rate = 1.0. However there are little better than initial prediction = 62.17.
- Hence, by scaling the tree with a learning rate results in a small step in the right direction.
- Empirical study shows that taking lots of small steps in the right direction results in better predictions with a testing dataset (lower variance).

**Repeat step2:** Build another tree based on the error of previous tree.

Average value of  
Discipline score

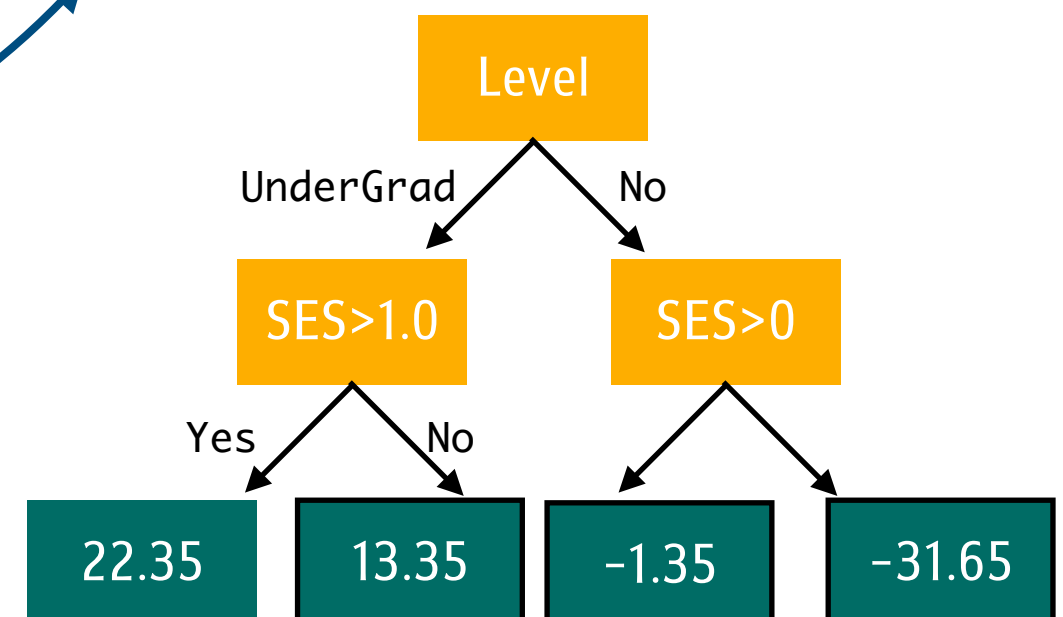


SES	Level	Gender	Discipline	Residual1	Residual2
-1.0	Primary	Male	27	-35.17	-31.65
0.5	Secondary	Female	59	-3.17	-3.02
1.0	UnderGrad	Female	77	14.83	13.35
1.5	Primary	Male	54	-8.17	-8.02
1.5	UnderGrad	Male	87	24.83	22.35
1.0	Secondary	Female	69	6.83	6.98

**Repeat step2:** Build another tree based on the error of previous tree.

SES	Level	Gender	Discipline	Residual2
-1.0	Primary	Male	27	-31.65
0.5	Secondary	Female	59	-3.02
1.0	UnderGrad	Female	77	13.35
1.5	Primary	Male	54	-8.02
1.5	UnderGrad	Male	87	22.35
1.0	Secondary	Female	69	6.98

Predict

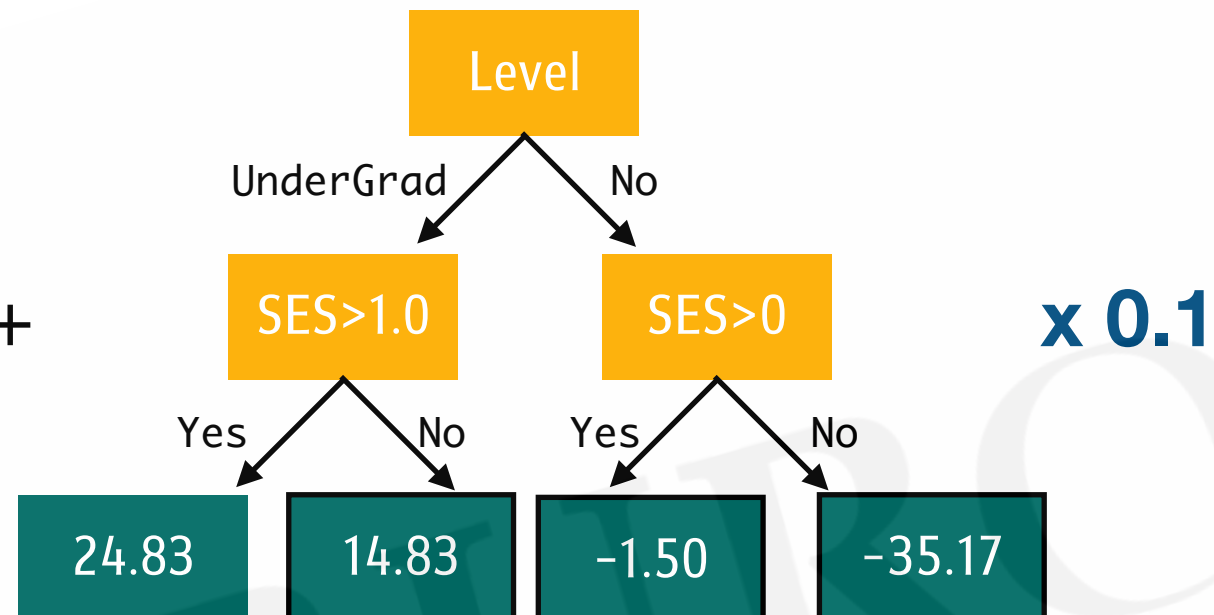


Combine the original leaf with the new tree to make a new prediction of an individual discipline from the training data

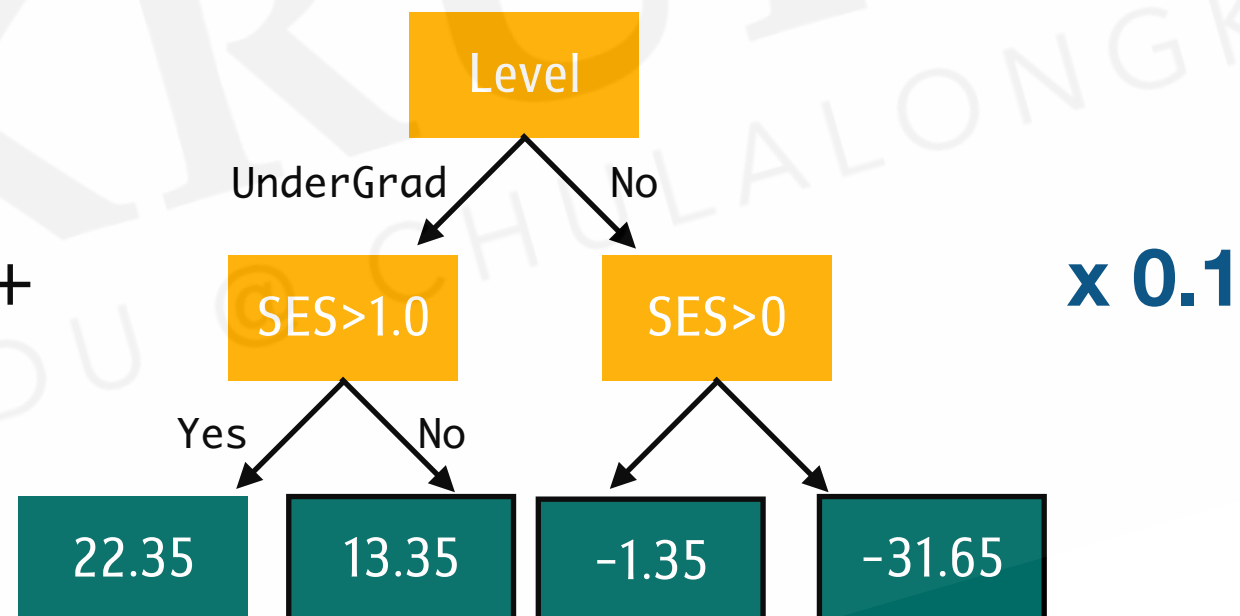
Average value of  
Discipline score

62.17

+



+



= Prediction value



Combine the original leaf with the new tree to make a new prediction of an individual discipline from the training data

SES	Level	Gender	Discipline	Prediction1	Prediction2
-1.0	Primary	Male	27	58.65	$62.17 + 0.1 * (-35.17) + (0.1) * (-31.65) = 55.49$
0.5	Secondary	Female	59	62.02	
1.0	UnderGrad	Female	77	63.65	
1.5	Primary	Male	54	62.02	
1.5	UnderGrad	Male	87	64.65	
1.0	Secondary	Female	69	62.02	

repeat ...

# GBM Hyperparameters

- `n.trees`: number of trees
- `bag.fraction`: proportion of observations to be sampled in each tree
- `n.minobsinnode`: minimum number of observations in the trees terminal nodes
- `interaction.depth`: maximum nodes per tree
- `shrinkage`: learning rate

# Gradient Boosting for Regression

Let  $(X, y) = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$  be a training dataset, and  $L(y_i, F(\mathbf{x}_i))$  be differentiable Loss function.

**Step 1:** fit the null model (model with only constant term)  $\implies F_0(\mathbf{X}) = \operatorname{argmin}_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$

**Step2:** for  $m$  in  $1 : M$

1. Compute  $e_{im} = - \left[ \frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(x)=F_{m-1}(x)}$  for  $i = 1, 2, \dots, n$   
Pseudo residual
2. Fit a regression tree to the  $e_{im}$  values and create terminal region  $R_{jm}$  for  $j = 1, \dots, J_m$
3. For  $j = 1, 2, \dots, J_m$  compute  $\gamma_{jm} = \operatorname{argmin}_x \sum_{x_i \in R_{ij}} L(y_i, F_{m-1}(x_i) + \gamma)$
4. Update  $F_m(\mathbf{x}) = F_{m-1}(\mathbf{x}) + \underset{\substack{\text{Learning rate} \\ [0,1]}}{\nu} \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$

**Output:**  $F_M(\mathbf{X})$

# Gradient Boosting for Classification

Let  $(X, y) = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$  be a training dataset, and  $L(y_i, F(\mathbf{x}_i))$  be differentiable Loss function.

**Step 1:** fit the null model (model with only constant term)  $\implies F_0(\mathbf{X}) = \operatorname{argmin}_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$

**Step2:** for  $m$  in  $1 : M$

1. Compute  $e_{im} = - \left[ \frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(x)=F_{m-1}(x)}$  for  $i = 1, 2, \dots, n$   
Pseudo residual
2. Fit a regression tree to the  $e_{im}$  values and create terminal region  $R_{jm}$  for  $j = 1, \dots, J_m$
3. For  $j = 1, 2, \dots, J_m$  compute  $\gamma_{jm} = \operatorname{argmin}_x \sum_{x_i \in R_{ij}} L(y_i, F_{m-1}(x_i) + \gamma)$
4. Update  $F_m(\mathbf{x}) = F_{m-1}(\mathbf{x}) + \underset{\substack{\text{Learning rate} \\ [0,1]}}{\nu} \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$

**Output:**  $F_M(\mathbf{X})$

Let  $(X, y) = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$  be a training dataset, and  $L(y_i, F(\mathbf{x}_i))$  be differentiable Loss function.

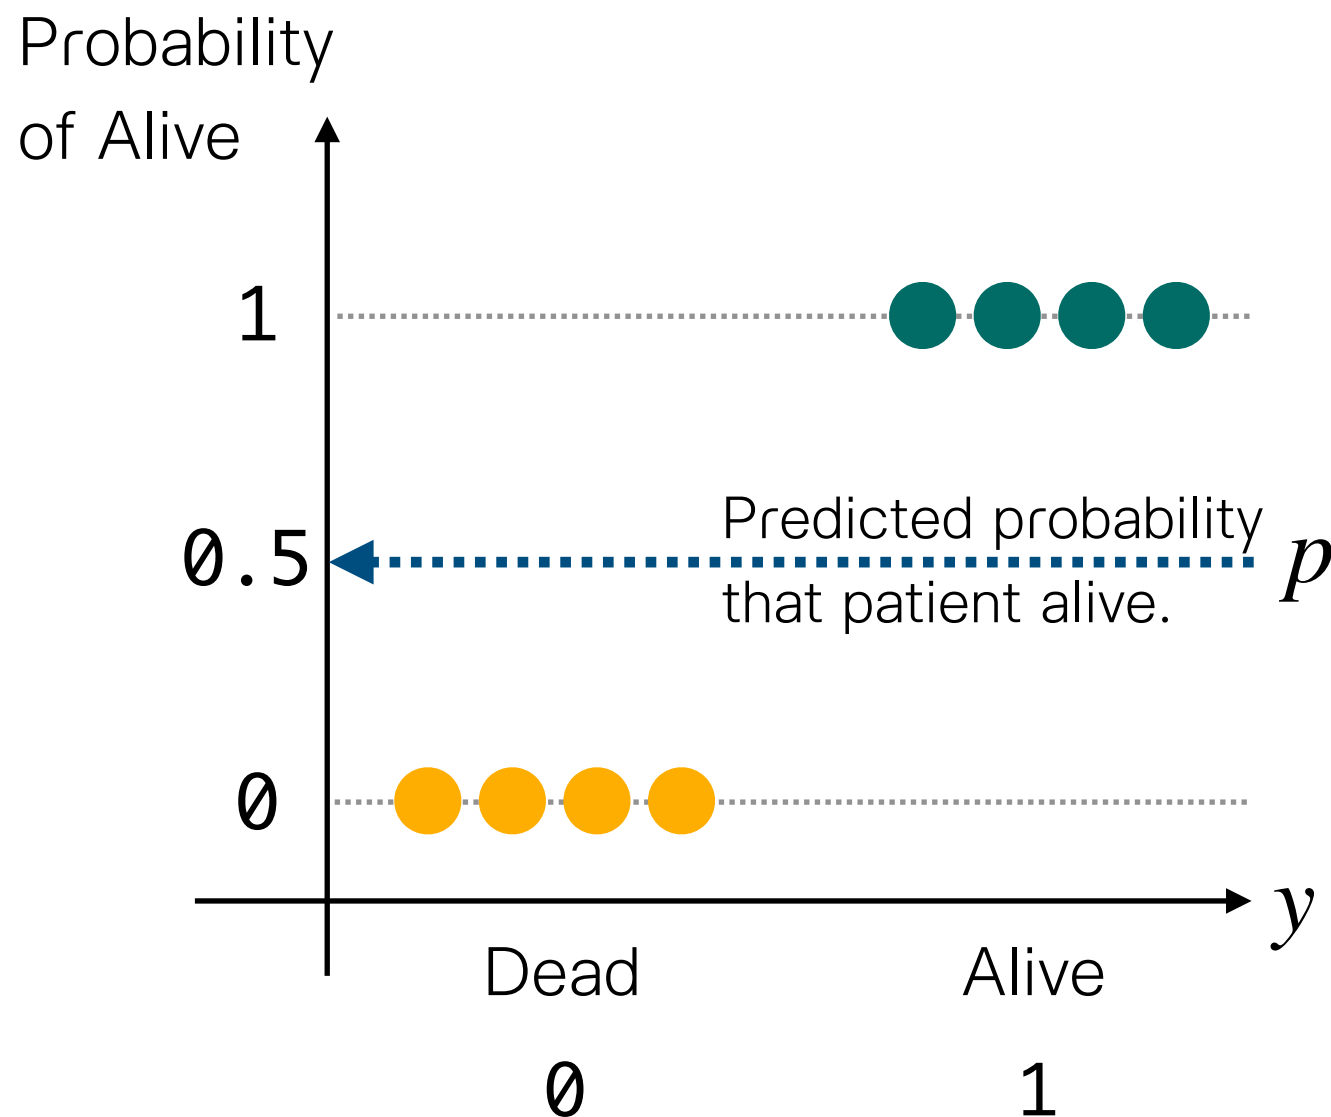
## Training Dataset

LOC	SYS	AGE	STA
No	36	27	Dead
No	48	59	Dead
Yes	44	77	Dead
No	62	54	Dead
Yes	112	87	Alive
Yes	108	69	Alive
No	140	63	Alive
Yes	138	30	Alive

- STA - life status
- LOC - level of consciousness
- Sys - systolic blood pressure
- Age - patient's age

Let  $(X, y) = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$  be a training dataset, and  $L(y_i, F(\mathbf{x}_i))$  be differentiable Loss function.

## How to define **the Loss function**?



Let  $y_i$  be patient status (STA), and by using probability theory, we have

$$y_i \sim \text{Ber}(p)$$

where  $p$  is the predicted probability, and  $y_i$  is the observed value of STA (0 = dead or 1 = alive)

So, the observation model is

$$p(y_i) = p^{y_i}(1 - p)^{1-y_i} \quad ; \quad y_i = 0, 1$$

Hence **the likelihood function** is

$$p(\mathbf{y} | p) = \prod_{i=1}^n p^{y_i}(1 - p)^{1-y_i}$$

Consistency between  $y$  and  $p$

Let  $(X, y) = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$  be a training dataset, and  $L(y_i, F(\mathbf{x}_i))$  be differentiable Loss function.

## How to define **the Loss function?**

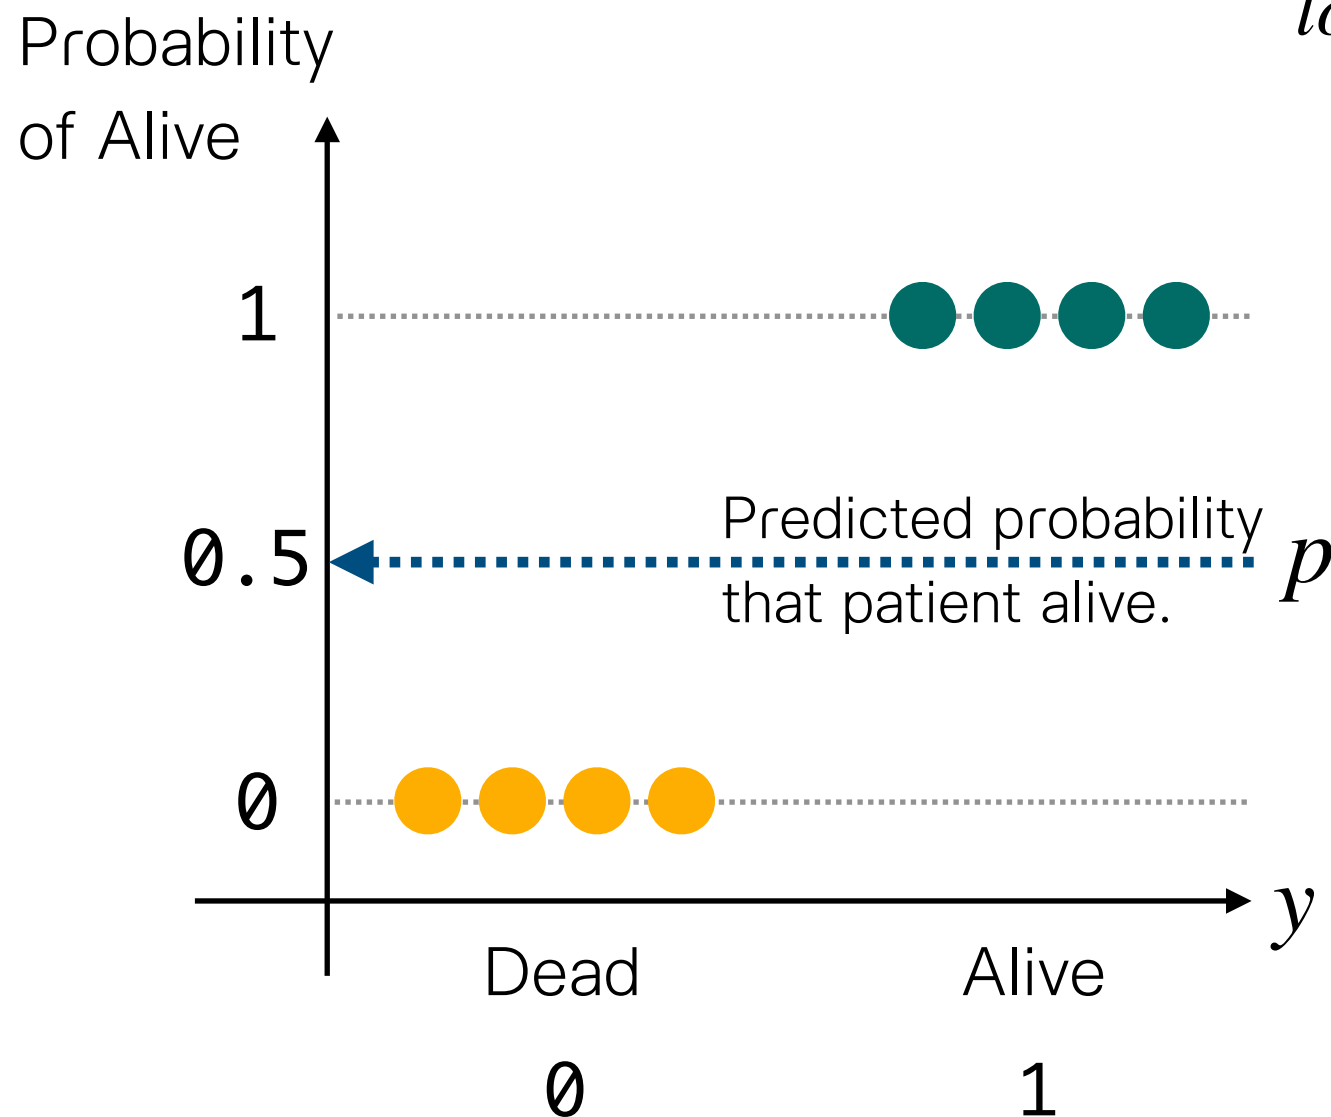
$$p(\mathbf{y} | p) = \prod_{i=1}^n p^{y_i} (1 - p)^{1-y_i}$$

$$\log[p(\mathbf{y} | p)] = \log[\prod_{i=1}^n p^{y_i} (1 - p)^{1-y_i}]$$

$$\log(p^{y_i} (1 - p)^{1-y_i})$$

$$\log(p^{y_i}) + \log((1 - p)^{1-y_i})$$

$$y_i \log(p) + (1 - y_i) \log((1 - p))$$



- The Log-likelihood of the data given the prediction is

$$\sum_{i=1}^n [y_i \log(p) + (1 - y_i) \log(1 - p)]$$

per observation  
log-likelihood



Let  $(X, y) = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$  be a training dataset, and  $L(y_i, F(\mathbf{x}_i))$  be differentiable Loss function.

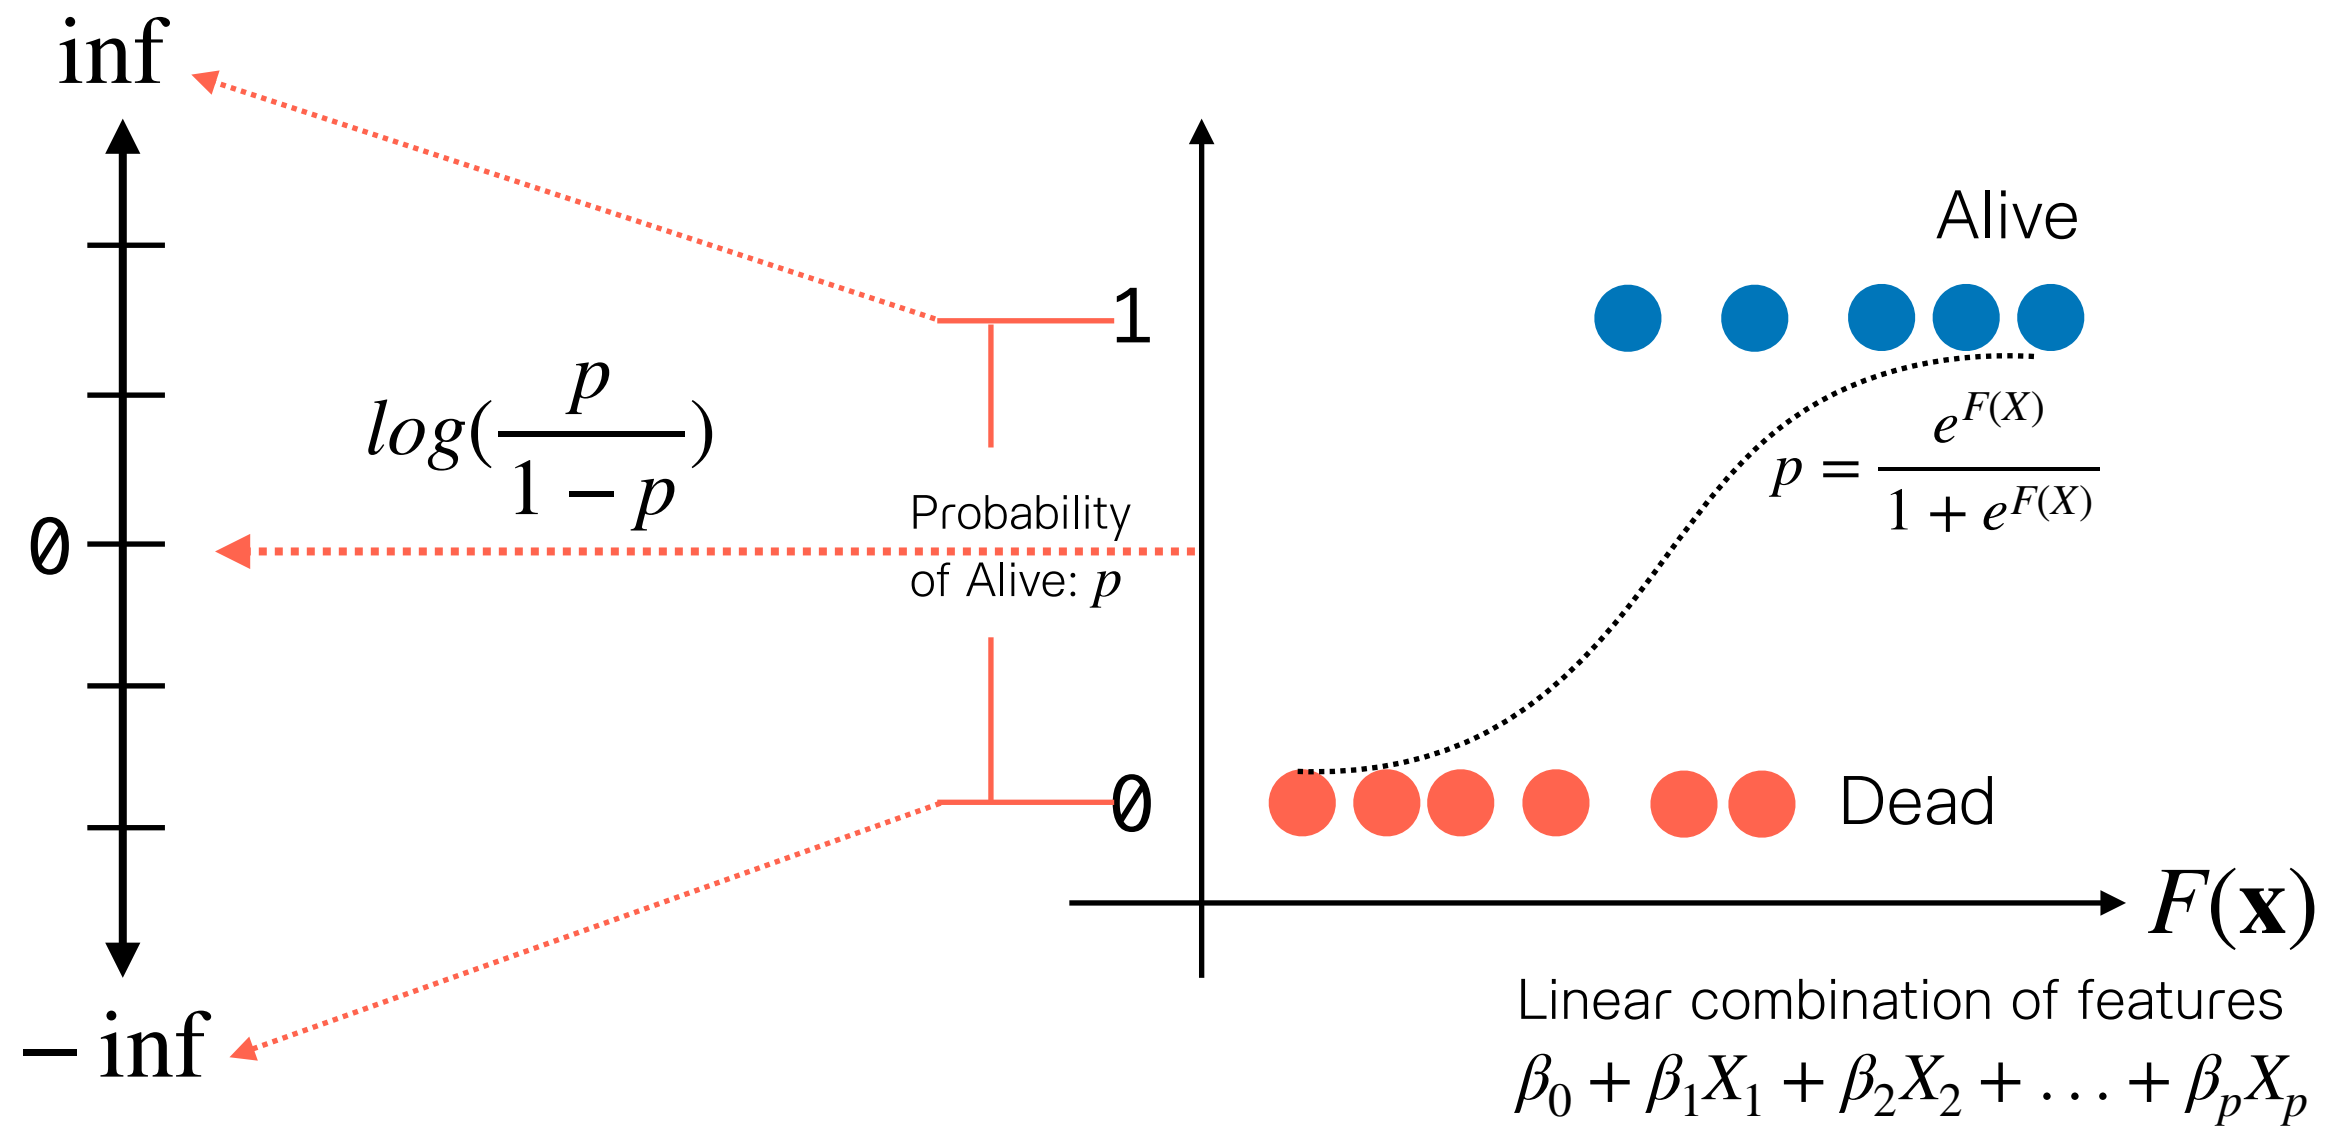
## log-likelihood function

$$\sum_{i=1}^n \{y_i \log(p) + (1 - y_i) \log(1 - p)\}$$

LOC	SYS	AGE	STA	Log-likelihood	
No	36	27	Dead	$\log(0.5) = -0.6931$	← $(0)\log(0.5) + (1 - 0)\log(1 - 0.5)$
No	48	59	Dead	$\log(0.5) = -0.6931$	
Yes	44	77	Dead	$\log(0.5) = -0.6931$	
No	62	54	Dead	$\log(0.5) = -0.6931$	
Yes	112	87	Alive	$\log(0.5) = -0.6931$	← $(1)\log(0.5) + (1 - 1)\log(1 - 0.5)$
Yes	108	69	Alive	$\log(0.5) = -0.6931$	
No	140	63	Alive	$\log(0.5) = -0.6931$	
Yes	138	30	Alive	$\log(0.5) = -0.6931$	

- Since Likelihood is a **goodness of fit** of the model and **log()** function is an **one to one increasing function**, then **larger value of log-likelihood means better prediction**.
- That is why, when building logistic regression, the goal is to maximize the log-likelihood.

# Review



$$\text{Since } \frac{p}{1-p} = e^{F(X)} \implies \log\left(\frac{p}{1-p}\right) = \log(\text{odds}) = F(X)$$

Let  $(X, y) = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$  be a training dataset, and  $L(y_i, F(\mathbf{x}_i))$  be differentiable Loss function.

## Loss function

- Hence, if we want to use the log-likelihood as a Loss function (where a smaller values represent better fitting model), then we have to multiply the log-likelihood by (-1) that is:

$$L(y, F(\mathbf{X})) = - \sum_{i=1}^n \{y_i \log(p_i) + (1 - y_i) \log(1 - p_i)\}$$

Sometimes called  
logLoss or Cross-entropy

Consider per observation of  $(-1) \times \log$ -likelihood

$$\begin{aligned} -[y_i \log(p_i) + (1 - y_i) \log(1 - p_i)] &\implies -y_i \log(p_i) - (1 - y_i) \log(1 - p_i) \\ &\implies -y_i (\log(p_i) - \log(1 - p_i)) - \log(1 - p_i) \\ &\implies -y_i \left[ \log\left(\frac{p_i}{1 - p_i}\right) \right] - \log(1 - p_i) \\ &\implies -y_i [\log(odds_i)] - \log\left[1 - \frac{e^{\log(odds_i)}}{1 + e^{\log(odds_i)}}\right] \end{aligned}$$

$$-y_i[\log(odds)] - \log\left[1 - \frac{e^{\log(odds)}}{1 + e^{\log(odds)}}\right]$$

$$\Rightarrow -y_i[\log(odds)] - \log\left[\frac{1 + e^{\log(odds)} - e^{\log(odds)}}{1 + e^{\log(odds)}}\right]$$

$$\Rightarrow -y_i[\log(odds)] - \log\left[\frac{1}{1 + e^{\log(odds)}}\right]$$

$$\Rightarrow -y_i[\log(odds)] - \log\left[\frac{1}{1 + e^{\log(odds)}}\right]$$

$$\Rightarrow -y_i[\log(odds)] - [\log(1) - \log(1 + e^{\log(odds)})]$$

$$L(y_i, F(\mathbf{x}_i)) = -y_i \log(odds) + \log(1 + e^{\log(odds)}) \leftarrow \text{This is our **Loss function**.}$$

(per observation)

**Note:** the prediction function  $F_i(\mathbf{x}_i) = \log(odds)$

Next step, we need to show that the Loss function is differentiable.

$$\frac{\partial}{\partial \log(odds)} [-y_i \log(odds) + \log(1 + e^{\log(odds)})]$$

$$= -y_i + \frac{e^{\log(odds_i)}}{1 + e^{\log(odds_i)}}$$

$$= -y_i + p_i$$

## Loss function

$$L(y, F(\mathbf{X})) = \sum_{i=1}^n [-y_i \log(odds_i) + \log(1 + e^{\log(odds_i)})]$$

$$\frac{\partial}{\partial \log(odds)} L(y, F(\mathbf{X})) = \sum_{i=1}^n [-y_i + p_i]$$

Let  $(X, y) = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$  be a training dataset, and  $L(y_i, F(\mathbf{x}_i))$  be differentiable Loss function.

**Step 1:** fit the null model (model with only constant term)  $\Rightarrow F_0(\mathbf{X}) = \operatorname{argmin}_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$

Like when using a gradient boost for regression, we have to specify the initial prediction via

Since  $L(y, \gamma) = \sum_{i=1}^n [-y_i \gamma + \log(1 + e^{\gamma})]$ , and  $\frac{\partial}{\partial \gamma} L(y, \gamma) = -\sum_{i=1}^n y_i + np$

$$\text{Let } \frac{\partial}{\partial \gamma} L(y, \gamma) = -\sum_{i=1}^n y_i + np = 0$$

$$\Rightarrow p = \frac{\sum_{i=1}^n y_i}{n}$$

$$p = \frac{0 + 0 + 0 + 0 + 1 + 1 + 1 + 1}{8} = \frac{1}{2}$$

LOC	SYS	AGE	STA
No	36	27	Dead
No	48	59	Dead
Yes	44	77	Dead
No	62	54	Dead
Yes	112	87	Alive
Yes	108	69	Alive
No	140	63	Alive
Yes	138	30	Alive

From the initial predicted probability, we have  $\log(\text{odds}) = \log\left(\frac{p}{1-p}\right) = 0 = F_0(\mathbf{x})$

Let  $(X, y) = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$  be a training dataset, and  $L(y_i, F(\mathbf{x}_i))$  be differentiable Loss function.

**Step 1:** fit the null model (model with only constant term)  $\implies F_0(\mathbf{X}) = \operatorname{argmin}_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$   
 $\implies F_0(\mathbf{X}) = 0$

**Step2:** for  $m$  in  $1 : M$

1. Compute  $e_{im} = - \left[ \frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(x)=F_{m-1}(x)}$  for  $i = 1, 2, \dots, n$
2. Fit a regression tree to the  $e_{im}$  values and create terminal region  $R_{jm}$  for  $j = 1, \dots, J_m$
3. For  $j = 1, 2, \dots, J_m$  compute  $\gamma_{jm} = \operatorname{argmin}_x \sum_{x_i \in R_{ij}} L(y_i, F_{m-1}(x_i) + \gamma)$
4. Update  $F_m(\mathbf{x}) = F_{m-1}(\mathbf{x}) + \nu \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$

Let  $(X, y) = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$  be a training dataset, and  $L(y_i, F(\mathbf{x}_i))$  be differentiable Loss function.

**Step 1:** fit the null model (model with only constant term)  $\implies F_0(\mathbf{X}) = \operatorname{argmin}_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$   
 $\implies F_0(\mathbf{X}) = 0$

**Step2:** for  $m = 1$

1. Compute  $e_{im} = - \left[ \frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(x)=F_{m-1}(x)}$  for  $i = 1, 2, \dots, n$

$$\begin{aligned} e_{i1} &= - \left[ \frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(x)=F_0(x)} \\ &= - \left[ -y_i + \frac{e^0}{1 + e^0} \right] \\ &= y_i - 0.5 \end{aligned}$$

LOC	SYS	AGE	STA	$e_{i1}$
No	36	27	Dead	-0.5
No	48	59	Dead	-0.5
Yes	44	77	Dead	-0.5
No	62	54	Dead	-0.5
Yes	112	87	Alive	0.5
Yes	108	69	Alive	0.5
No	140	63	Alive	0.5
Yes	138	30	Alive	0.5

←  $0 - 0.5$

←  $1 - 0.5$



Let  $(X, y) = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$  be a training dataset, and  $L(y_i, F(\mathbf{x}_i))$  be differentiable Loss function.

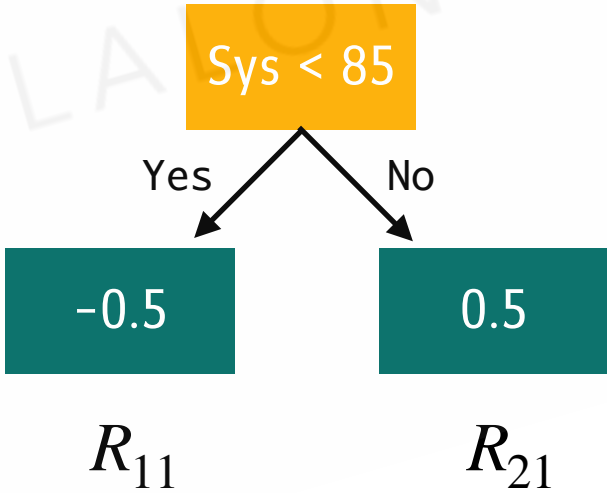
**Step 1:** fit the null model (model with only constant term)  $\implies F_0(\mathbf{X}) = \operatorname{argmin}_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$

$\implies F_0(\mathbf{X}) = 0$

**Step 2:** for  $m = 1$

1. Compute  $e_{im} = - \left[ \frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(x)=F_{m-1}(x)}$  for  $i = 1, 2, \dots, n$
2. Fit a regression tree to the  $e_{im}$  values and create terminal region  $R_{jm}$  for  $j = 1, \dots, J_m$

LOC	SYS	AGE	STA	$e_{i1}$
No	36	27	Dead	-0.5
No	48	59	Dead	-0.5
Yes	44	77	Dead	-0.5
No	62	54	Dead	-0.5
Yes	112	87	Alive	0.5
Yes	108	69	Alive	0.5
No	140	63	Alive	0.5
Yes	138	30	Alive	0.5



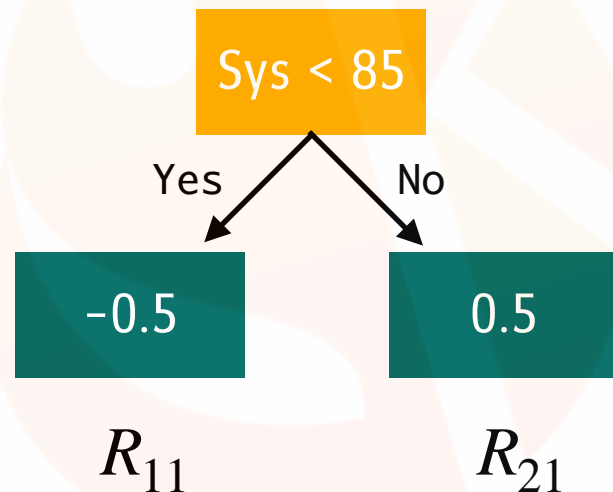
Predict

Let  $(X, y) = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$  be a training dataset, and  $L(y_i, F(\mathbf{x}_i))$  be differentiable Loss function.

**Step 1:** fit the null model (model with only constant term)  $\Rightarrow F_0(\mathbf{X}) = \operatorname{argmin}_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$   
 $\Rightarrow F_0(\mathbf{X}) = 0$

**Step 2:** for  $m = 1$

1. Compute  $e_{im} = - \left[ \frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(x)=F_{m-1}(x)}$  for  $i = 1, 2, \dots, n$
2. Fit a regression tree to the  $e_{im}$  values and create terminal region  $R_{jm}$  for  $j = 1, \dots, J_m$
3. For  $j = 1, 2, \dots, J_m$  compute  $\gamma_{jm} = \operatorname{argmin}_{\gamma} \sum_{x_i \in R_{ij}} L(y_i, F_{m-1}(x_i) + \gamma)$



$$\gamma_{11} = \frac{-0.5}{0.25} = -2 \quad \gamma_{21} = \frac{0.5}{0.25} = 2$$

$$\gamma = \frac{\sum_{i=1}^n e_{im}}{\sum_{i=1}^n p_i(1 - p_i)}, \text{ where } p = \frac{e^{\log(\text{odds})}}{1 + e^{\log(\text{odds})}}$$

Let  $(X, y) = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$  be a training dataset, and  $L(y_i, F(\mathbf{x}_i))$  be differentiable Loss function.

**Step 1:** fit the null model (model with only constant term)  $\implies F_0(\mathbf{X}) = \operatorname{argmin}_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$

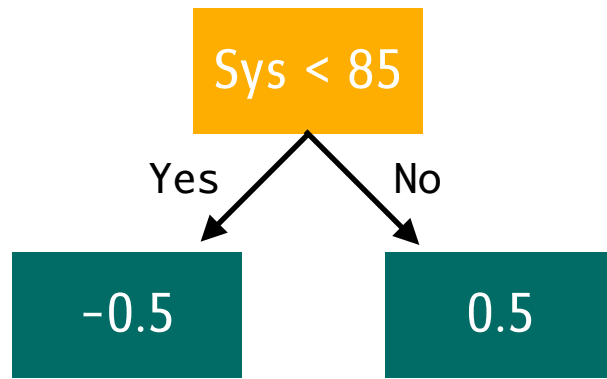
$$\implies F_0(\mathbf{X}) = 0$$

**Step 2:** for  $m = 1$

1. Compute  $e_{im} = - \left[ \frac{\partial L(y_i, F(\mathbf{x}_i))}{\partial F(\mathbf{x}_i)} \right]_{F(x)=F_{m-1}(x)}$  for  $i = 1, 2, \dots, n$
2. Fit a regression tree to the  $e_{im}$  values and create terminal region  $R_{jm}$  for  $j = 1, \dots, J_m$
3. For  $j = 1, 2, \dots, J_m$  compute  $\gamma_{jm} = \operatorname{argmin}_{\gamma} \sum_{x_i \in R_{ij}} L(y_i, F_{m-1}(x_i) + \gamma)$
4. Update  $F_m(\mathbf{x}) = F_{m-1}(\mathbf{x}) + \nu \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$

$$F_1(\mathbf{x}) = F_0(\mathbf{x}) + \nu \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$$

- $\nu \in (0,1)$  is the learning rate
- Learning rate is an effect of the regression tree on the final prediction


 $R_{11}$ 
 $R_{21}$ 

$$\gamma_{11} = \frac{-0.5}{0.25} = -2 \quad \gamma_{21} = \frac{0.5}{0.25} = 2$$

$$F_1(\mathbf{x}) = 0 + (0.9) \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$$

LOC	SYS	AGE	STA	New predicted log(odds)	prob(Alive)
No	36	27	Dead	-1.8	0.141851064900488
No	48	59	Dead	-1.8	0.141851064900488
Yes	44	77	Dead	-1.8	0.141851064900488
No	62	54	Dead	-1.8	0.141851064900488
Yes	112	87	Alive	1.8	0.858148935099512
Yes	108	69	Alive	1.8	0.858148935099512
No	140	63	Alive	1.8	0.858148935099512
Yes	138	30	Alive	1.8	0.858148935099512

repeat ...

# Component-wise Gradient boosting

The diagram illustrates the process of component-wise gradient boosting. It shows three stages of adding weak learners  $b_1$ ,  $b_2$ , and  $b_3$  to a base function  $f_0$ . In each stage, one component is highlighted in orange to show its contribution to the next function. The final equation is  $f_3 = f_0 + \beta \sum_{(3,3,1)} b_i$ .



1. Initialize the function estimate  $\hat{f}^{[0]}$  with offset values. Note that  $\hat{f}^{[0]}$  is a vector of length  $n$ . In the following paragraphs, we will generally denote the vector of function estimates at iteration  $m$  by  $\hat{f}^{[m]}$ .

2. Specify a set of *base-learners*. Base-learners are simple regression estimators with a fixed set of input variables and a univariate response. The sets of input variables are allowed to differ among the base-learners. Usually, the input variables of the base-learners are small subsets of the set of predictor variables  $x_1, \dots, x_p$ . For example, in the simplest case, there is exactly one base-learner for each predictor variable, and the base-learners are just simple linear models using the predictor variables as input variables. Generally, the base-learners considered in this paper are either penalized or unpenalized least squares estimators using small subsets of the predictor variables as input variables (see Section 3.2.1 for details and examples). Each base-learner represents a modeling alternative for the statistical model. Denote the number of base-learners by  $P$  and set  $m = 0$ .

3. Increase  $m$  by 1, where  $m$  is the number of iterations.

4. a) Compute the negative gradient  $-\frac{\partial \rho}{\partial f}$  of the loss function and evaluate it at  $\hat{f}^{[m-1]}(\mathbf{x}_i^\top)$ ,  $i = 1, \dots, n$  (i.e., at the estimate of the previous iteration). This yields the negative gradient vector

$$\mathbf{u}^{[m]} = \left( u_i^{[m]} \right)_{i=1, \dots, n} := \left( -\frac{\partial}{\partial f} \rho \left( y_i, \hat{f}^{[m-1]}(\mathbf{x}_i^\top) \right) \right)_{i=1, \dots, n}.$$

b) Fit each of the  $P$  base-learners to the negative gradient vector, i.e., use each of the regression estimators specified in step 2 separately to fit the negative gradient. The resulting  $P$  regression fits yield  $P$  vectors of predicted values, where each vector is an estimate of the negative gradient vector  $\mathbf{u}^{[m]}$ .

c) Select the base-learner that fits  $\mathbf{u}^{[m]}$  best according to the residual sum of squares (RSS) criterion and set  $\hat{\mathbf{u}}^{[m]}$  equal to the fitted values of the best-fitting base-learner.

d) Update the current estimate by setting  $\hat{f}^{[m]} = \hat{f}^{[m-1]} + \nu \hat{\mathbf{u}}^{[m]}$ , where  $0 < \nu \leq 1$  is a real-valued step length factor.

5. Iterate Steps 3 and 4 until the stopping iteration  $m_{\text{stop}}$  is reached (the choice of  $m_{\text{stop}}$  is discussed below).

# Package ‘compboost’

October 28, 2018

**Type** Package

**Title** C++ Implementation of Component-Wise Boosting

**Version** 0.1.0

**Maintainer** Daniel Schalk <daniel.schalk@stat.uni-muenchen.de>

**Description** C++ implementation of component-wise boosting implementation of component-wise boosting written in C++ to obtain high runtime performance and full memory control. The main idea is to provide a modular class system which can be extended without editing the source code. Therefore, it is possible to use R functions as well as C++ functions for custom base-learners, losses, logging mechanisms or stopping criteria.

**License** MIT + file LICENSE

**Encoding** UTF-8

**LazyData** true

**RoxygenNote** 6.1.0

**Imports** Rcpp (>= 0.11.2), methods, glue, R6, checkmate

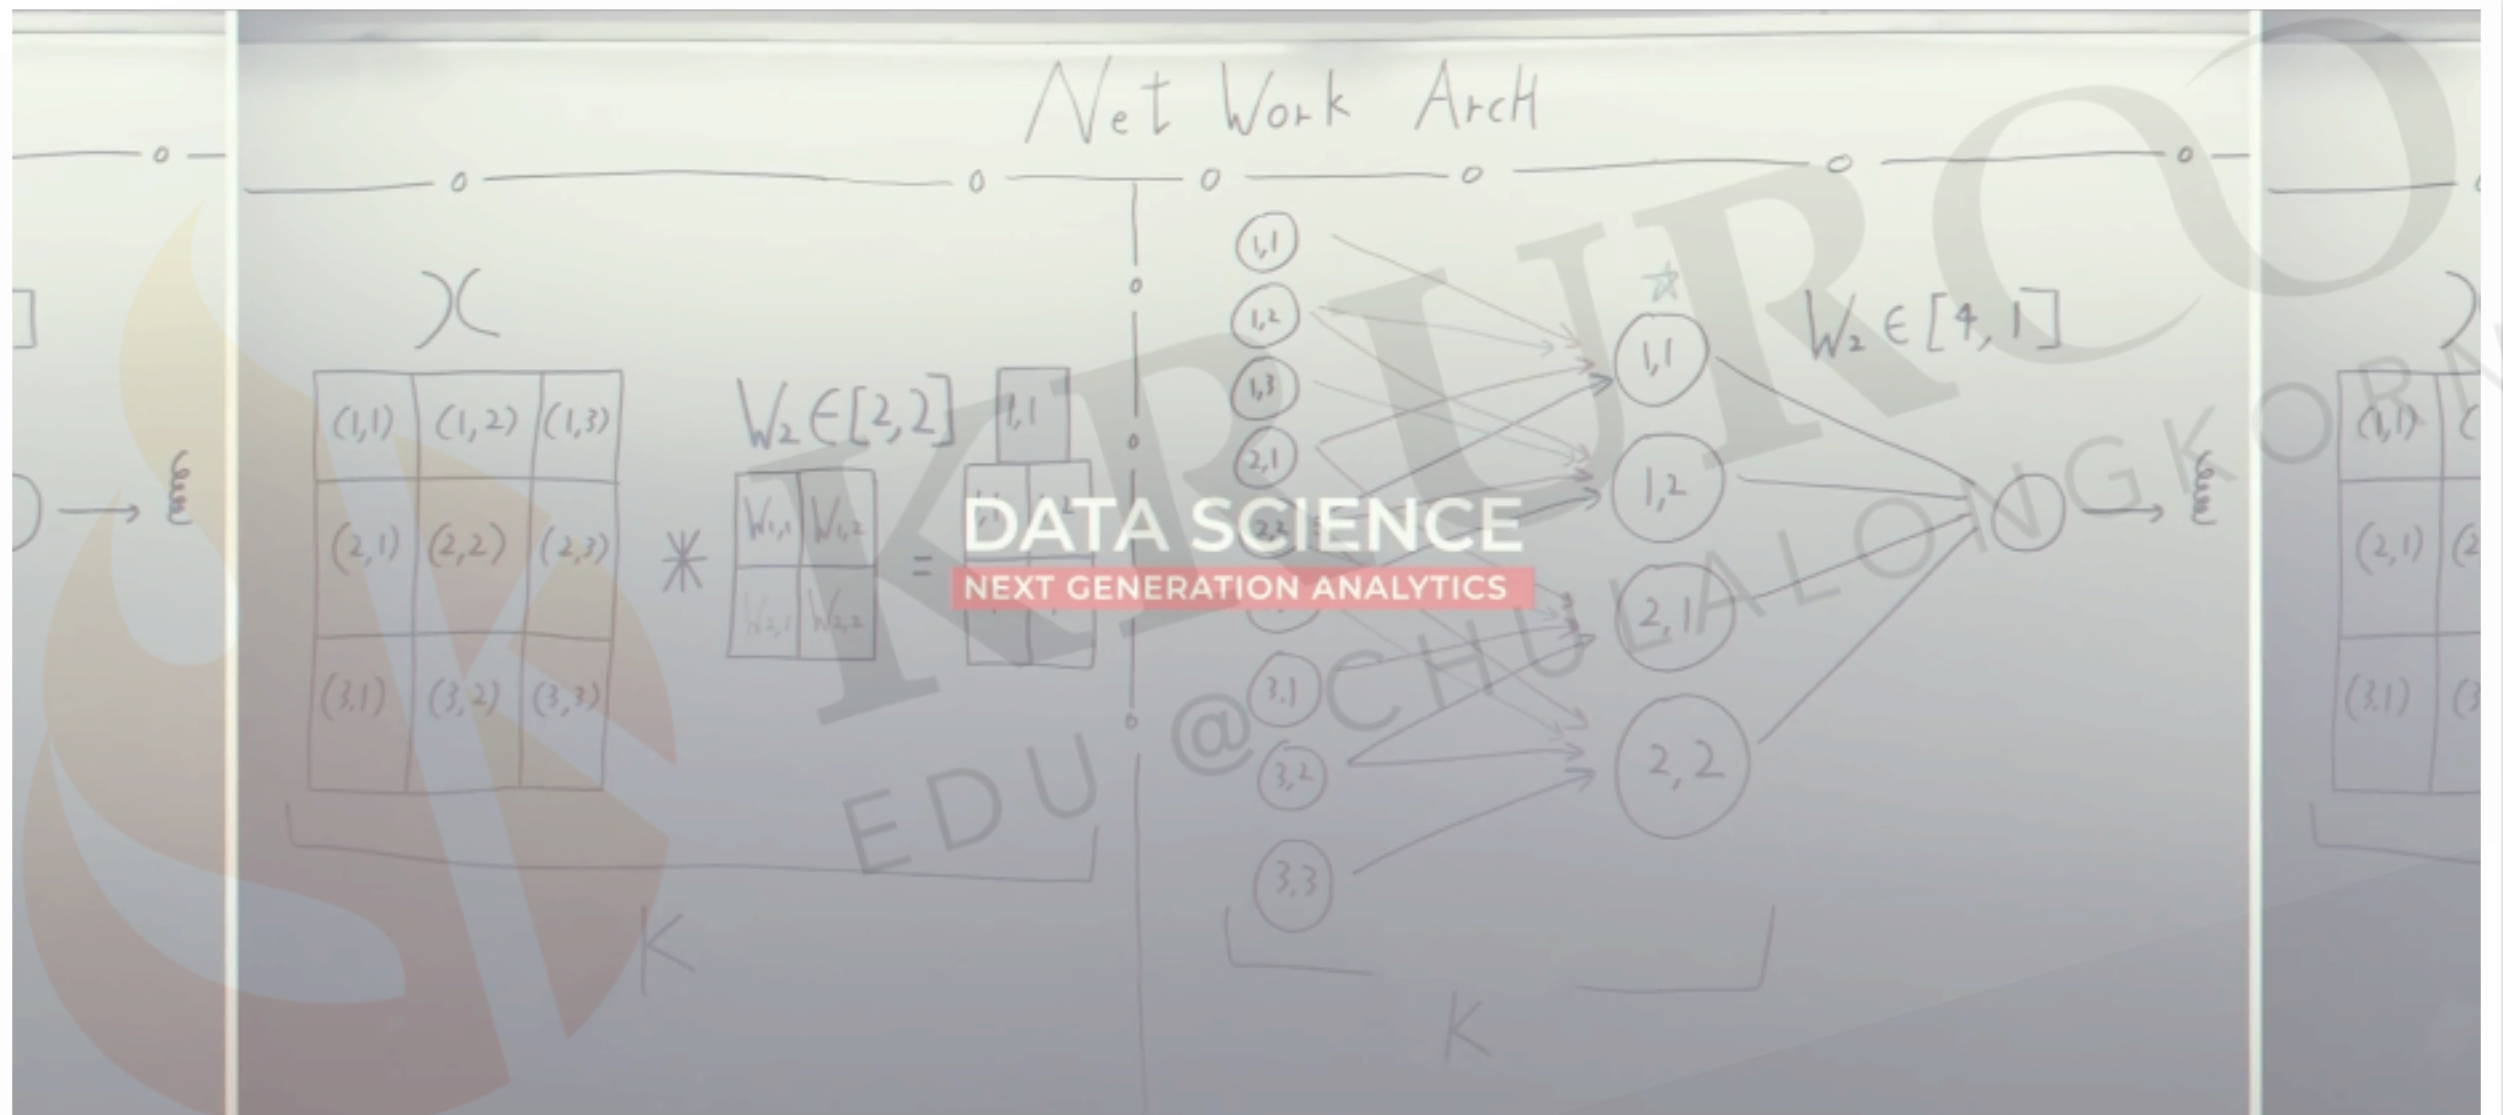
**LinkingTo** Rcpp, RcppArmadillo

**Suggests** RcppArmadillo (>= 0.9.100.5.0), ggplot2, testthat, rpart,

```
> boostLinear()  
> boostSplines()
```



# Automated Feature Selection via component-wise boosting algorithm



<https://www.youtube.com/watch?v=ucSt28PPUPY&t=1264s>

<https://github.com/STATWORX/bounceR>

# Feature Selection

- Features contain information about the outcome variable.
- More features = more information?  
= better prediction performance?



(STATWORX, 2018)

# Feature Selection

- Irrelevant features
- Redundant feature



KRUROO  
EDU @ CHULALONGKORN

# Feature Selection methods

- **Filter methods** - using bivariate statistic
- **Wrapper methods**
  - iteratively searching for the best feature set
  - Model-based approach

<https://www.youtube.com/watch?v=ucSt28PPUPY&t=1264s>

<https://topepo.github.io/caret/variable-importance.html#model-specific-metrics>

**Result:** Stability Score Distribution for Feature Space  $S$

**foreach**  $n\_rounds$  in  $N\_ROUNDS$  **do**

    initialize Stability Score Distribution  $\Omega^j$ ;

**foreach**  $n\_mods$  in  $N\_MODS$  **do**

        draw  $s_i$  features from  $S^*$ ;

        bootstrap  $q_i$  obs from  $Q$  obs;

        run componentwise boosting on set  $(q_i, s_i)$ ;

        penalize score in  $\Omega^j$  for  $s_i^-$  for not surviving boosting run ;

        reward score in  $\Omega^j$  for  $s_i^+$  for surviving boosting run ;

**if** *score for feature  $k$  < threshold* **then**

            | cut feature  $k$  from  $S^*$

**return** Stability Score Distribution  $\Omega_{updated}^j$

**end**

**return** average stability score distribution  $\bar{\Omega}$

**end**



```
# load devtools
install.packages(devtools)
library(devtools)

# download from our public repo
devtools::install_github("STATWORX/bounceR")

# source it
library(bounceR)
```

## About Our Algorithm: Usage

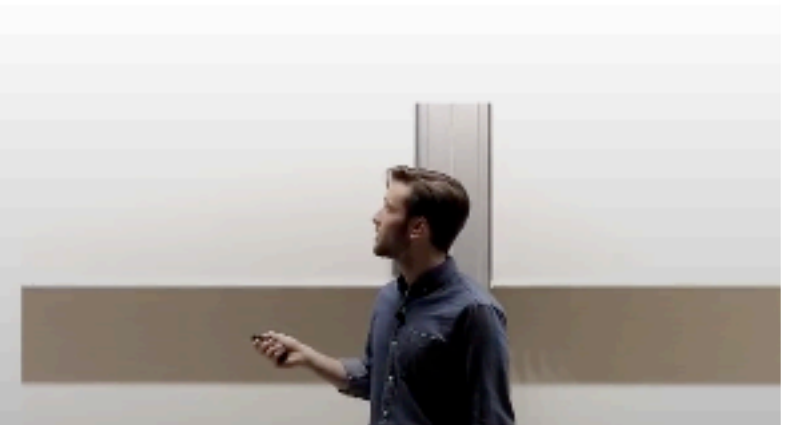
Sure, you have a lot of tuning parameters, however we put them all together in a nice and handy little interface. By the way, we set the defaults based on several simulation studies, so you can - sort of - trust them - sometimes.

```
# Feature Selection using bounceR-----
selection <- featureSelection(data = train_df,
                             target = "target",
                             index = NULL,
                             selection = selectionControl(n_rounds = 100,
                                                         n_mods = 1000,
                                                         p = NULL,
                                                         reward = 0.2,
                                                         penalty = 0.3,
                                                         max_features = NULL),
                             bootstrap = "regular",
                             boosting = boostingControl(mstop = 100, nu = 0.1),
                             early_stopping = "aic",
                             n_cores = 6)
```

## About Our Package: Content

The package contains a variety of useful functions surrounding the topic of feature selection, such as:

- Convenience:
  - `sim_data`: a function simulating regression and classification data, where the true feature space is known
- Filtering:
  - `featureFiltering`: a function implementing several popular filter methods for feature selection
- Wrapper:
  - `featureSelection`: a function implementing our home grown algorithm for feature selection
- Methods:
  - `print.sel_obj`: an S4 printing method for the object class "sel\_obj"
  - `plot.sel_obj`: an S4 plotting method for the object class "sel\_obj"
  - `summary.sel_obj`: an S4 summary method for the object class "sel\_obj"
  - `builder`: method to extract a formula with n features from a "sel\_obj"



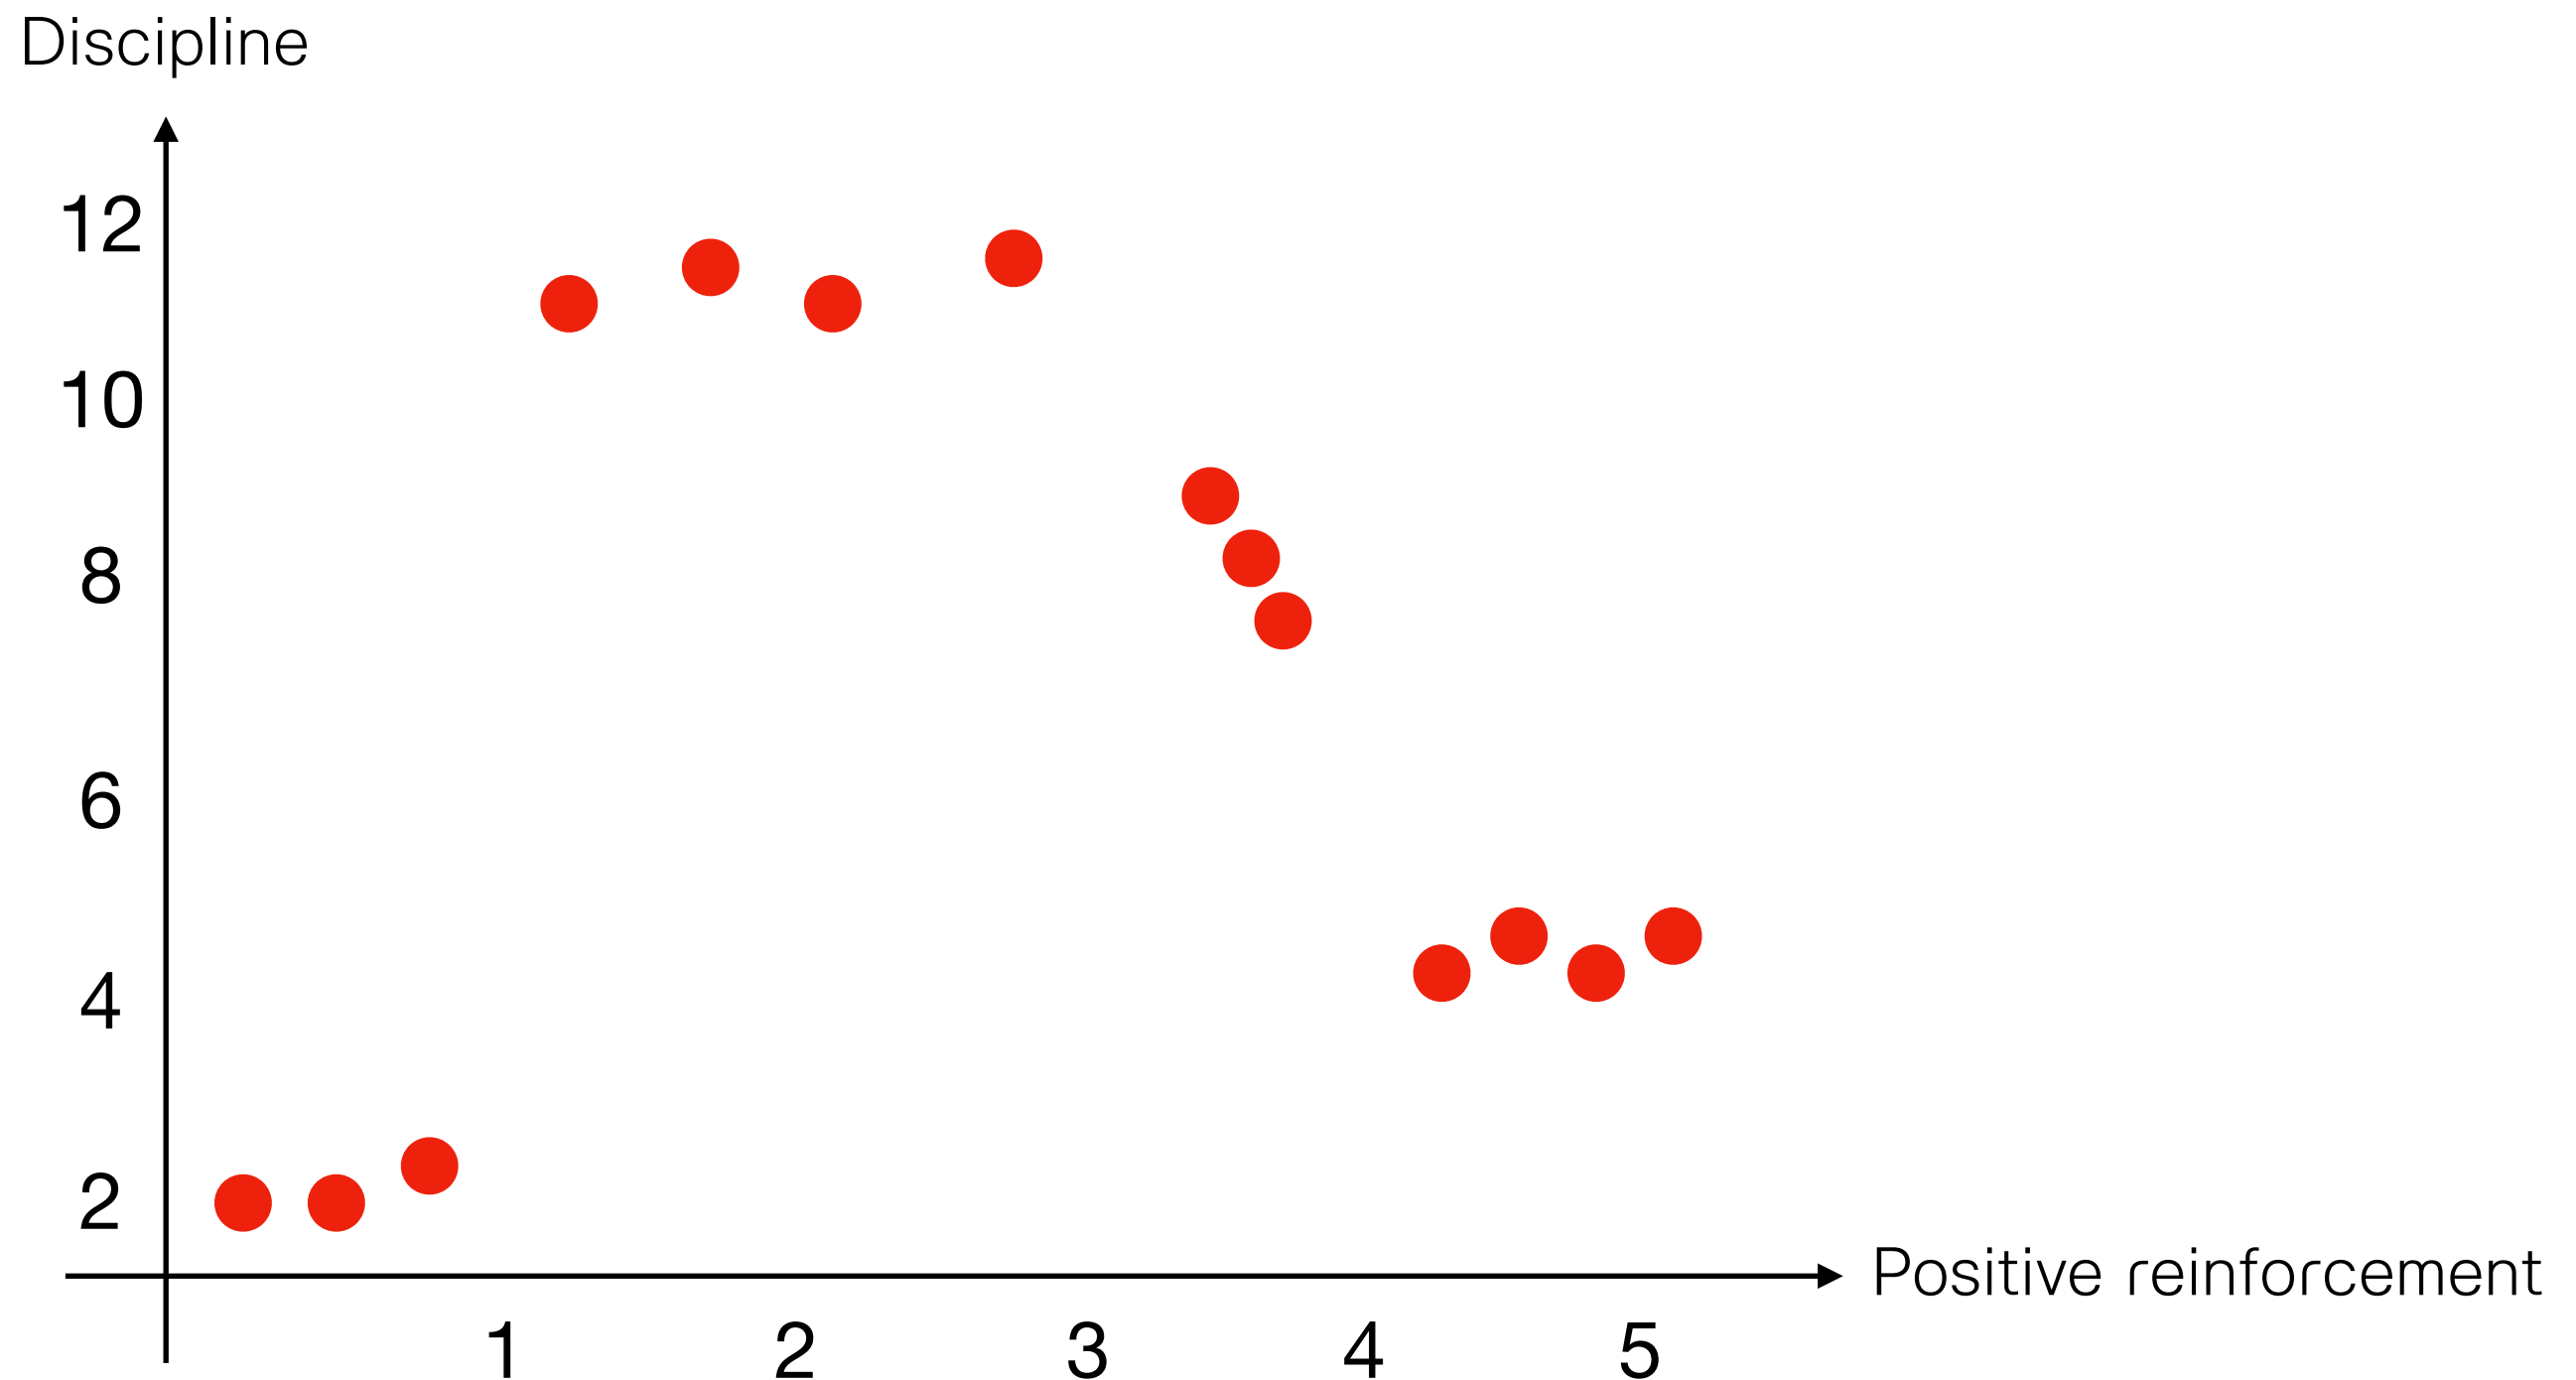
# XGBoost Trees

---

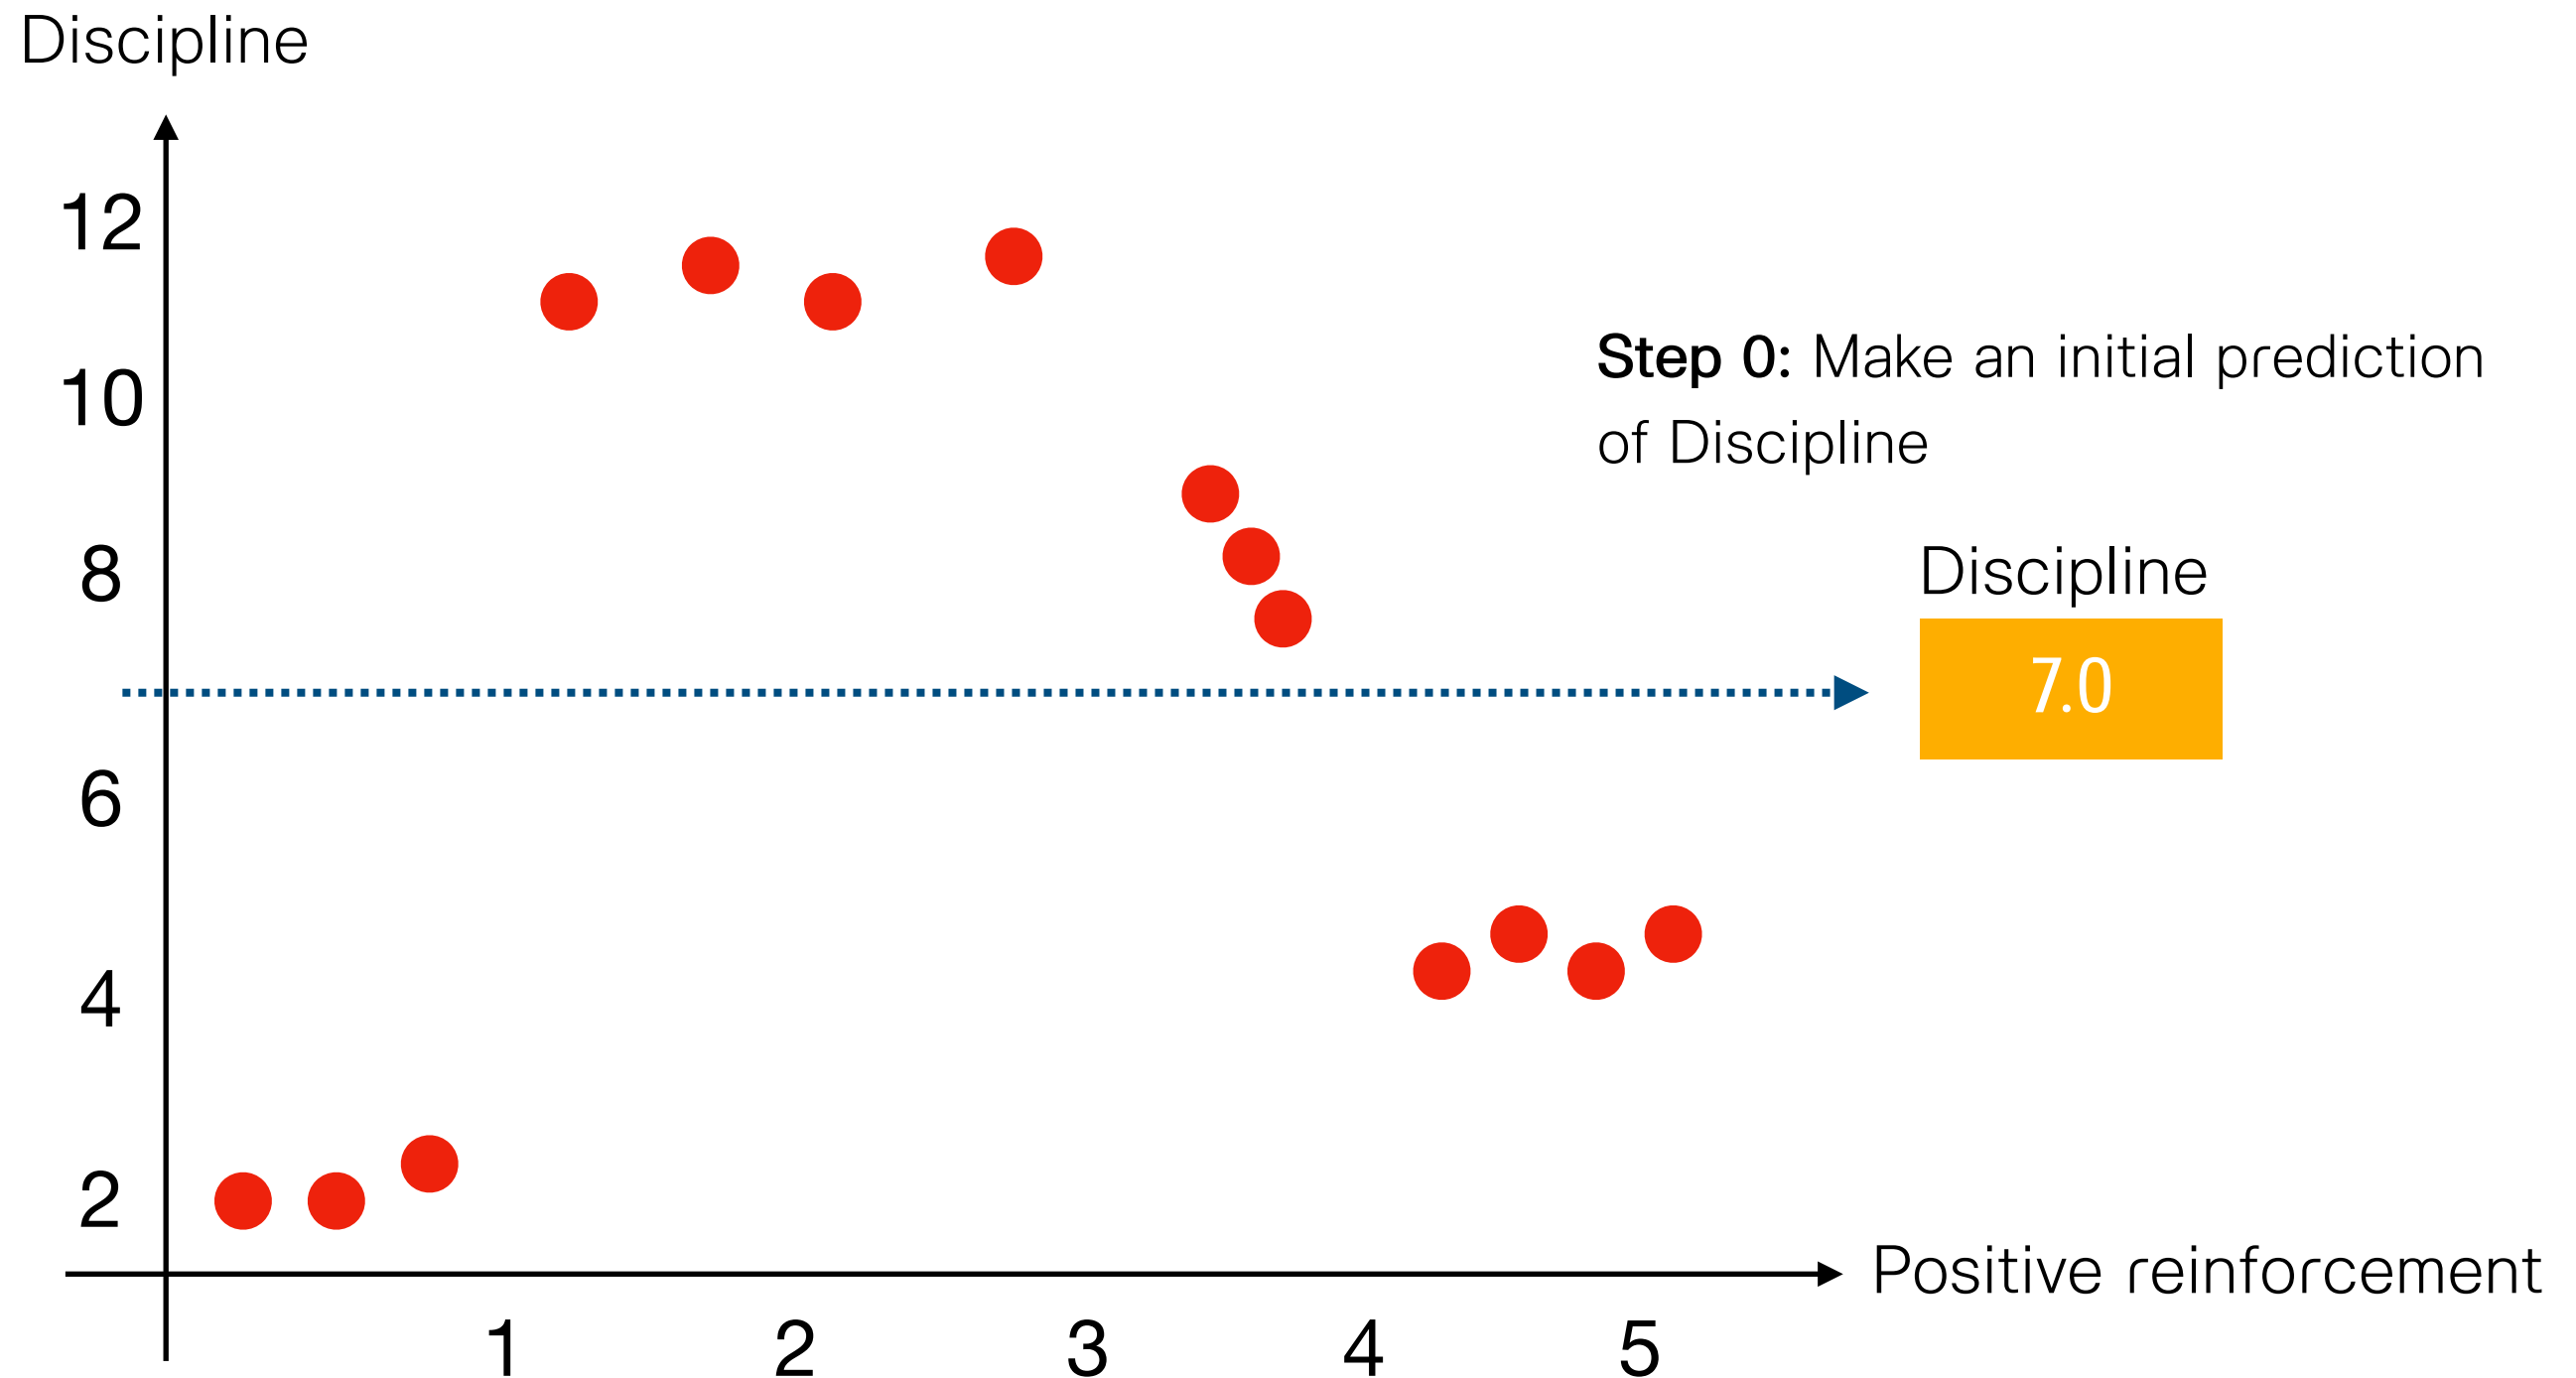
- Regression
- Classification



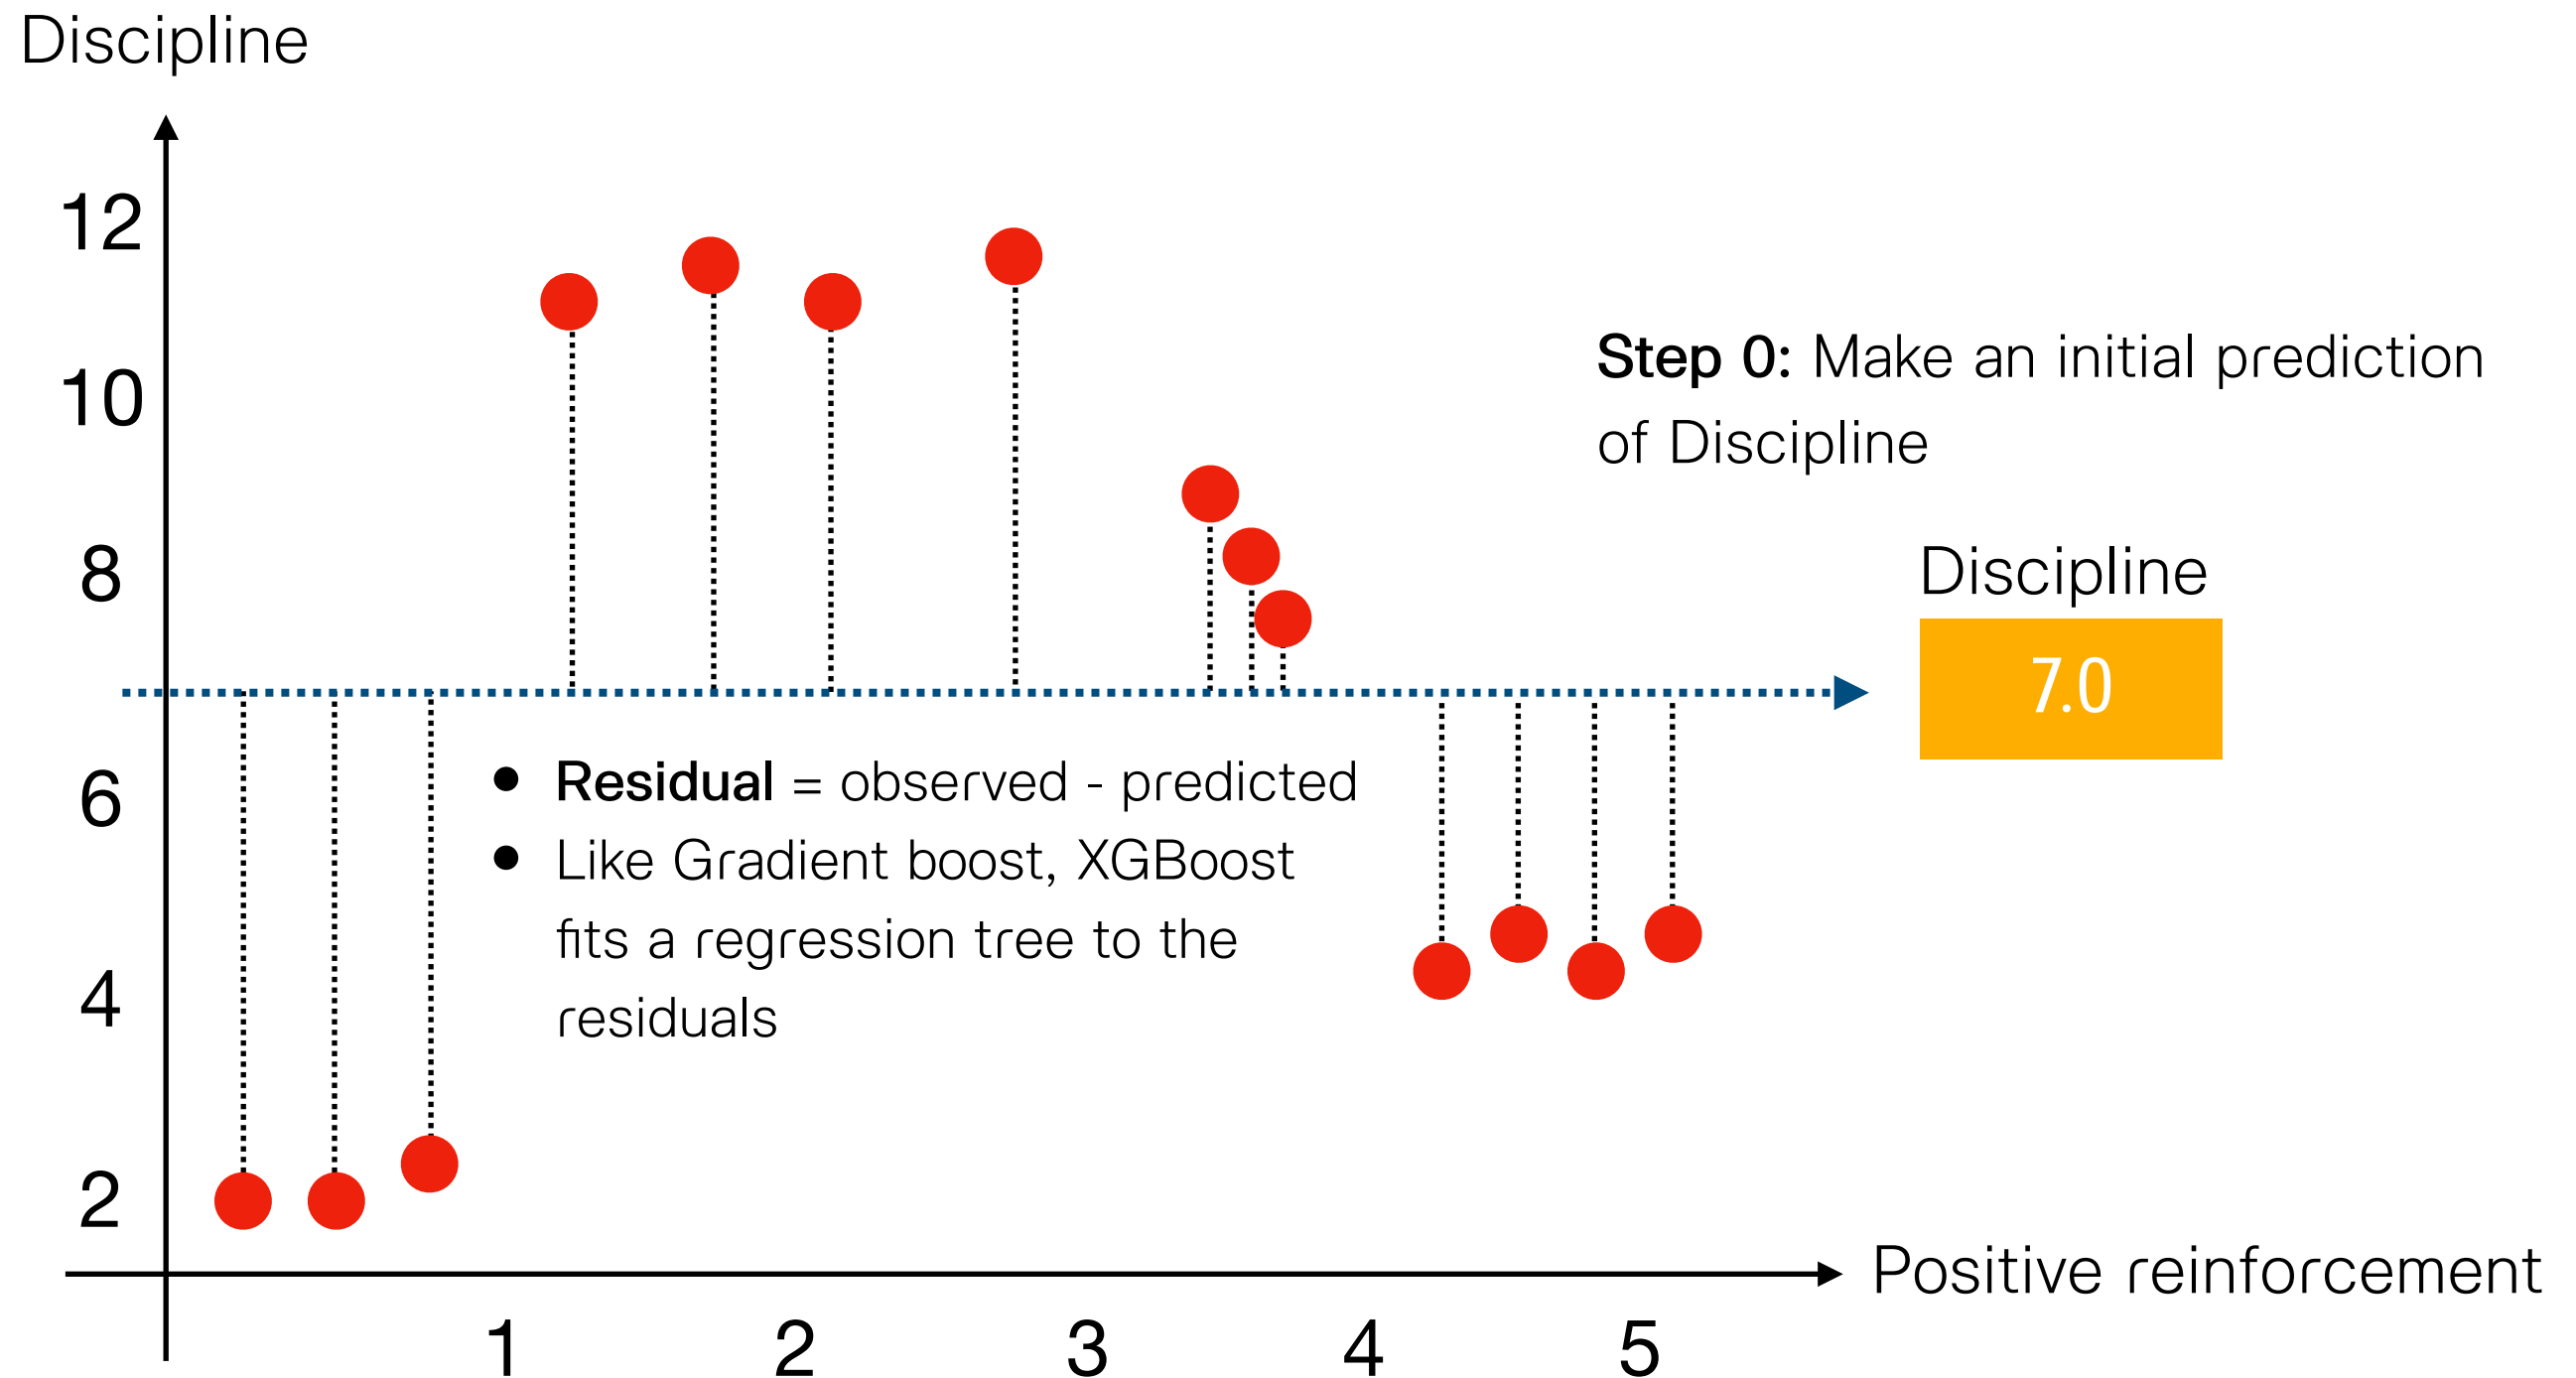
# XGBoost Trees for Regression



# XGBoost Trees for Regression

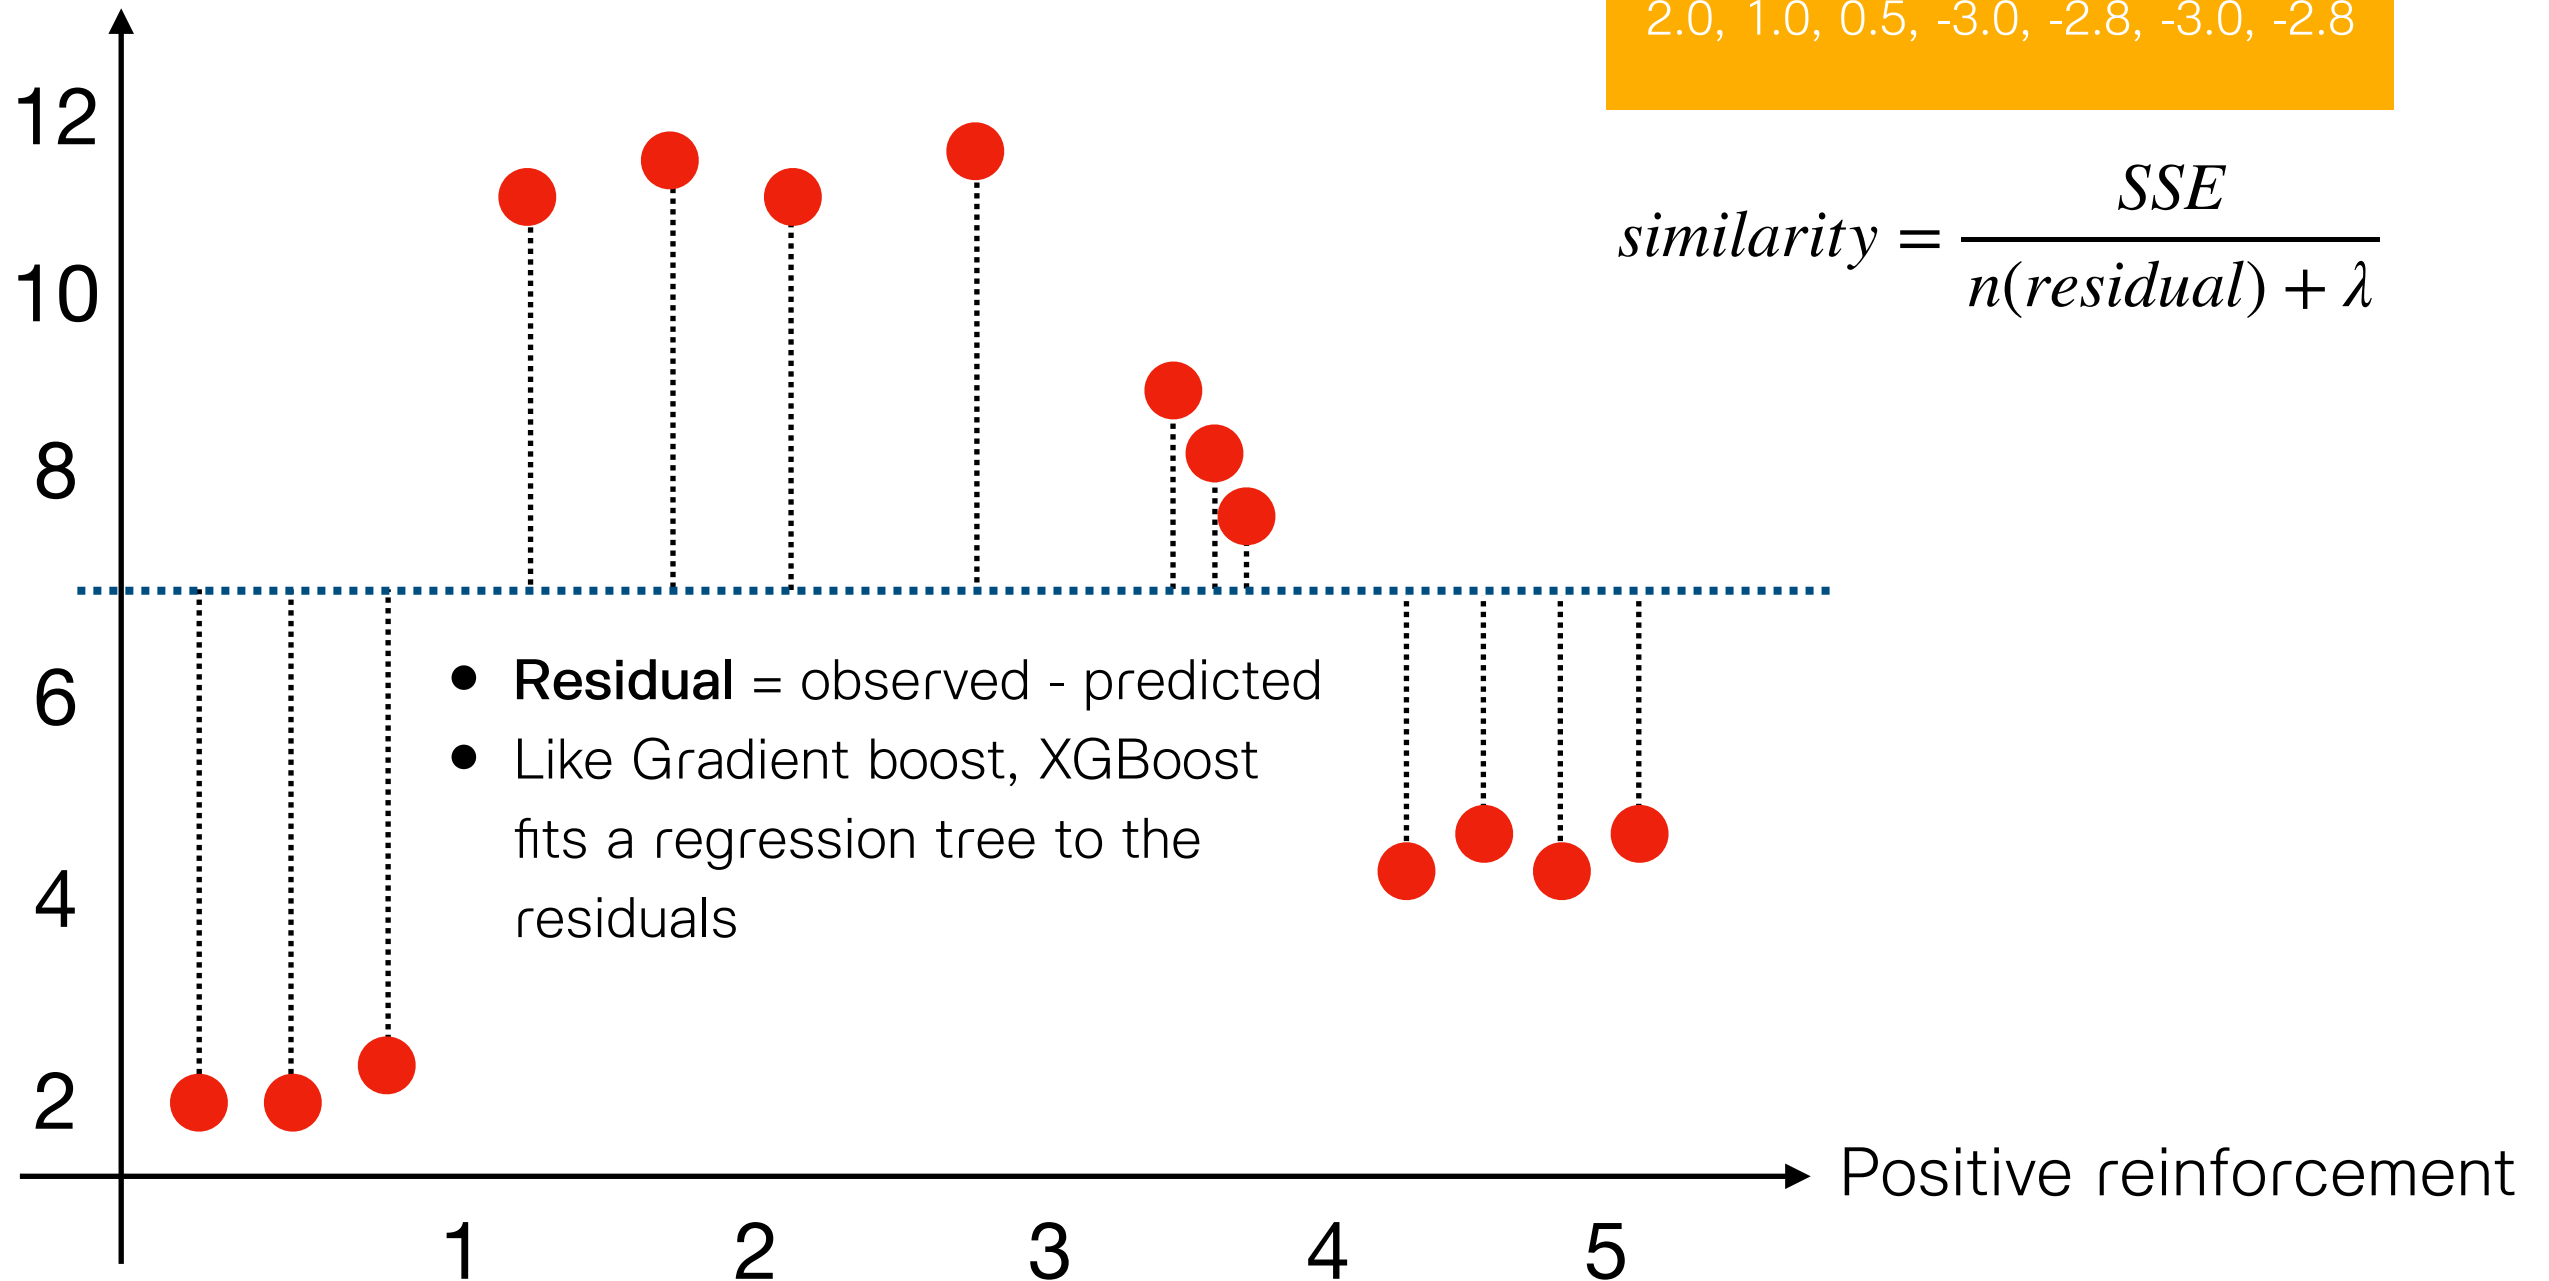


# XGBoost Trees for Regression



# XGBoost Trees for Regression

Discipline

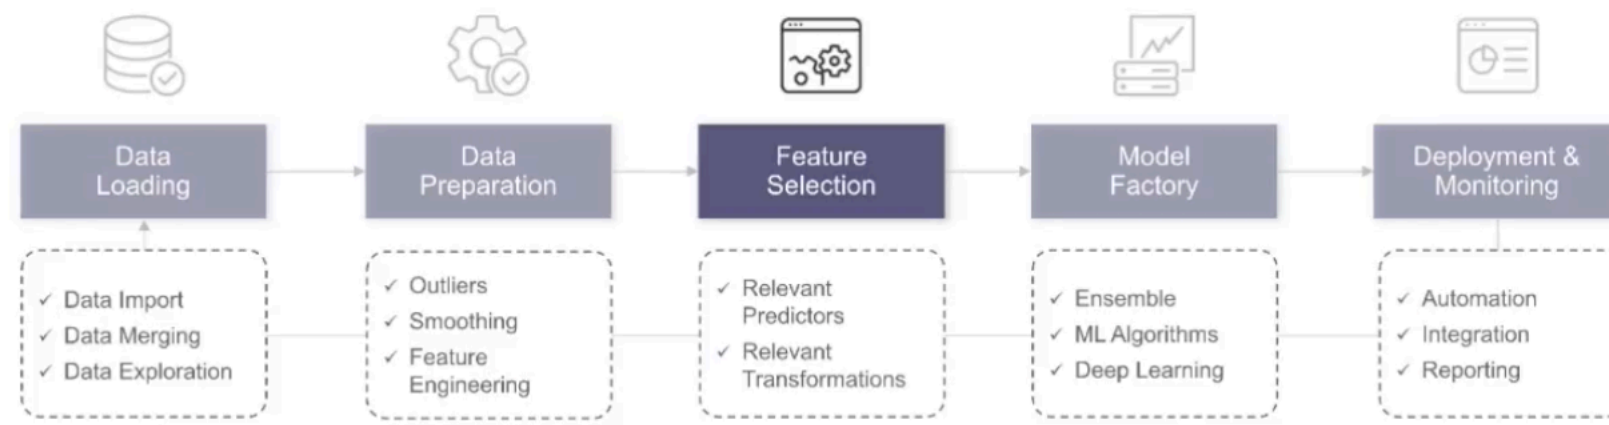
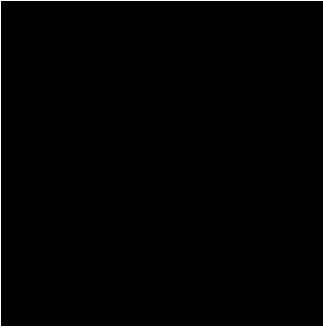


# Support Vector Machine

---

- Introduction
- Support Vector Classification - Linear Kernels
- Polynomial Kernels
- Radius Basis Function Kernels

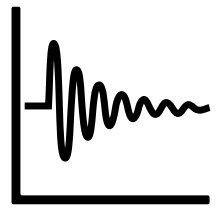




# Feature Selection Problem



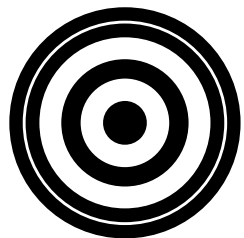
# Why do we need feature selection?



With too many features, most models can't be identified or are noisy



The right combination of features boosts performance, not algorithm.



Features are all information you have to predict your outcome.

# How can we do feature selection?

- Manually selecting (testing) all feature set.
- Using correlation matrix to filter the relevant features.
- Loop over all feature set combinations and evaluate them based on performance gain



# Gradient Boost for Regression

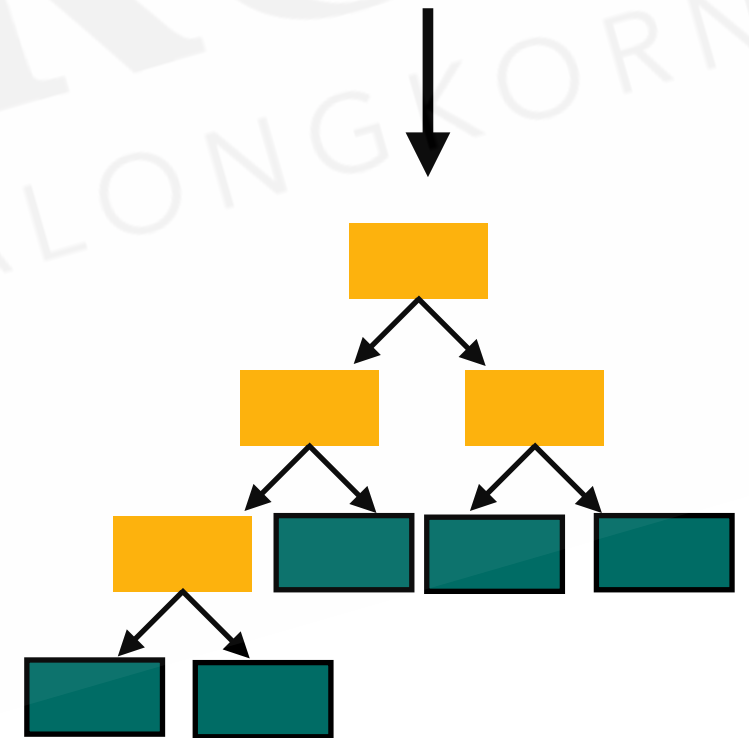
```
> head(dat)
```

	R.D.Spend	Administration	Marketing.Spend	State	Profit
1	165349.2	136897.80	471784.1	New York	192261.8
2	162597.7	151377.59	443898.5	California	191792.1
3	153441.5	101145.55	407934.5	Florida	191050.4
4	144372.4	118671.85	383199.6	New York	182902.0
5	142107.3	91391.77	366168.4	Florida	166187.9
6	131876.9	99814.71	362861.4	New York	156991.1

Average  
value of  
profit

112,013.00

- Then Gradient Boost builds a tree.
- Like AdaBoost, this tree is based on the error made by the previous tree.
- In contrast with AdaBoost, this tree is usually larger than a stump.



## Step1: Build up the initial leaf

SES	Level	Gender	Discipline
-1.0	Primary	Male	27
0.5	Secondary	Female	59
1.0	UnderGrad	Female	77
1.5	Primary	Male	54
1.5	UnderGrad	Male	87
1.0	Secondary	Female	69

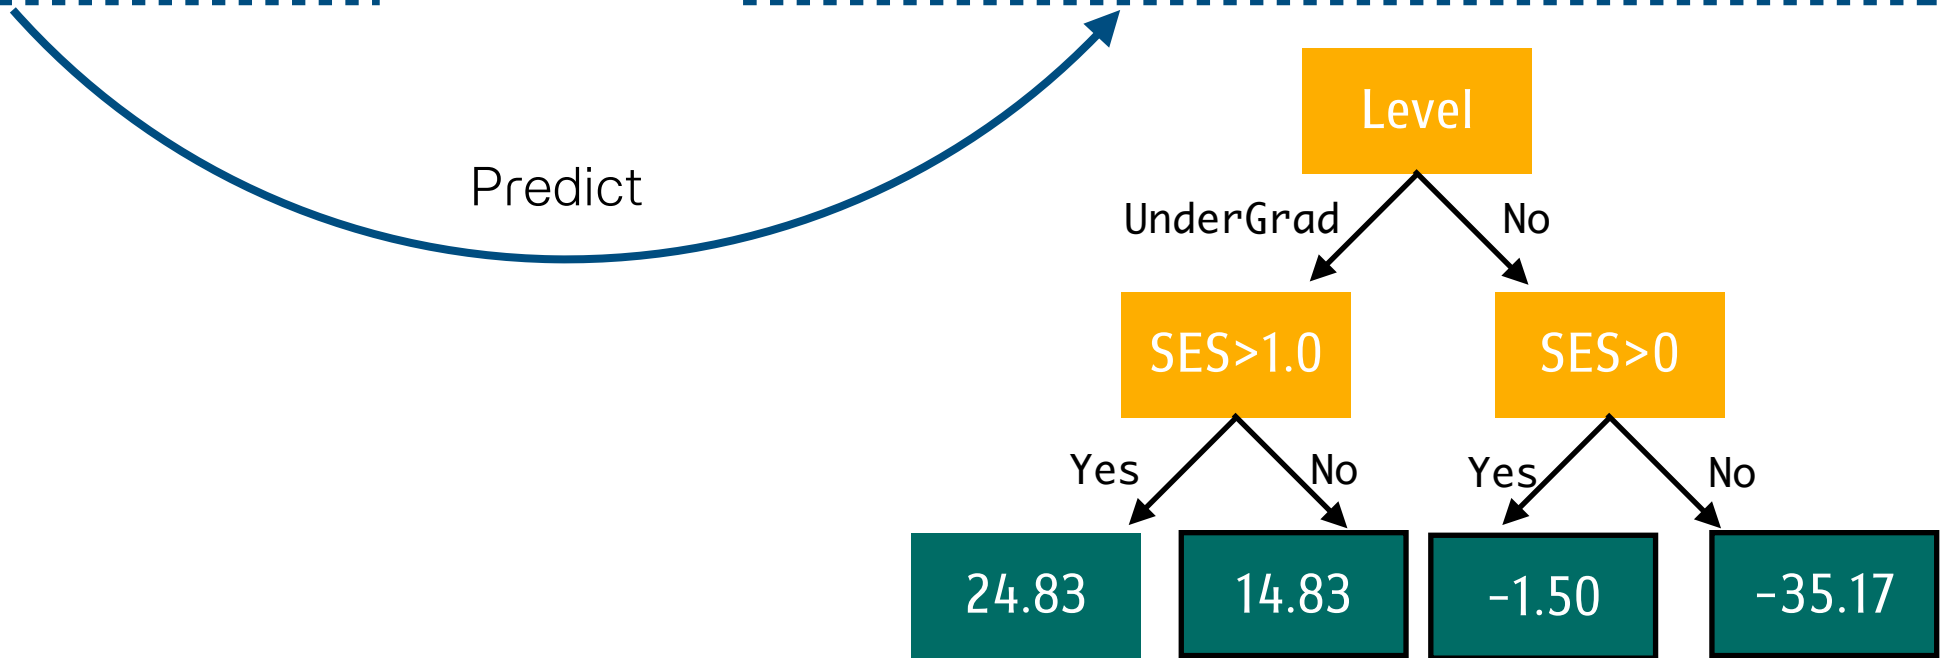


Average value of  
Discipline score

62.17

Step2: Build a tree based on the error from the first tree.

SES	Level	Gender	Discipline	(Pseudo)Residual = actual - predicted
-1.0	Primary	Male	27	-35.17
0.5	Secondary	Female	59	-3.17
1.0	UnderGrad	Female	77	14.83
1.5	Primary	Male	54	-8.17
1.5	UnderGrad	Male	87	24.83
1.0	Secondary	Female	69	6.83



# Gradient Boosting using R

```
> install.packages("gbm")  
> library(gbm)
```





**Step2:** Build a tree based on the error from the first tree.

SES	Level	Gender	Learning Discipline
-1.0	Primary	Male	27
0.5	Secondary	Female	59
1.0	UnderGrad	Female	77
1.5	Primary	Male	54
1.5	UnderGrad	Male	87
1.0	Secondary	Female	69



Average  
value of  
LD score

62.17

**Objective:** we will create a forest of stumps with AdaBoost to predict a patient status (1 = alive or 0 = dead)

1. Give an **equal weight** to each data point.

```
> head(dat,10)
```

	ID	STA	AGE	CAN	SYS	TYP	LOC	Weight1
1	8	0	27	0	142	1	0	0.005
2	12	0	59	0	112	1	0	0.005
3	14	0	77	0	100	0	0	0.005
4	28	0	54	0	142	1	0	0.005
5	32	0	87	0	110	1	0	0.005
6	38	0	69	0	110	1	0	0.005
7	40	0	63	0	104	0	0	0.005
8	41	0	30	0	144	1	0	0.005
9	42	0	35	0	108	1	0	0.005
10	50	0	70	1	138	0	0	0.005

**Objective:** we will create a forest of stumps with AdaBoost to predict a patient status (1 = alive or 0 = dead)

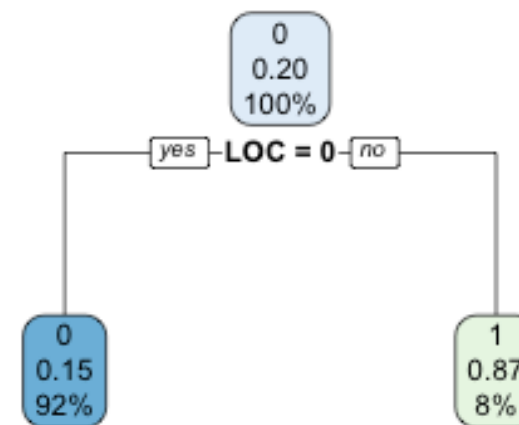
```
> head(dat,10)
```

	ID	STA	AGE	CAN	SYS	TYP	LOC	Weight1
1	8	0	27	0	142	1	0	0.005
2	12	0	59	0	112	1	0	0.005
3	14	0	77	0	100	0	0	0.005
4	28	0	54	0	142	1	0	0.005
5	32	0	87	0	110	1	0	0.005
6	38	0	69	0	110	1	0	0.005
7	40	0	63	0	104	0	0	0.005
8	41	0	30	0	144	1	0	0.005
9	42	0	35	0	108	1	0	0.005
10	50	0	70	1	138	0	0	0.005

1. Give an **equal weight** to each data point.

$$weight = \frac{1}{m} = \frac{1}{200}$$

2. Make the first stump



**Objective:** we will create a forest of stumps with AdaBoost to predict a patient status (1 = alive or 0 = dead)

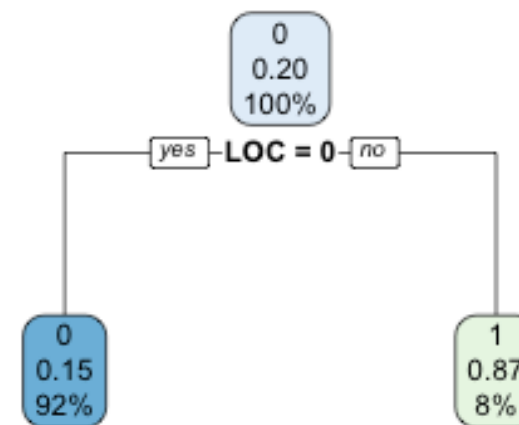
```
> head(dat,10)
```

	ID	STA	AGE	CAN	SYS	TYP	LOC	Weight1
1	8	0	27	0	142	1	0	0.005
2	12	0	59	0	112	1	0	0.005
3	14	0	77	0	100	0	0	0.005
4	28	0	54	0	142	1	0	0.005
5	32	0	87	0	110	1	0	0.005
6	38	0	69	0	110	1	0	0.005
7	40	0	63	0	104	0	0	0.005
8	41	0	30	0	144	1	0	0.005
9	42	0	35	0	108	1	0	0.005
10	50	0	70	1	138	0	0	0.005

1. Give an **equal weight** to each data point.

$$weight = \frac{1}{m} = \frac{1}{200}$$

2. Make the first stump



3. Evaluate classification error of LOC

Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	158	27
1	2	13



**Objective:** we will create a forest of stumps with AdaBoost to predict a patient status (1 = alive or 0 = dead)

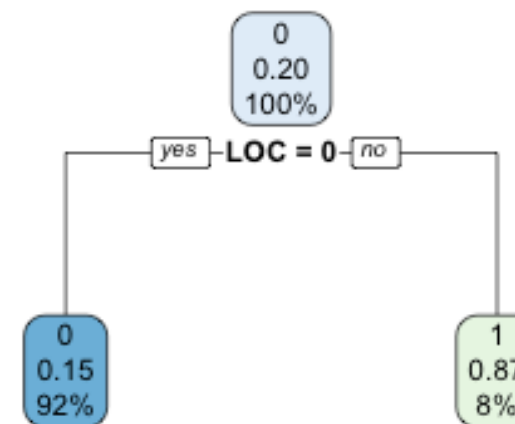
```
> head(dat,10)
```

	ID	STA	AGE	CAN	SYS	TYP	LOC	Weight1
1	8	0	27	0	142	1	0	0.005
2	12	0	59	0	112	1	0	0.005
3	14	0	77	0	100	0	0	0.005
4	28	0	54	0	142	1	0	0.005
5	32	0	87	0	110	1	0	0.005
6	38	0	69	0	110	1	0	0.005
7	40	0	63	0	104	0	0	0.005
8	41	0	30	0	144	1	0	0.005
9	42	0	35	0	108	1	0	0.005
10	50	0	70	1	138	0	0	0.005

1. Give an **equal weight** to each data point.

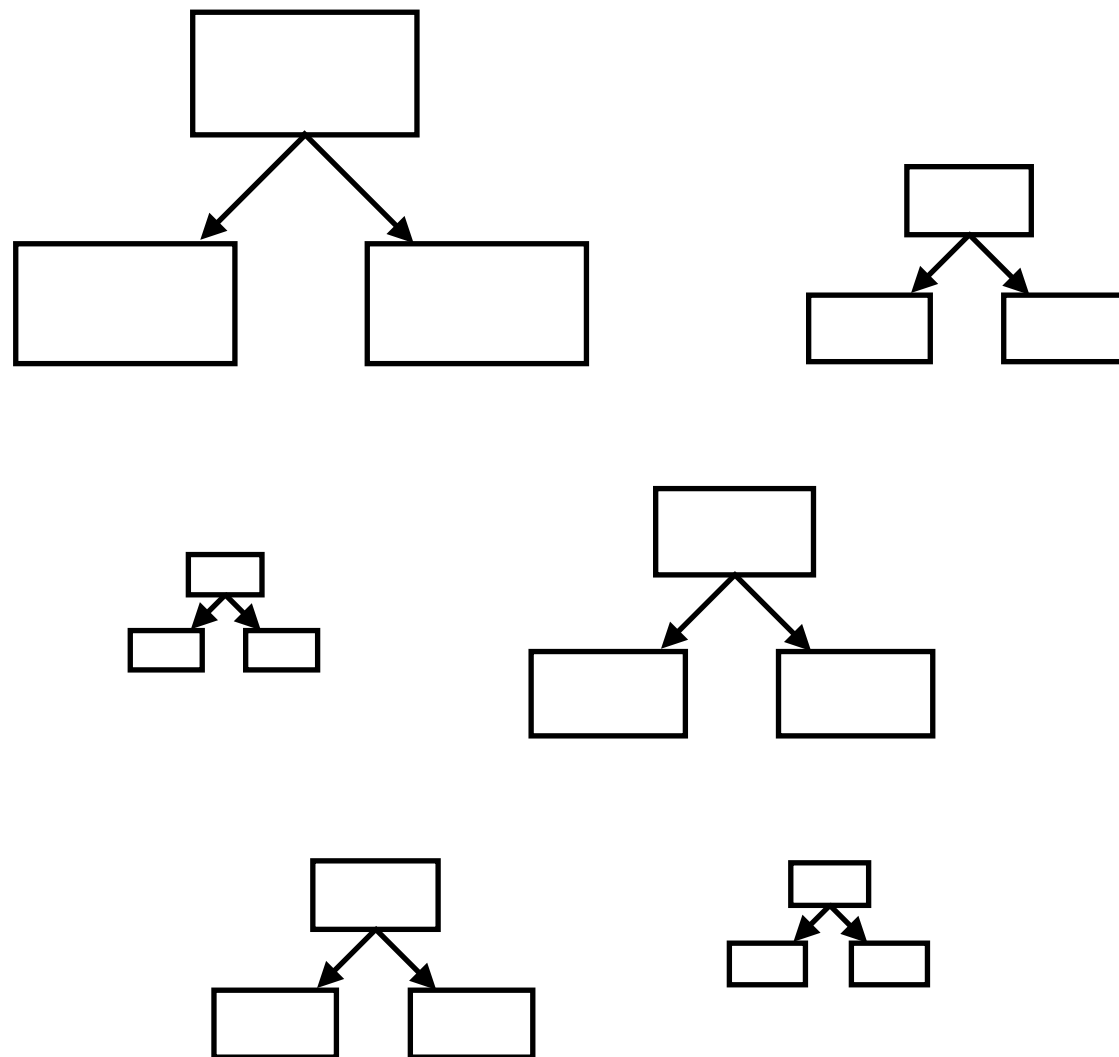
$$weight = \frac{1}{m} = \frac{1}{200}$$

2. Make the first stump



3. Evaluate classification error of LOC

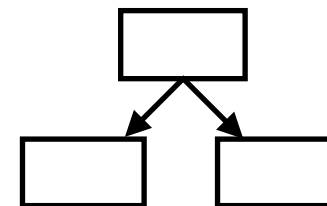
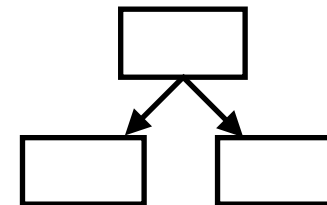
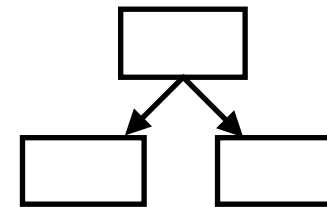
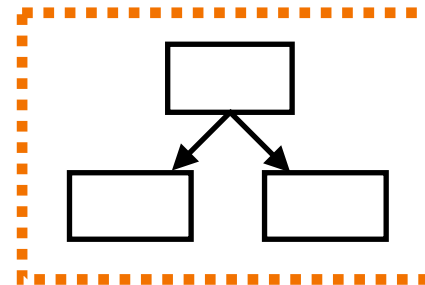
# Main concepts behind AdaBoost



Forest of stumps made with AdaBoost, some stumps get more weight in the final classification than others.

# Main concepts behind AdaBoost

Called “stump”



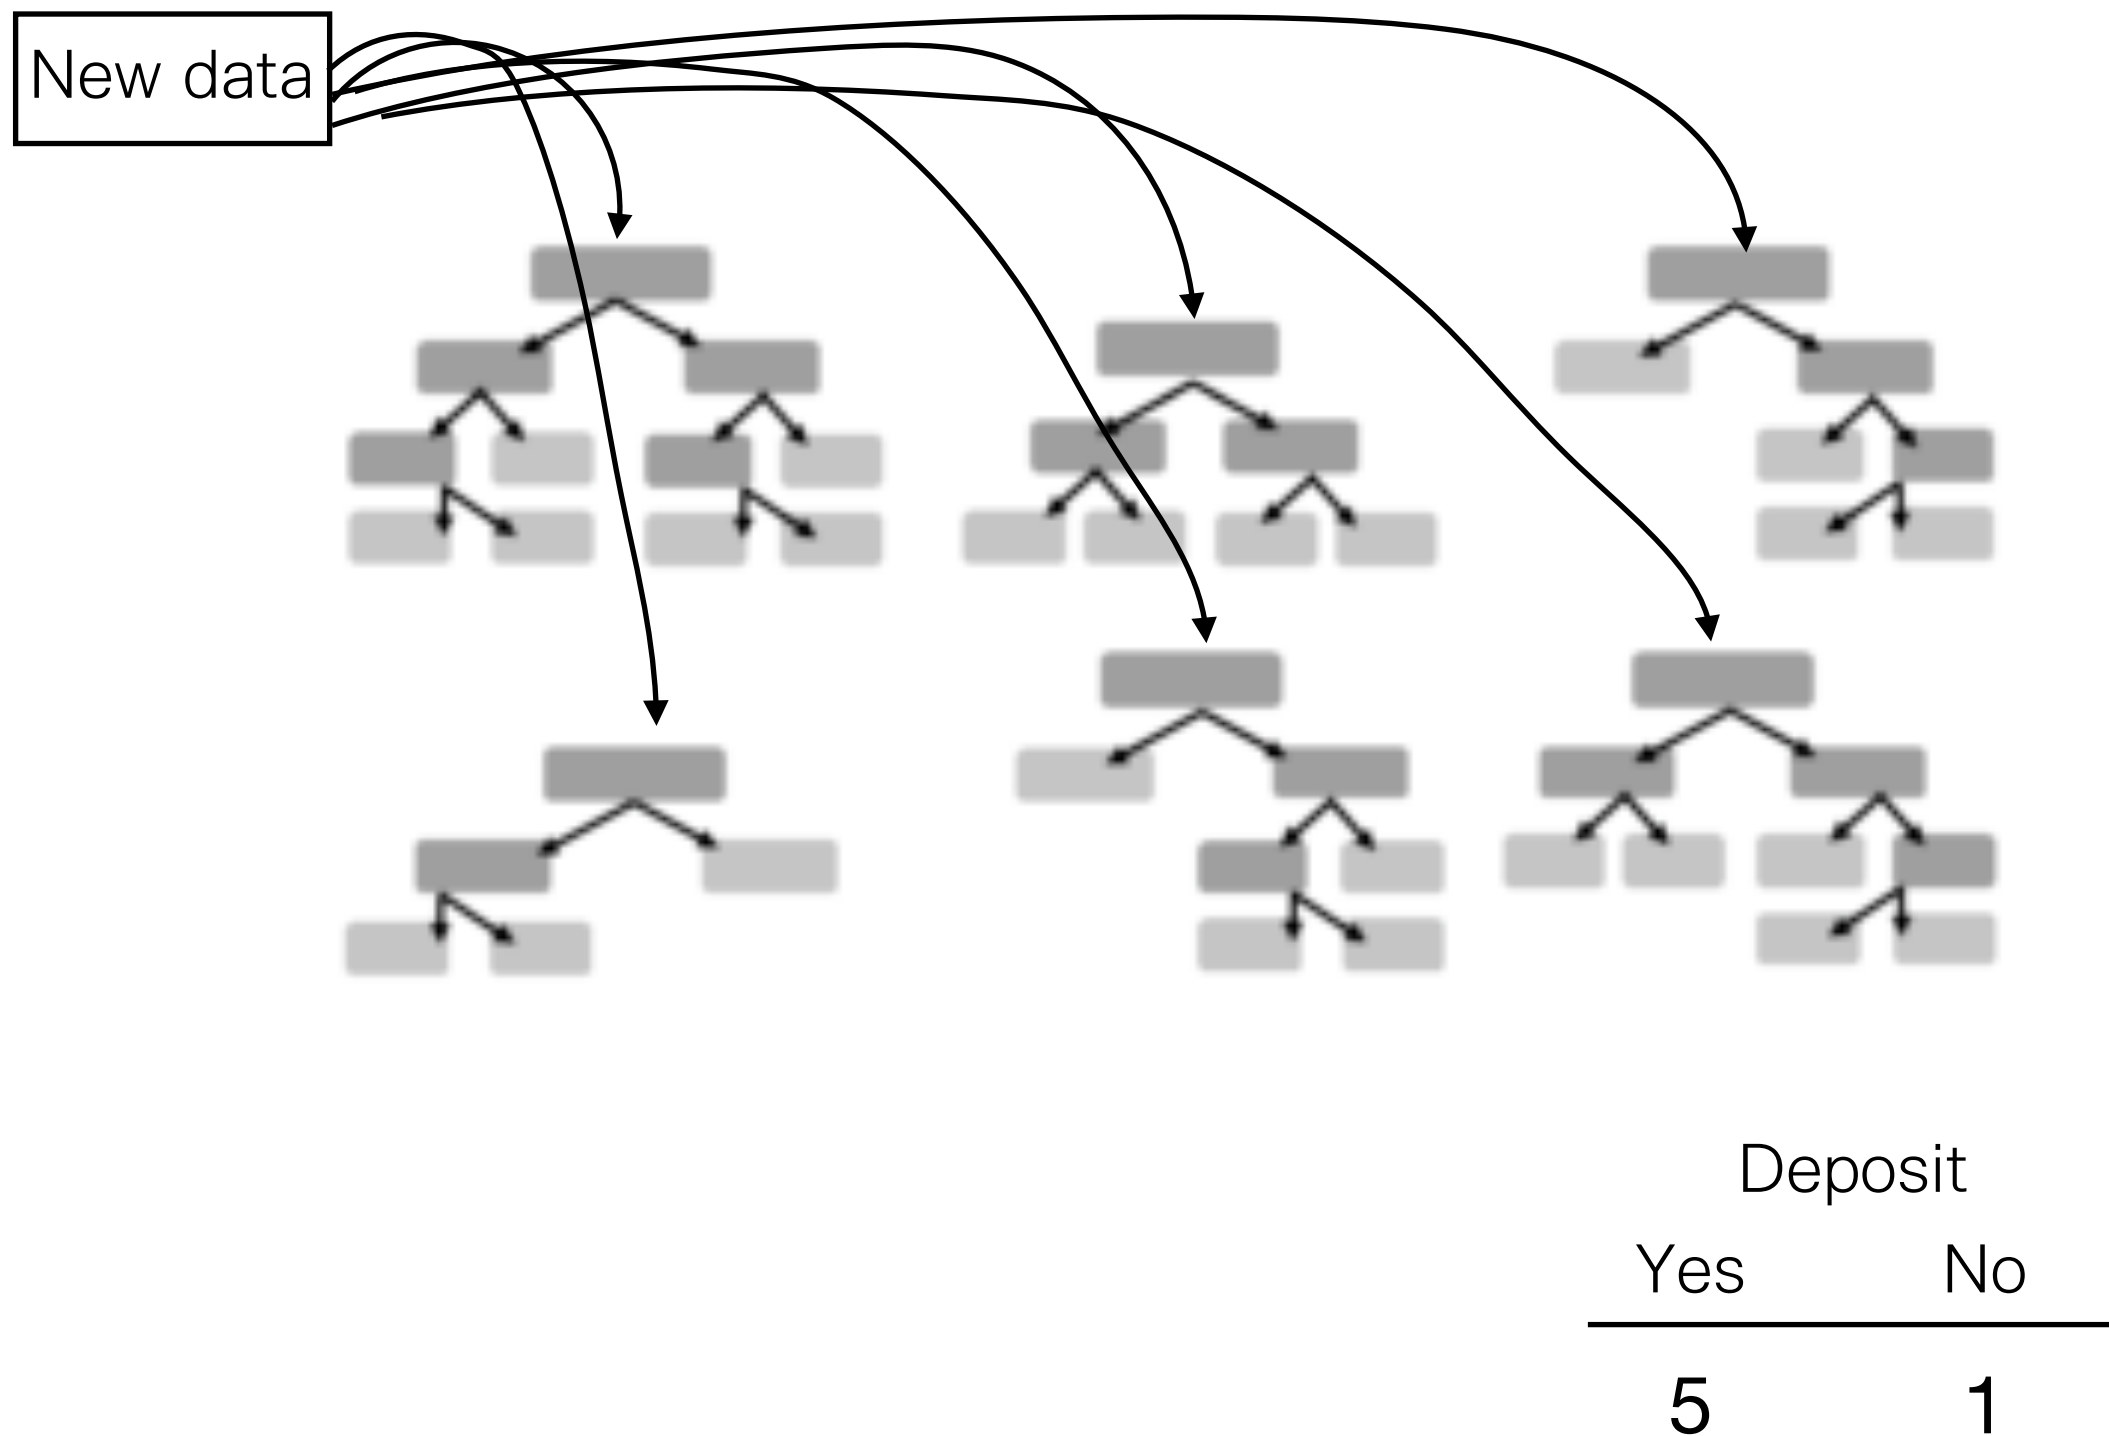
...

AdaBoost





# Random Forest



# A Formal view on AdaBoost

1. Given a training dataset  $(X, y)$  which contains data points  $(x_1, y_1), (x_2, y_2), \dots, (x_m, y_m)$
2. For iteration  $t = 1, 2, \dots, T$ 
  - Construct distribution of  $W_t$  on  $\{1, 2, 3, \dots, m\}$ , where  $W_t(i)$  is the weight attributed to subject  $i$  on the iteration  $t$ 
    - Let  $W_1(i) = 1/m$
  - Build up a classifier  $M_t : X \rightarrow \{0, 1\}$
  - Compute the error associated to the classifier, let  $e_t(i) = 1$  if  $M_t(i) \neq y_i$ , and  $e_t(i) = 0$  if  $M_t(i) = y_i$



KRUROO  
EDU @ CHULALONGKORN

In a random forest, each time you make a tree, you make a full sized tree.







# Out-of-bag error vs Out-sample error

- OOB error only estimate errors (not AUC, sens, spec,...)
- Can't compare OOB error to other types of models



# Estimate the accuracy of a random forest

- Out of bag error

1. Create a “bootstrapped” dataset

# Tree models

Root Node

Discipline score  
> 55.

True

False

Internal Nodes

Time to  
study

Time to  
study

a lot of time

less

less

a lot of time

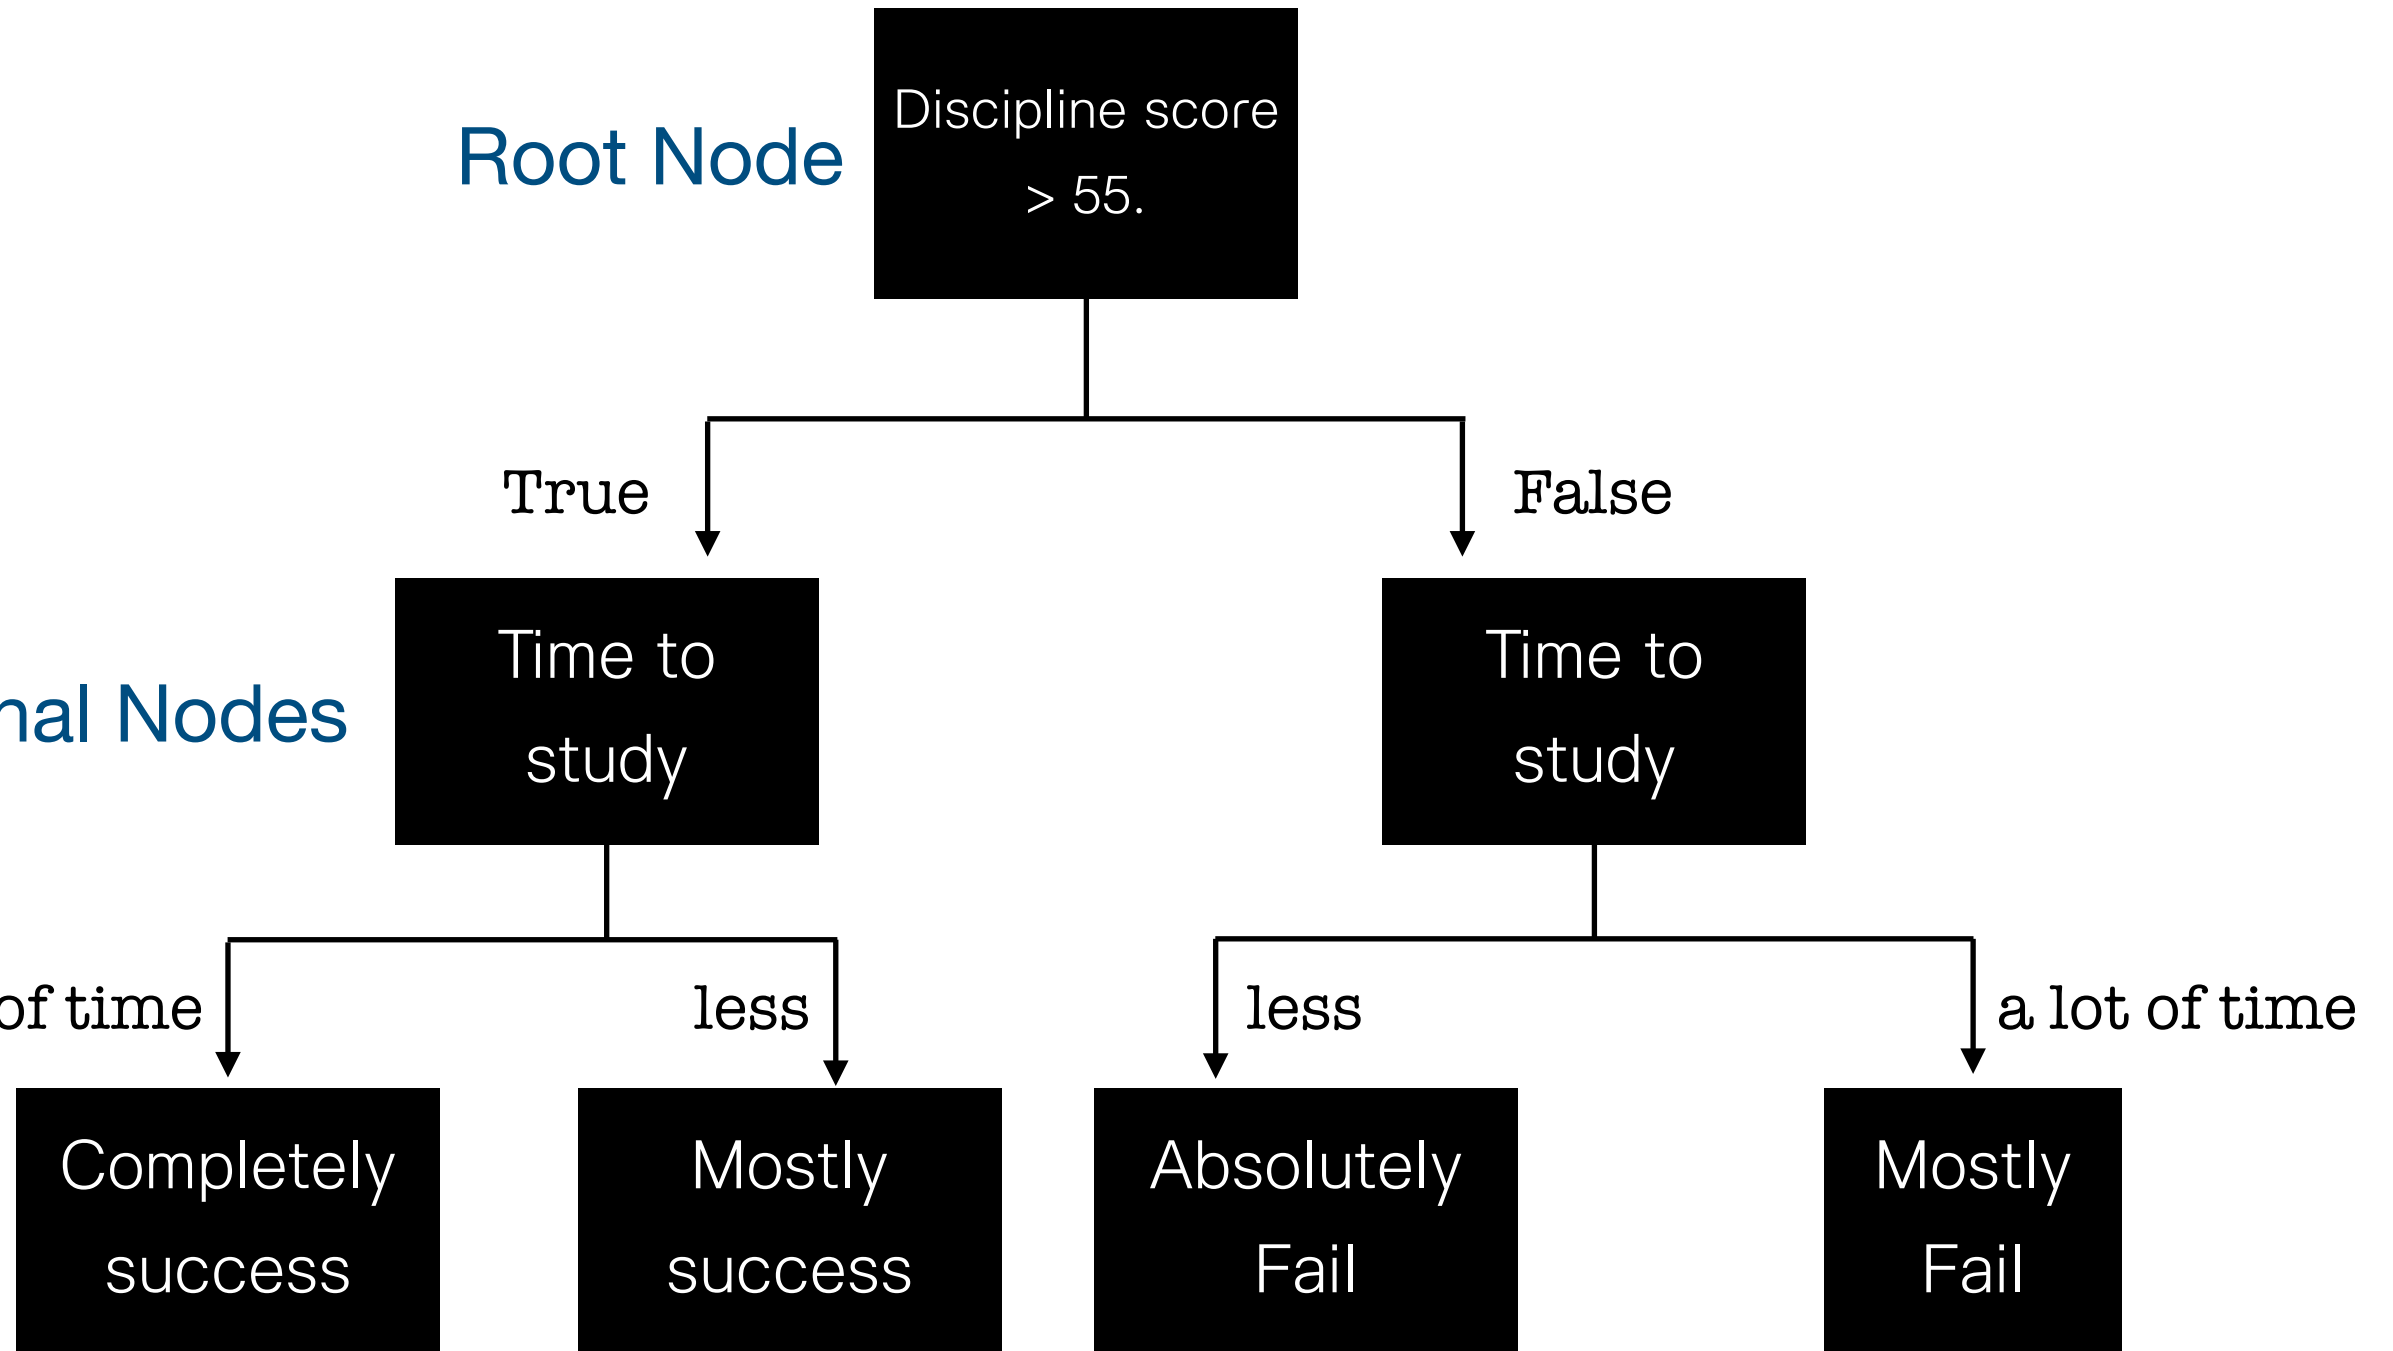
Completely  
success

Mostly  
success

Absolutely  
Fail

Mostly  
Fail

Leaf Nodes



# Advantages

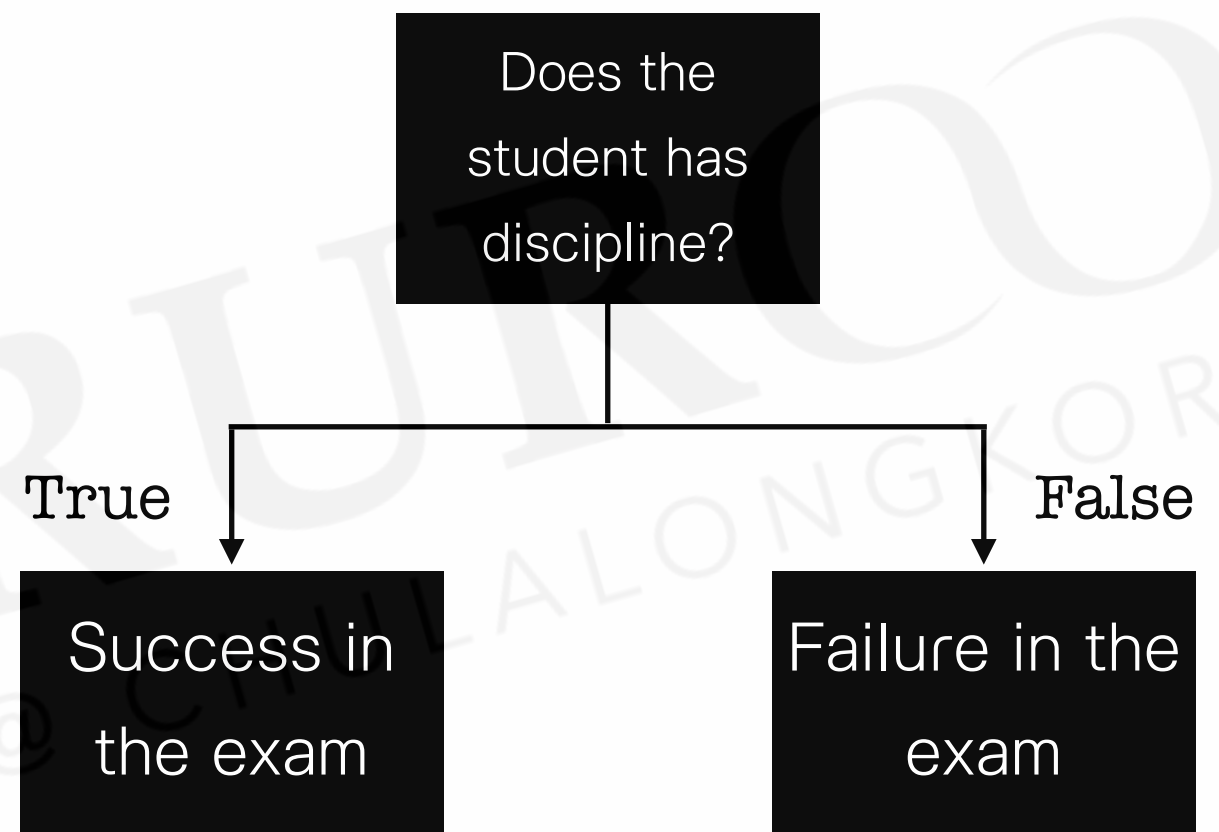
- Simple to understand, interpret, visualize
- Can handle both numerical and categorical features
  - No need to standardised or normalised
  - No need to create dummy variables
- Can handle missing data
- Robust to outliers, then trees require little or no data preparation
- Can model non-linearity in the data
- Trees can handle large datasets

# Disadvantages

- Large trees can be hard to interpret.
- Trees models tend to perform high variance. Hence the models need to be tuned up by choosing the most appropriate hyperparameter.

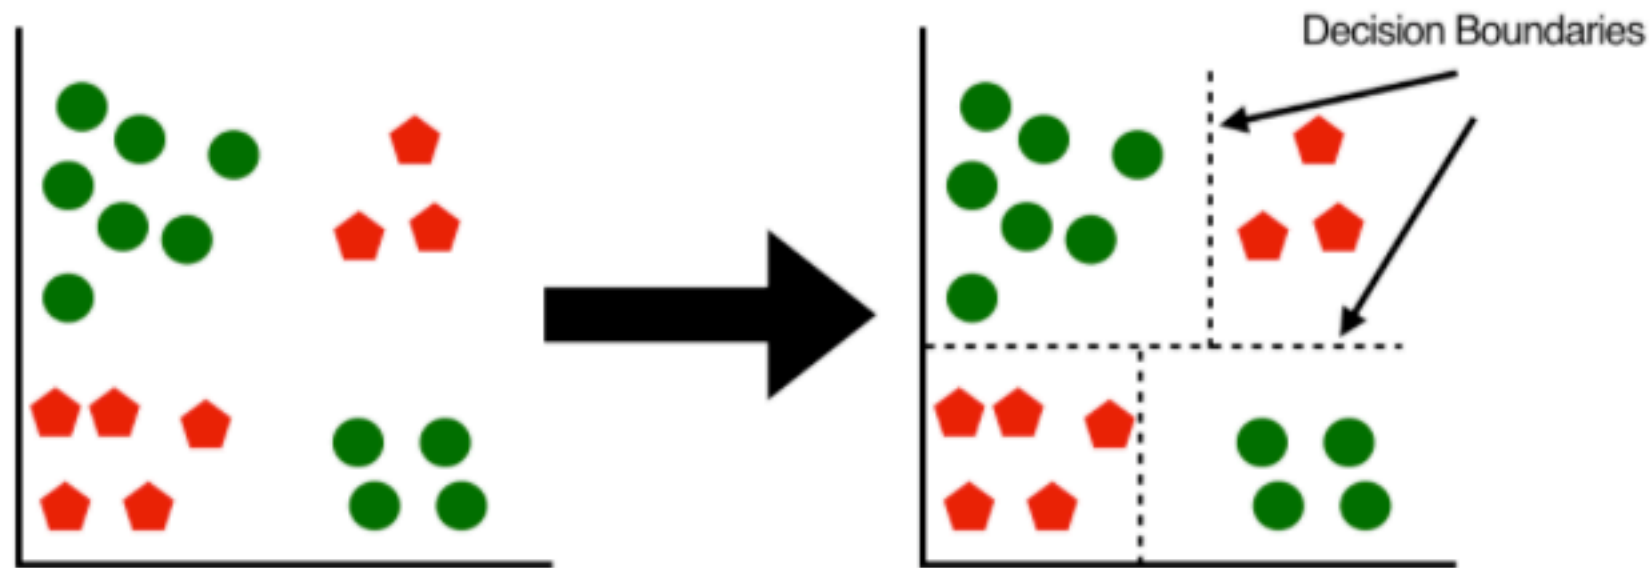
# Tree-based models

- Classification trees
- Regression trees
- Bagged Trees
- Random Forests
- Boosted Trees (GBM)



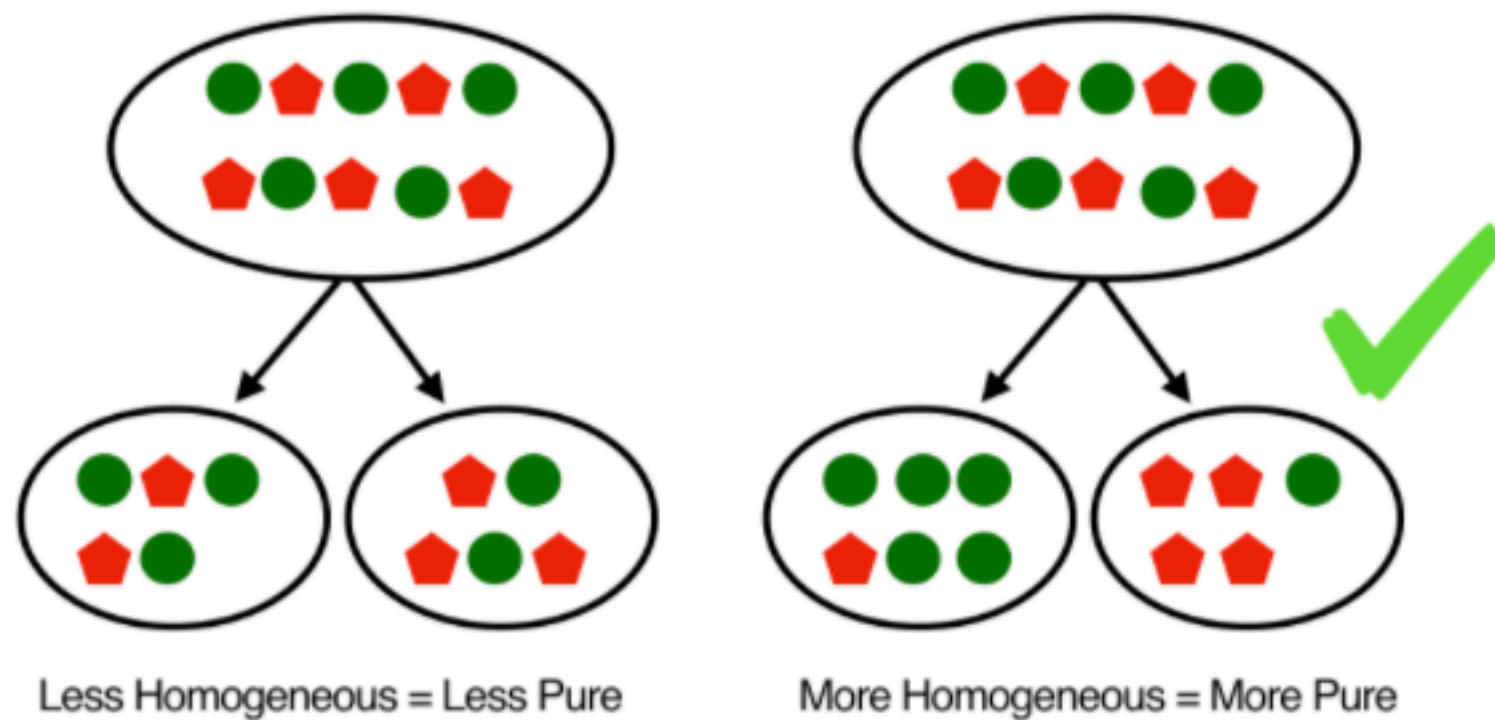
# Classification Trees

Split the data into "pure" regions



- Decision tree model makes classification decision based on the **decision boundaries**.
- The example above, there are 100% pure sub-region (only one class for each region)

# How to determine the best split?

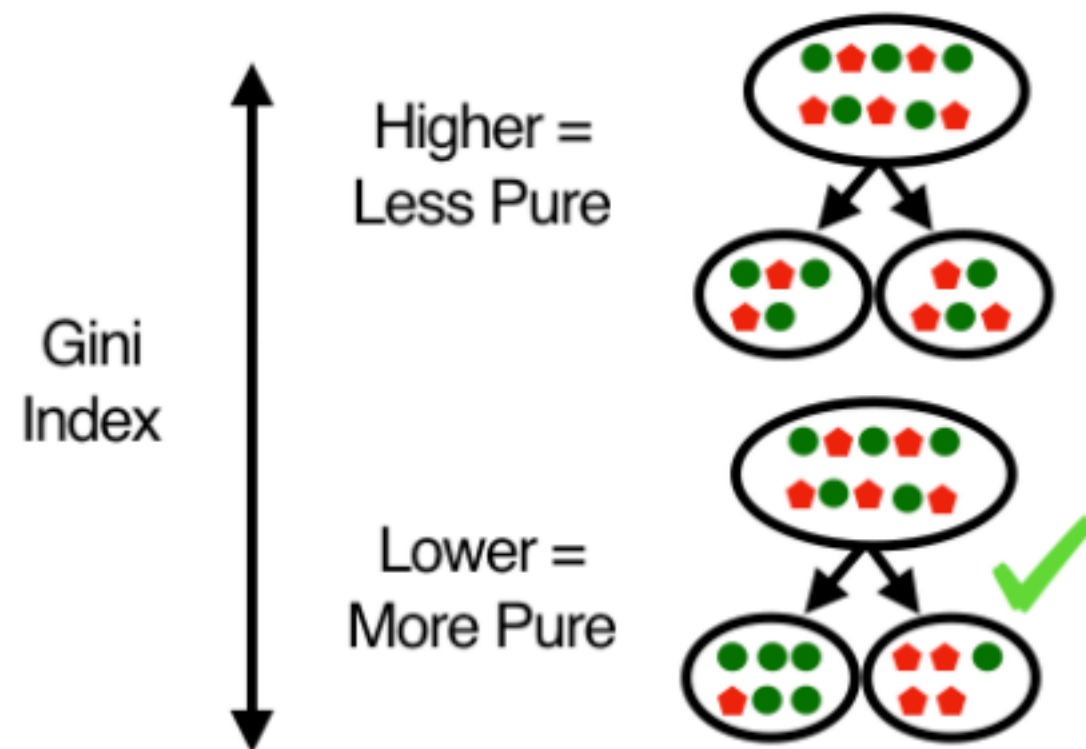


Goal of classification tree is to partition data at a node into subset that are as pure as possible.



# Impurity measure

- Measure of a node specifies how mixed the resulting subsets are.
- We want to find a node specifies that minimised the impurity measure
- Common impurity measures
  - Entropy
  - Gini impurity
  - Misclassification rate



Show some calculation

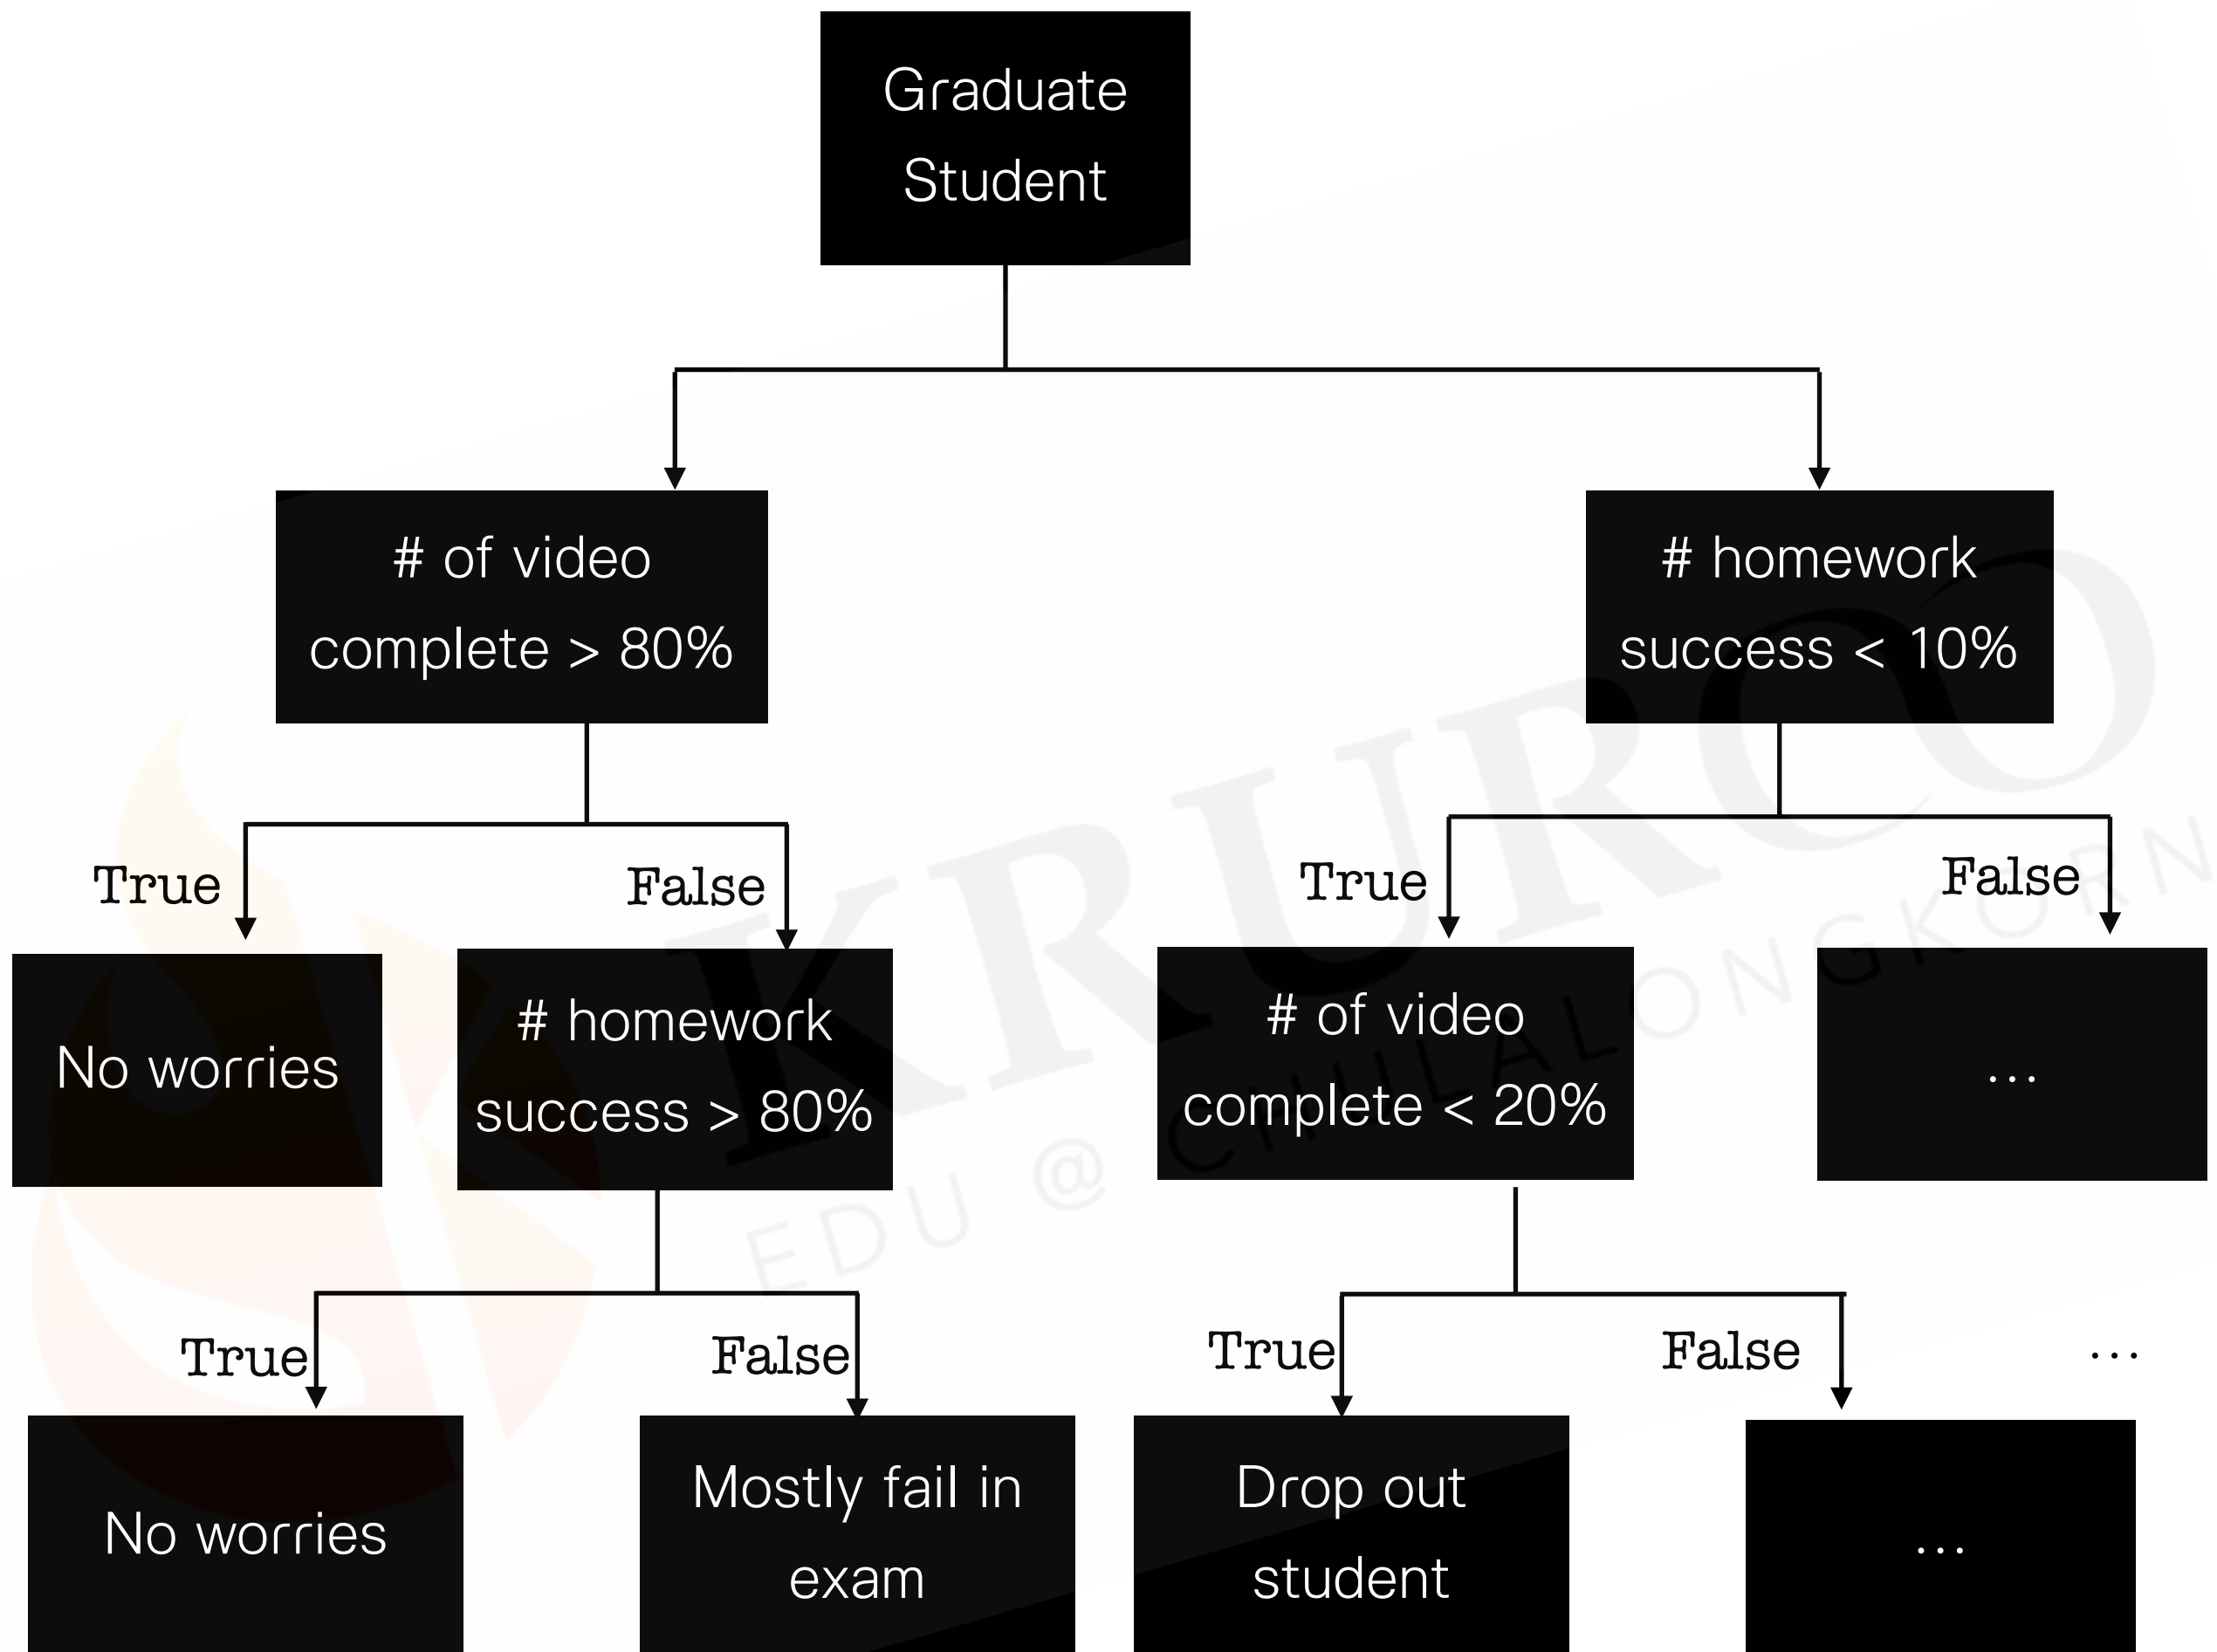
# Building up the model

1. Split data into train and test dataset. (80/20)
2. Train classification tree using train dataset

```
install.packages("rpart")  
install.packages("rpart.plot")  
library(rpart)  
library(rpart.plot)
```

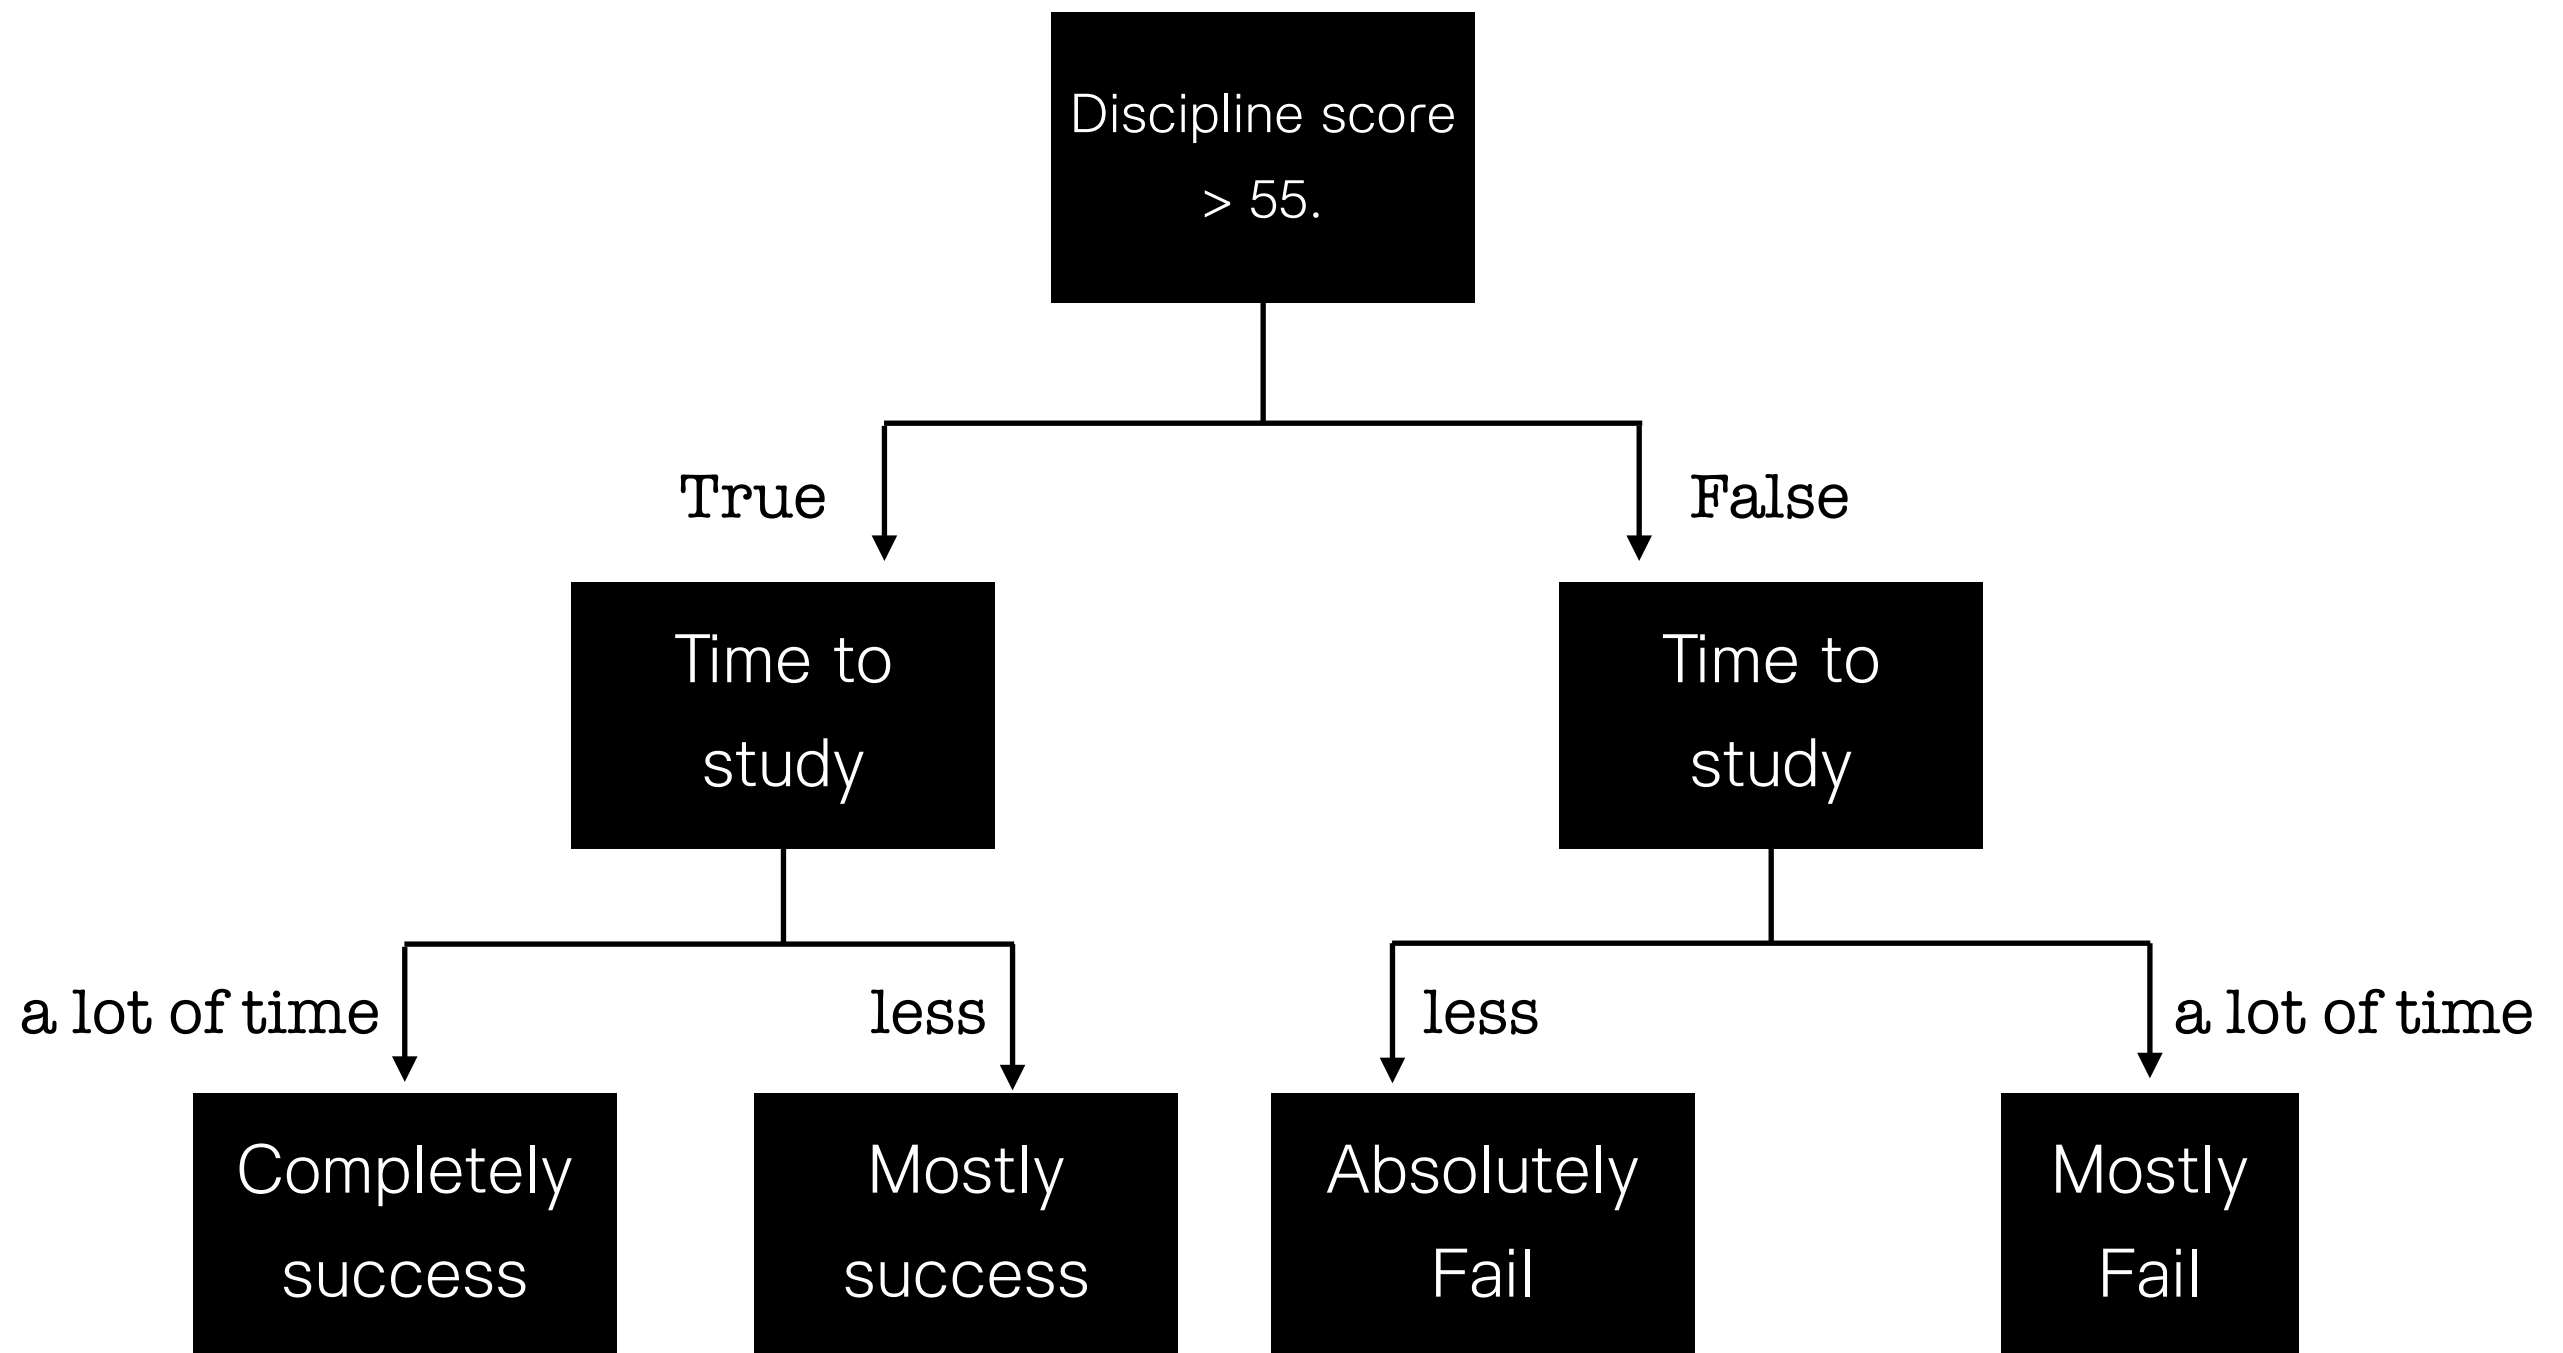
```
fit<-rpart(formula,data,method="class",  
           control = rpart.control(cp, minsplit,maxdepth,xval),  
           parms = list(split="information"))
```

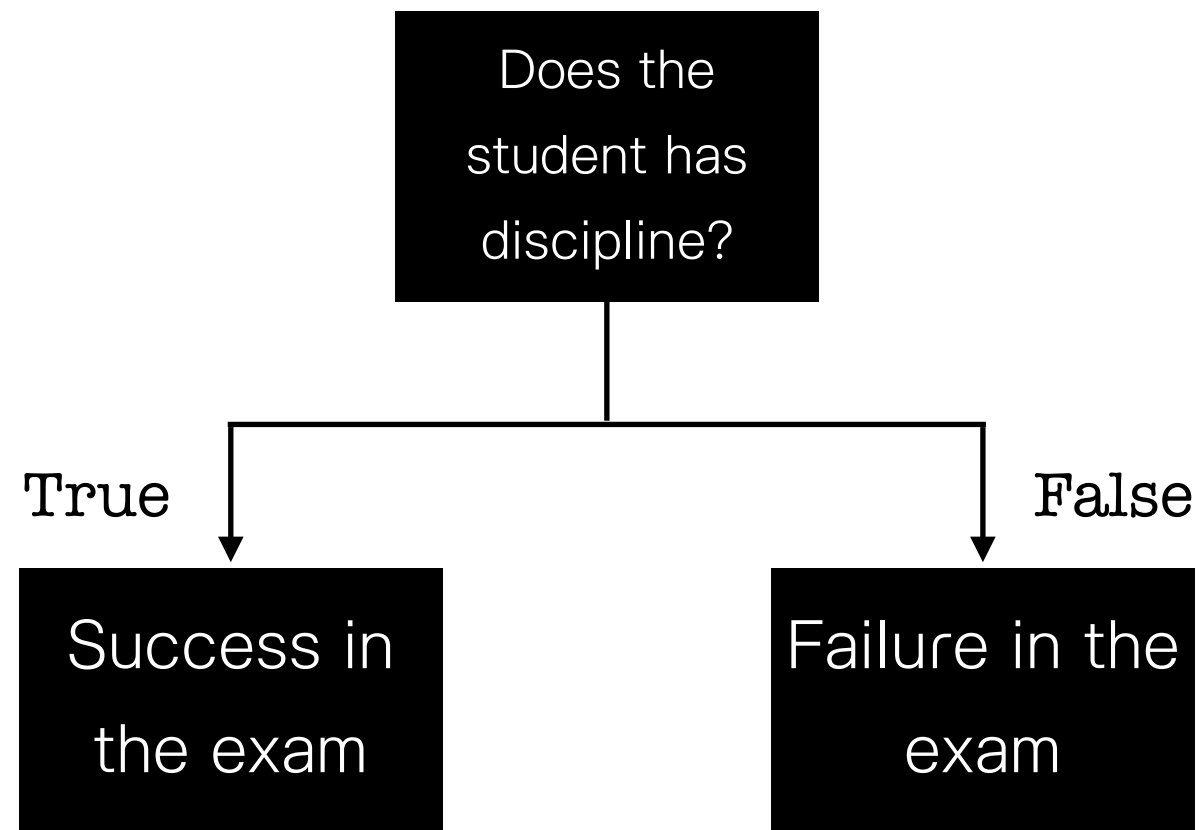
```
# For classification splitting, the list can contain any of: the vector of prior  
# probabilities (component prior), the loss matrix (component loss) or the splitting  
# index (component split). The priors must be positive and sum to 1. The loss matrix  
# must have zeros on the diagonal and positive off-diagonal elements. The splitting  
# index can be gini or information. The default priors are proportional to the data  
# counts, the losses default to 1, and the split defaults to gini.
```



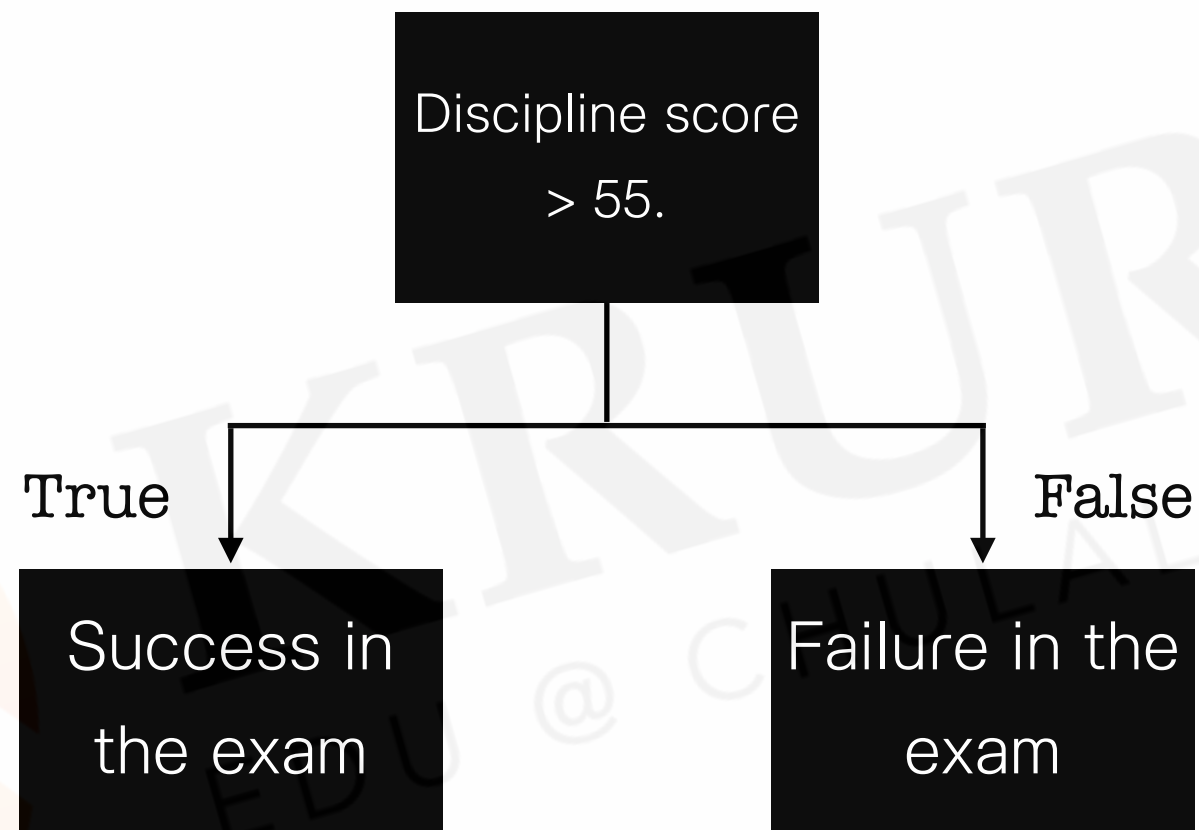
# Apple

- Discipline = 70
- No time to study

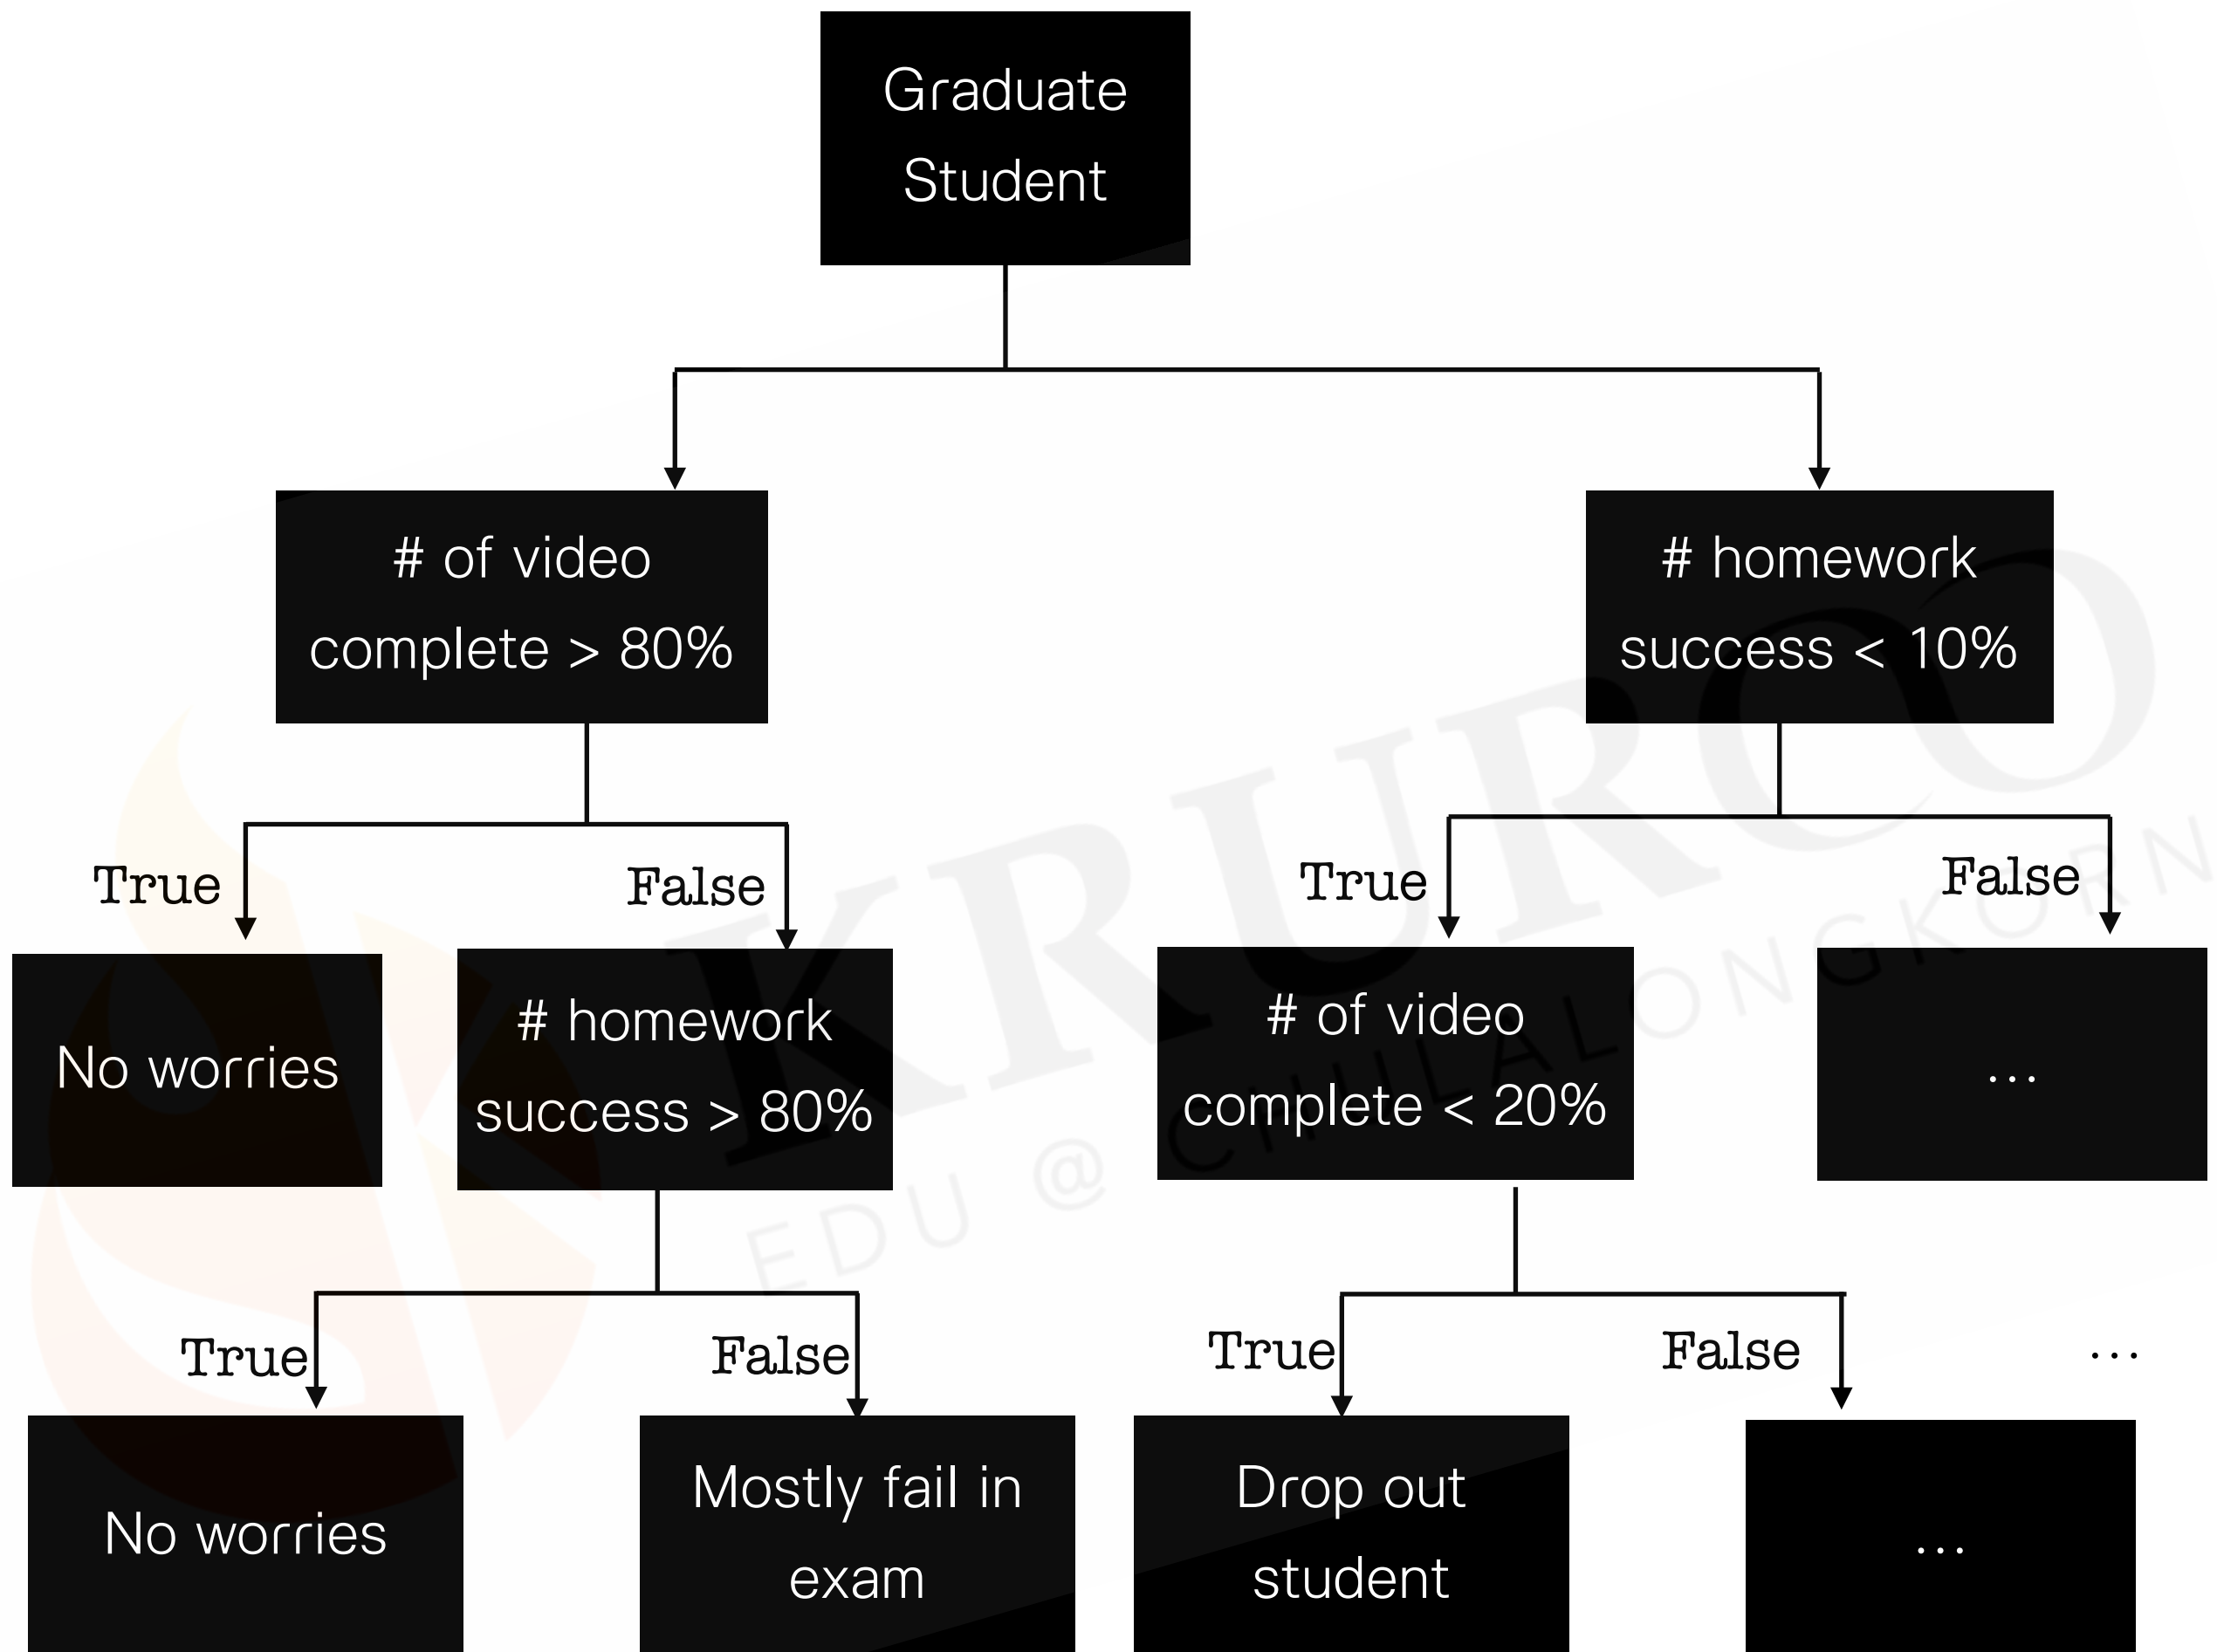




Decision tree ask a question, and then classifies the subject based on the answer.



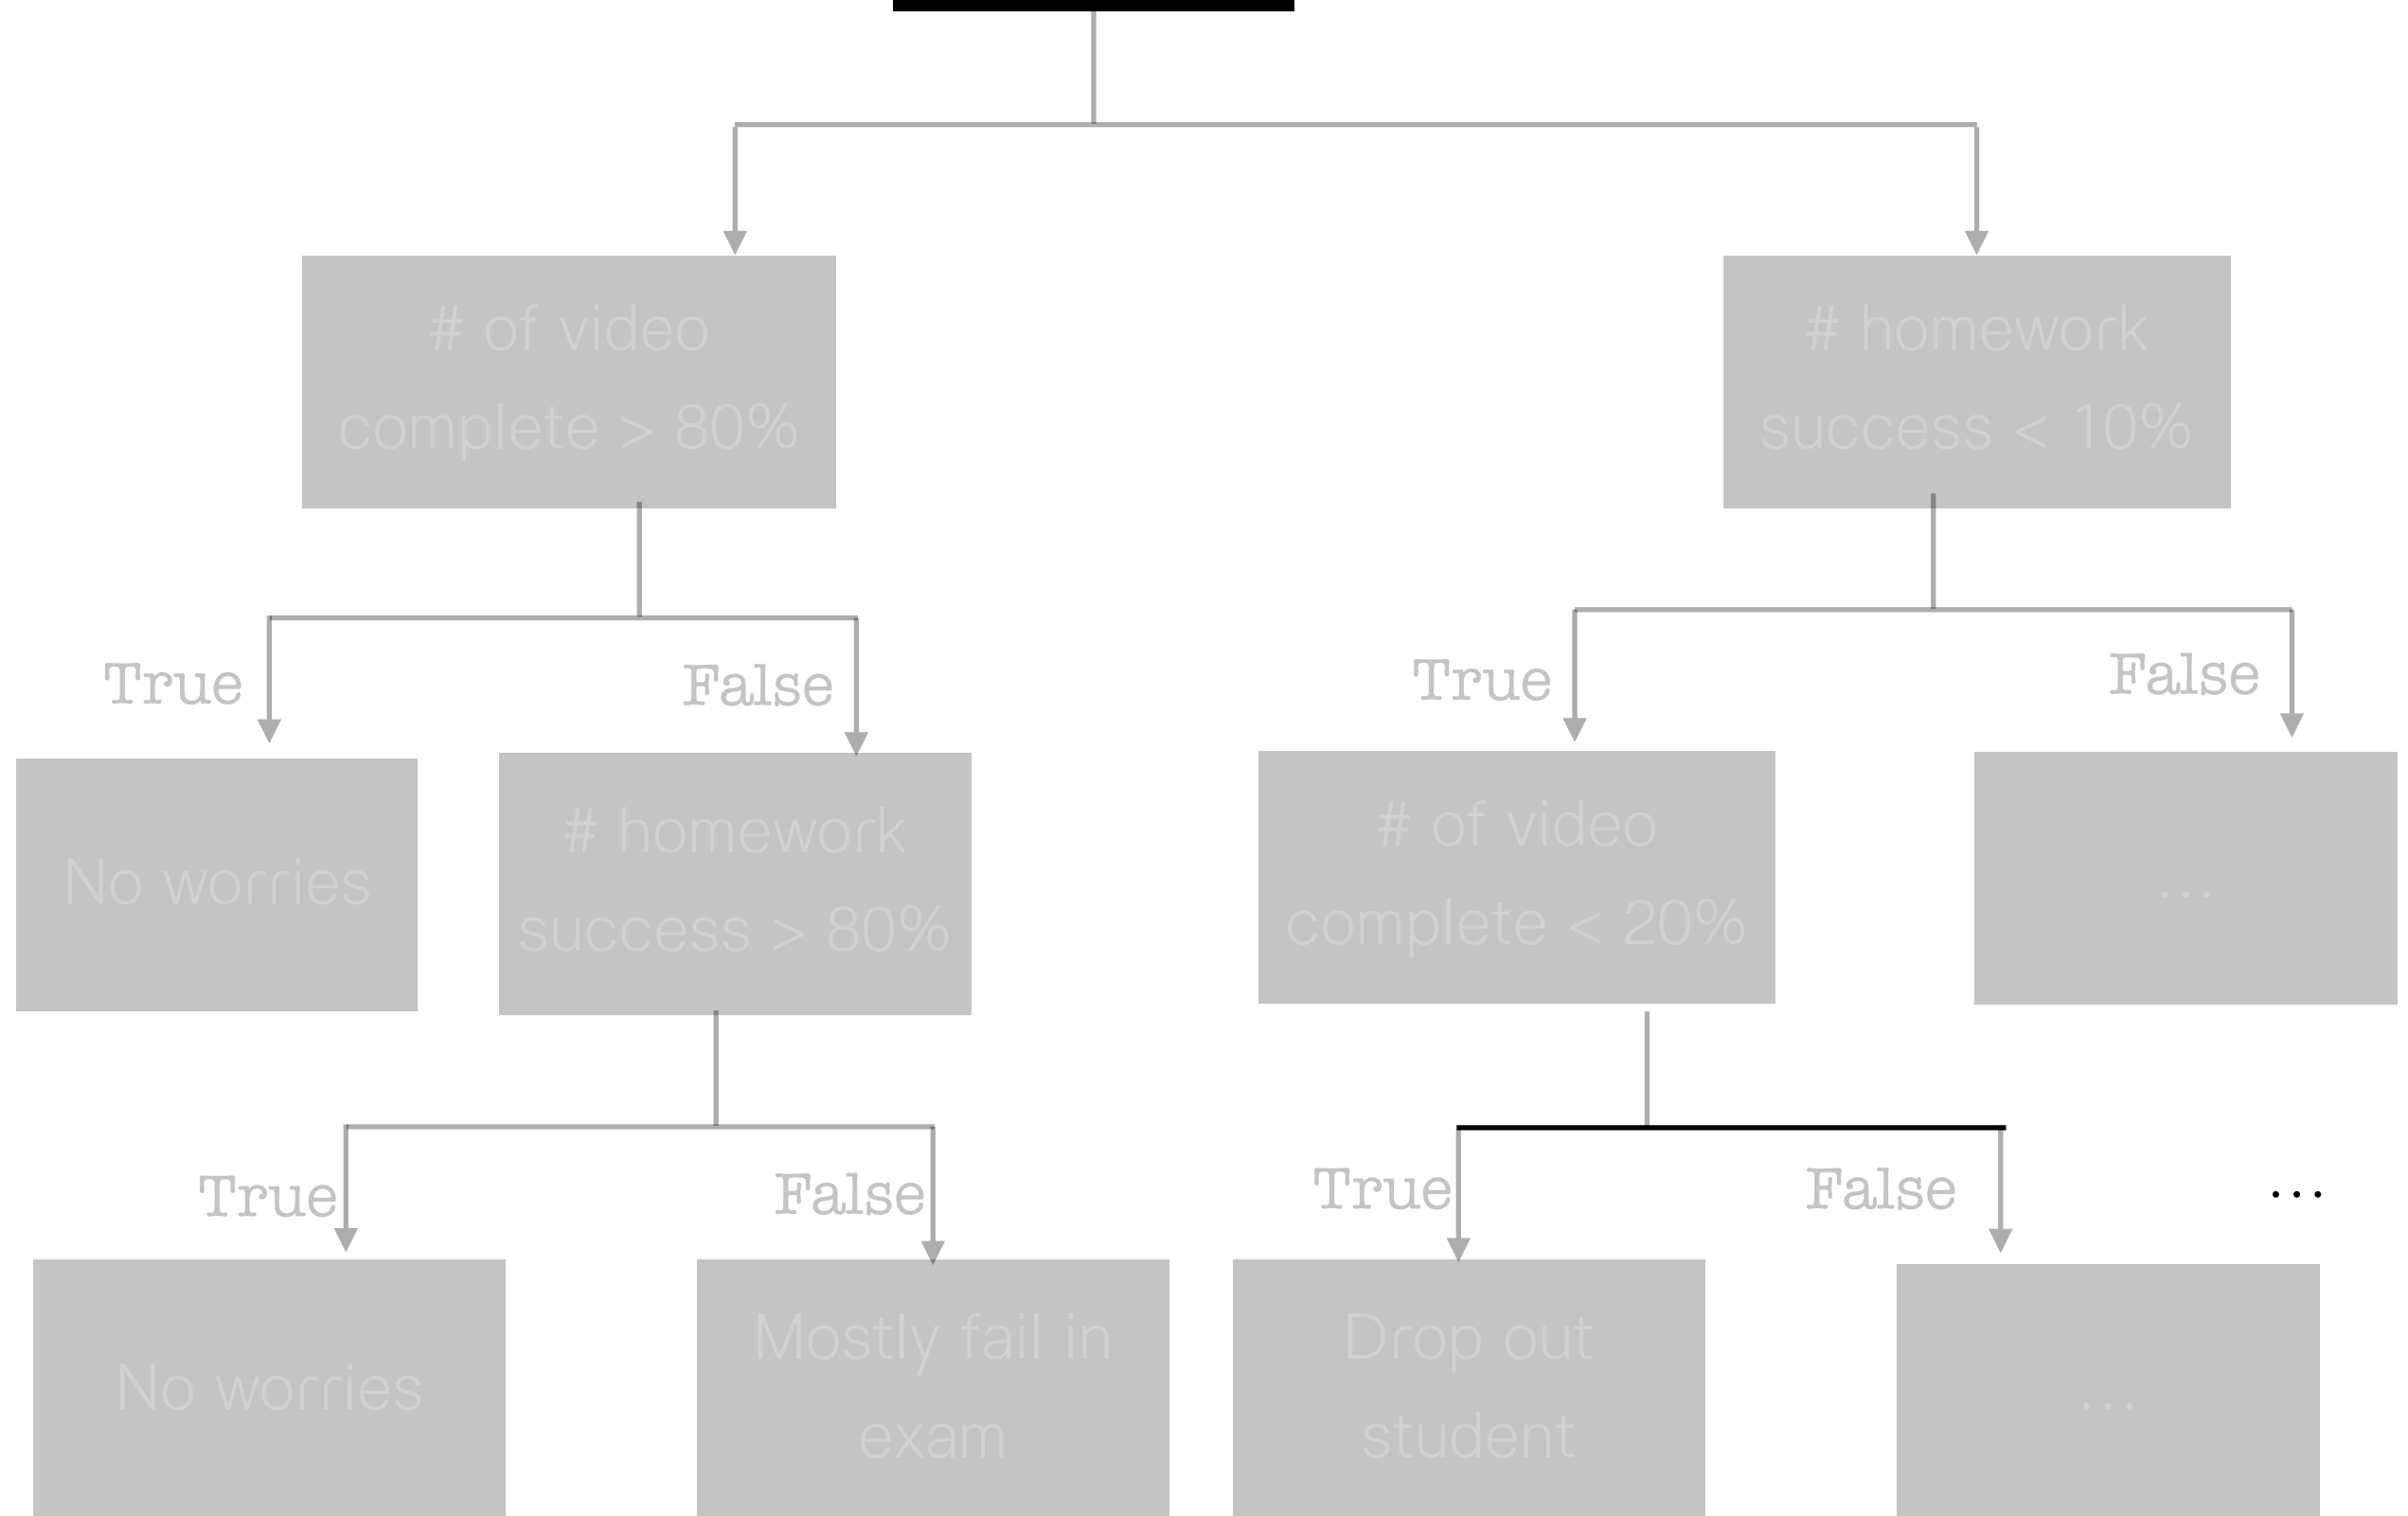
Decision tree for numeric data.





# Root Node

Graduate  
Student



Graduate  
Student

## Internal Nodes

# of video  
complete > 80%

# homework  
success < 10%

True

False

No worries

# homework  
success > 80%

True

False

# of video  
complete < 20%

...

True

False

No worries

Mostly fail in  
exam

True

False

Drop out  
student

...

...

